

CS2100

COMPUTER ORGANISATION

<http://www.comp.nus.edu.sg/~cs2100/>

Lecture #20

Pipelining

Part I: Introduction



NUS
National University
of Singapore

School of
Computing

Lecture #20: Pipelining I

1. Introduction
2. MIPS Pipeline Stages
3. Pipeline Datapath
4. Pipeline Control
5. Pipeline Performance


1. Introduction: Inspiration

Assembly Line


Simpler station tasks → more cars per hour.
Simple tasks take less time, clock is faster.



1. Introduction: Inspiration





EXAMPLE BASIC PRODUCTION LINE

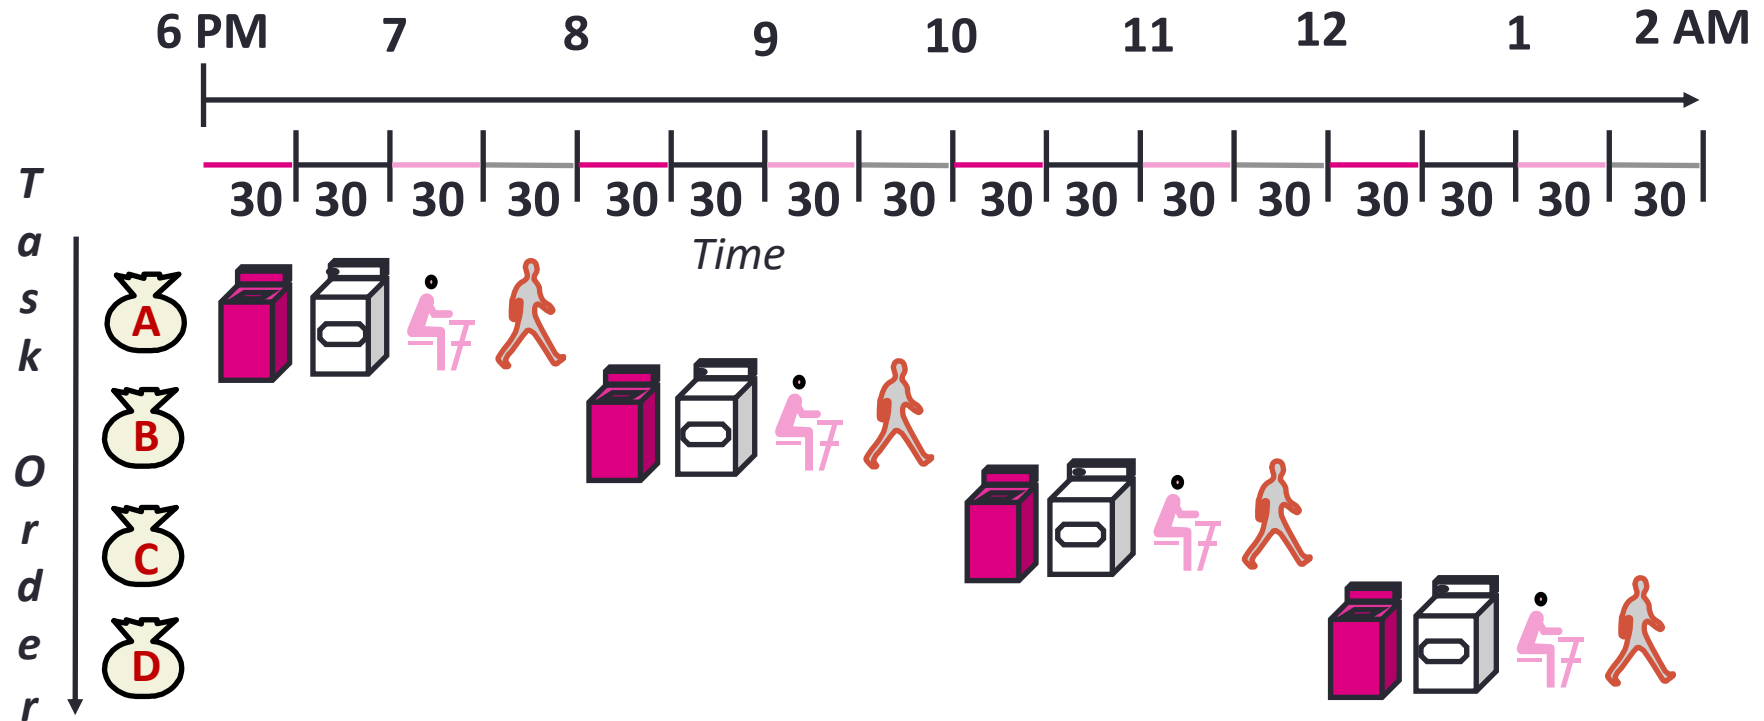


1. Introduction: Laundry

- **Ann, Brian, Cathy, Dave** each have one load of clothes to wash, dry, fold and stash
- **Washer** takes 30 minutes
- **Dryer** takes 30 minutes
- **“Folder”** takes 30 minutes
- **“Stasher”** takes 30 minutes to put clothes into drawers

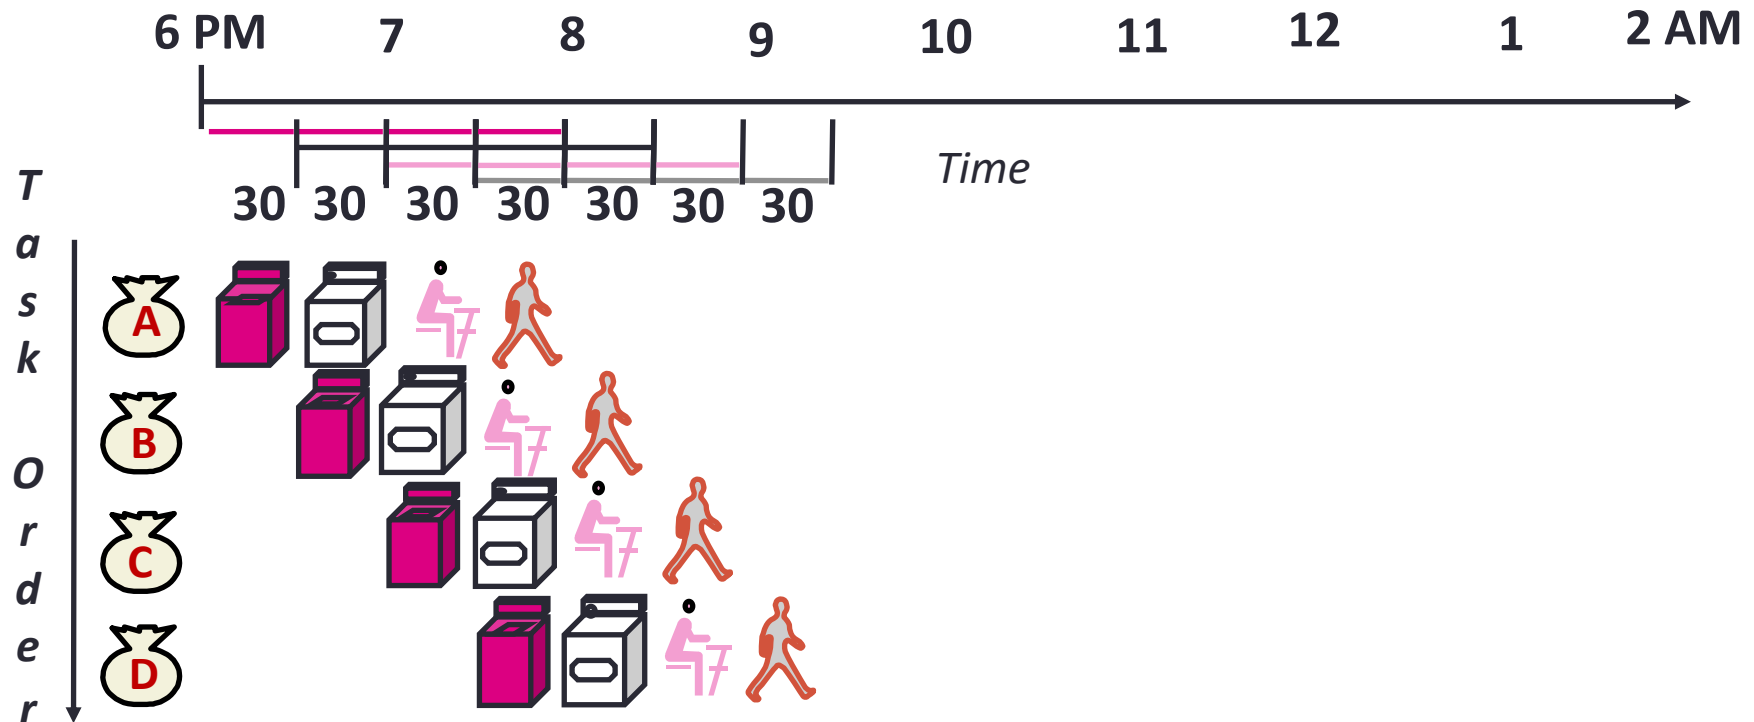


1. Introduction: Sequential Laundry




- Sequential laundry takes 8 hours for 4 loads
- Steady state: 1 load every 2 hours
- If they learned **pipelining**, how long would laundry take?

1. Introduction: Pipelined Laundry



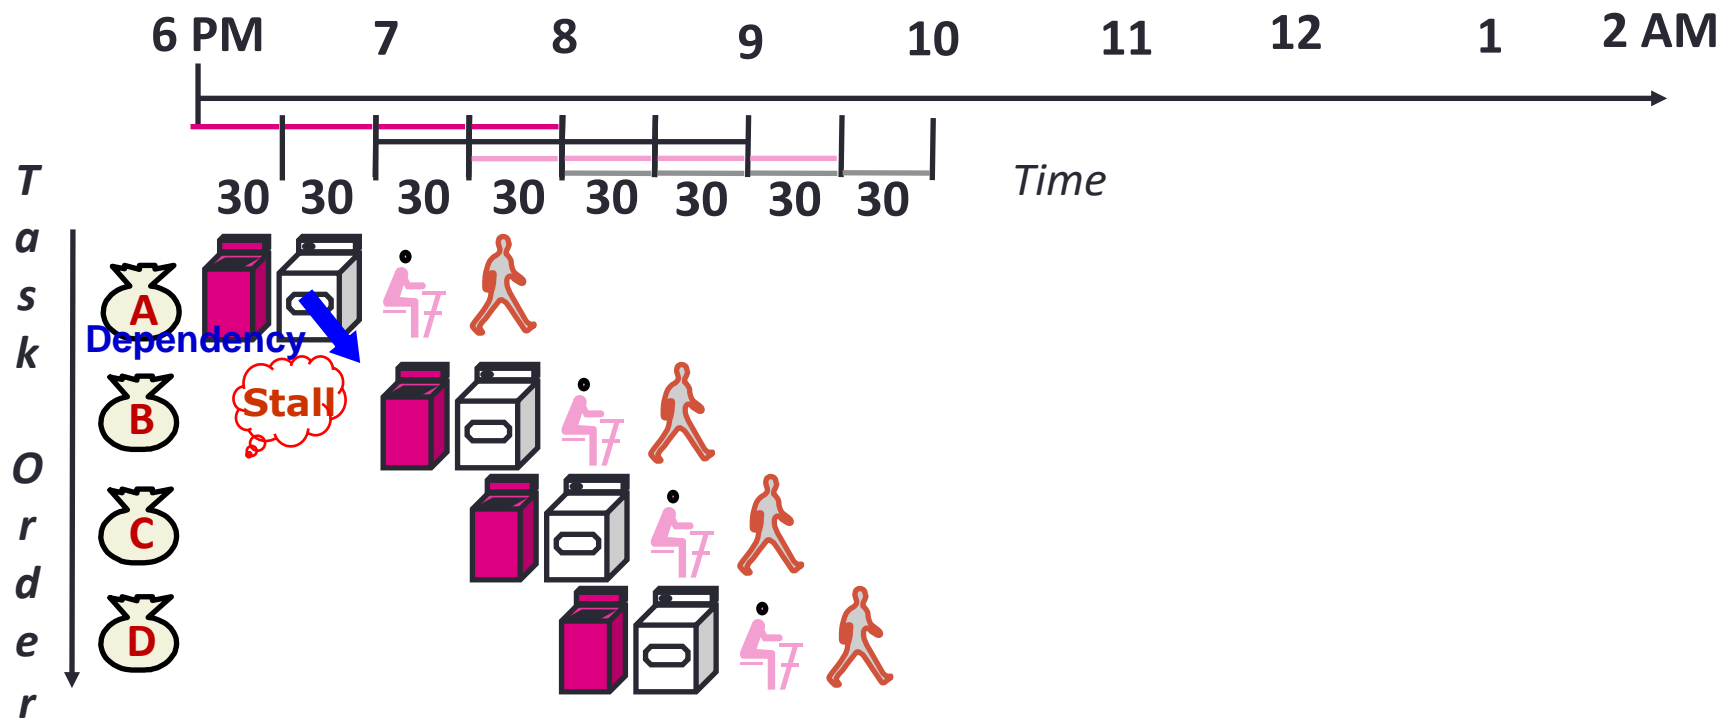
- Pipelined laundry takes 3.5 hours for 4 loads!
- Steady state: 1 load every 30 minutes
- Potential speedup = $2 \text{ hr} / 30 \text{ min} = 4$ (no. of stages)
- Time to fill pipeline takes 2 hours \rightarrow speedup \downarrow

1. Introduction: What If: Slow Dryer



- **Pipelined laundry now takes 5.5 hours!**
- Steady state: One load every 1 hour (dryer speed)
- **Pipeline rate is limited by the slowest stage**

1. Introduction: What If: Dependency



- Brian is using the laundry for the first time; he wants to see the outcome of one wash + dry cycle first before putting in his clothes
- **Pipelined laundry now takes 4 hours**

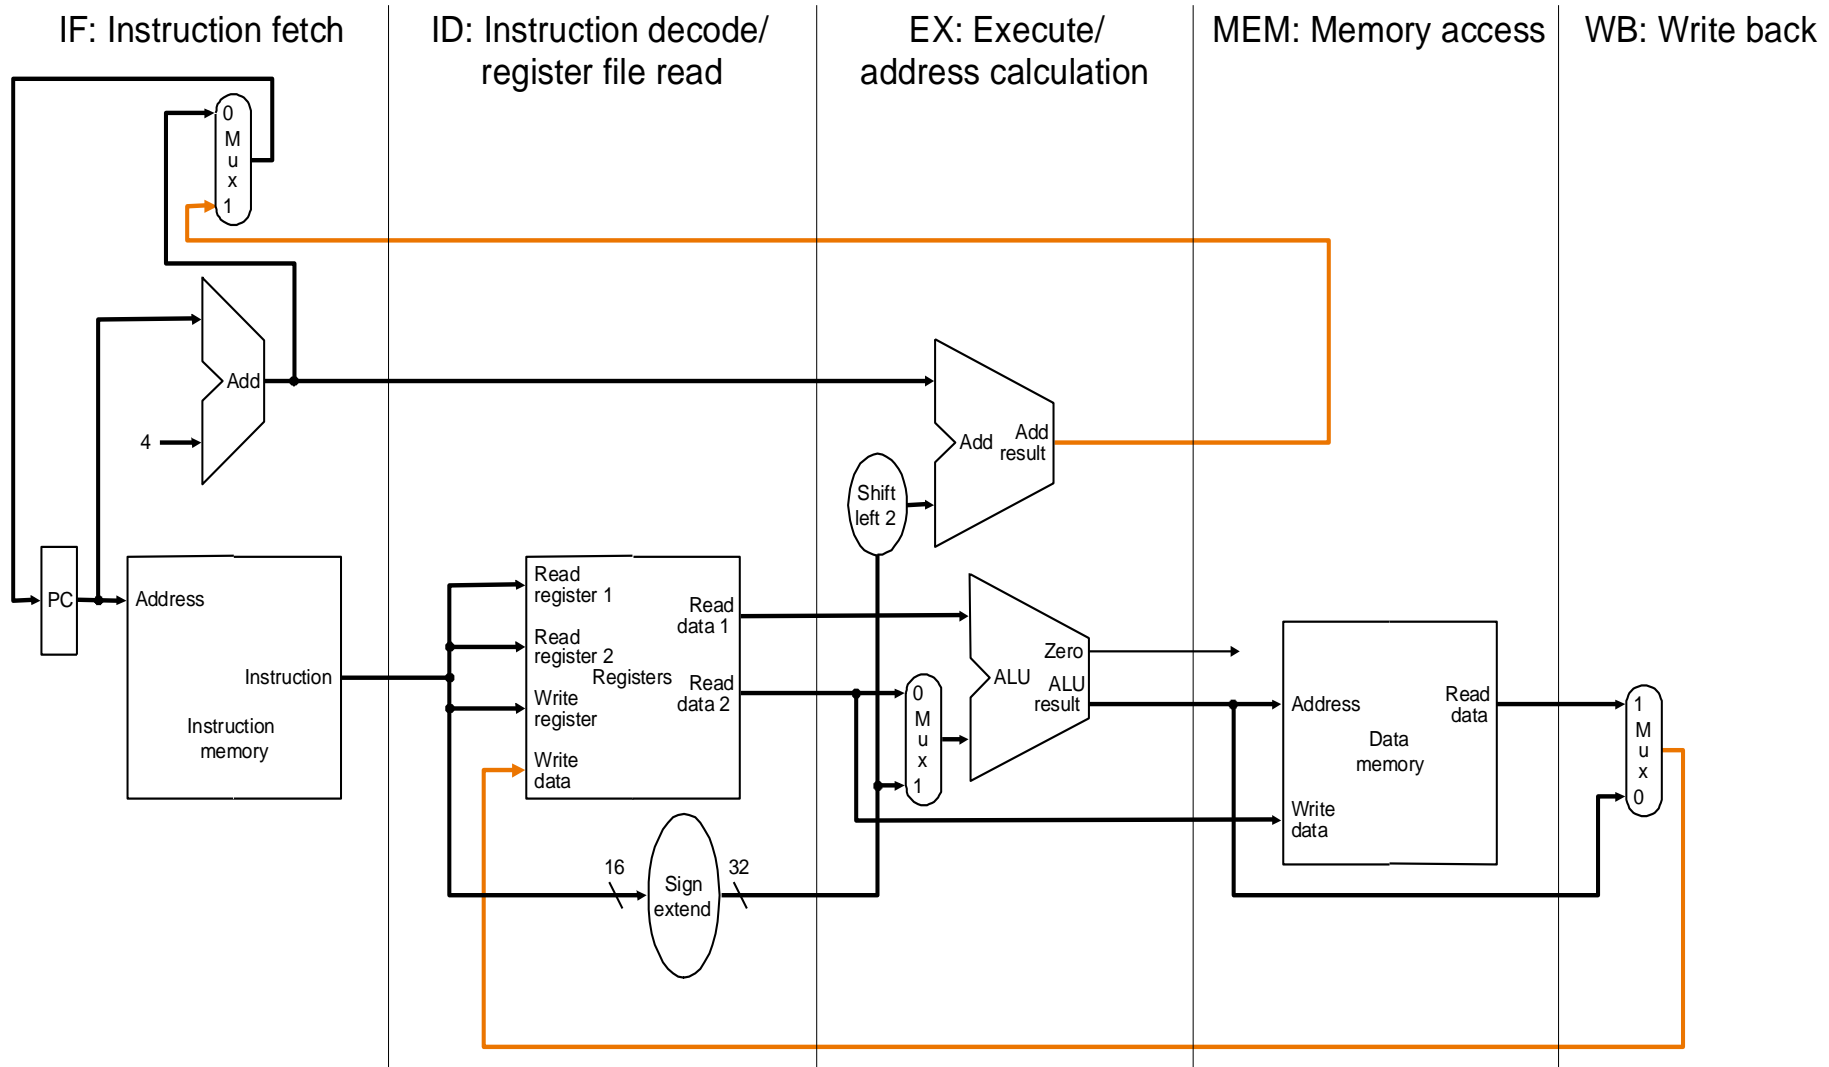
1. Introduction: Pipelining Lessons

- Pipelining doesn't help **latency** of single task:
 - It helps the **throughput** of entire workload
- **Multiple** tasks operating simultaneously using different resources
- Possible delays:
 - Pipeline rate limited by **slowest** pipeline stage
 - Stall for **dependencies**

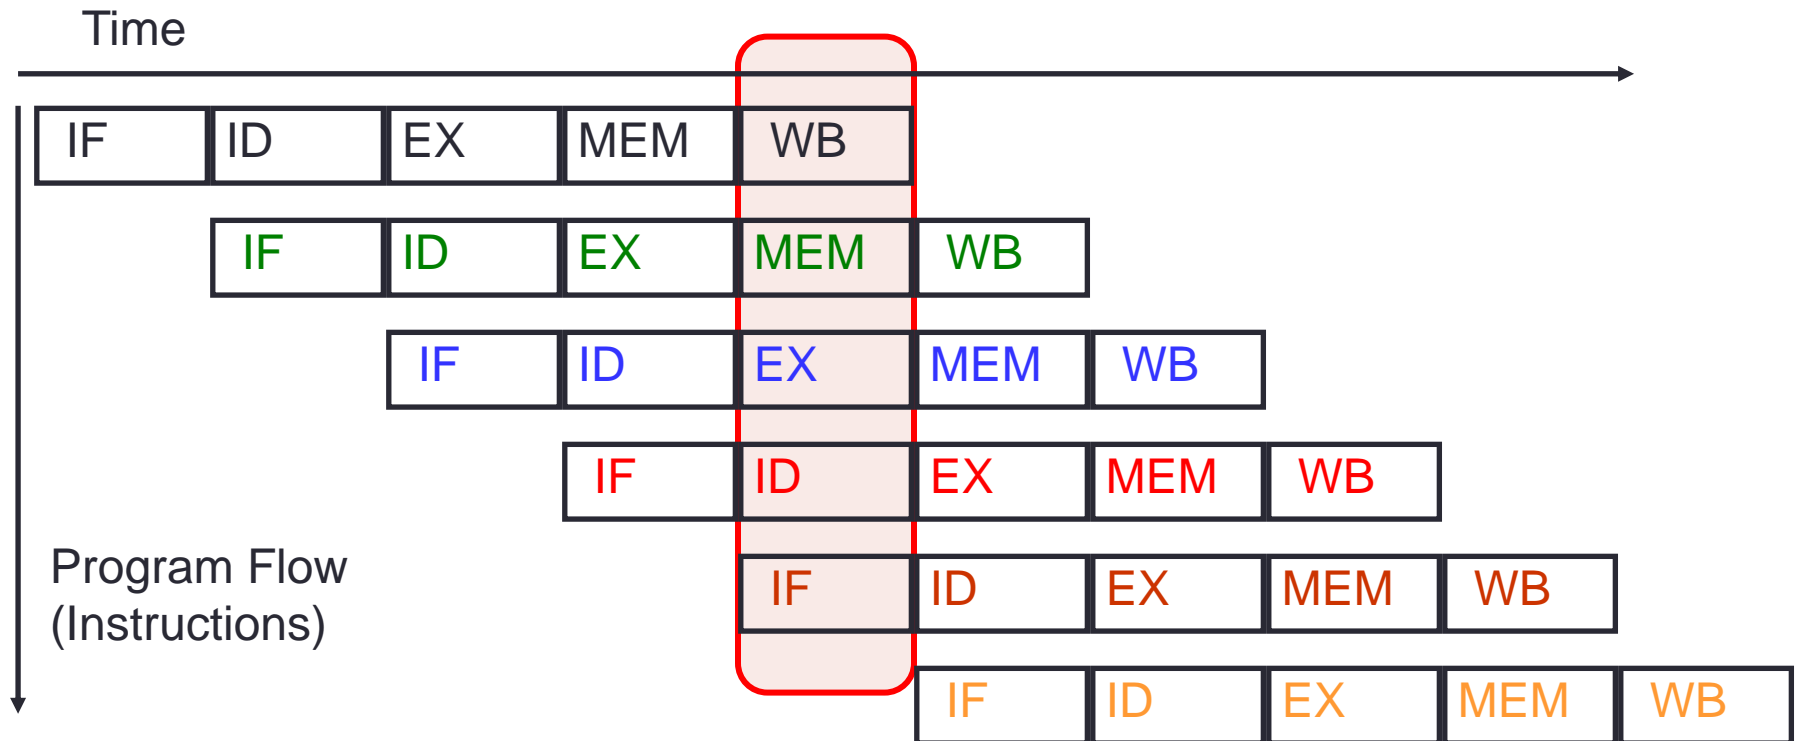
2. MIPS Pipeline Stages (1/2)

- **Five** Execution Stages
 - **IF**: Instruction Fetch
 - **ID**: Instruction Decode and Register Read
 - **EX**: Execute an operation or calculate an address
 - **MEM**: Access an operand in data memory
 - **WB**: Write back the result into a register
- **Idea:**
 - **Each execution stage takes 1 clock cycle**
 - General flow of data is from one stage to the next
- **Exceptions:**
 - Update of PC and write back of register file – more about this later...

2. MIPS Pipeline Stages (2/2)



2. Pipelined Execution: Illustration



Several instructions
are in the pipeline
simultaneously!

3. MIPS Pipeline: Datapath (1/3)

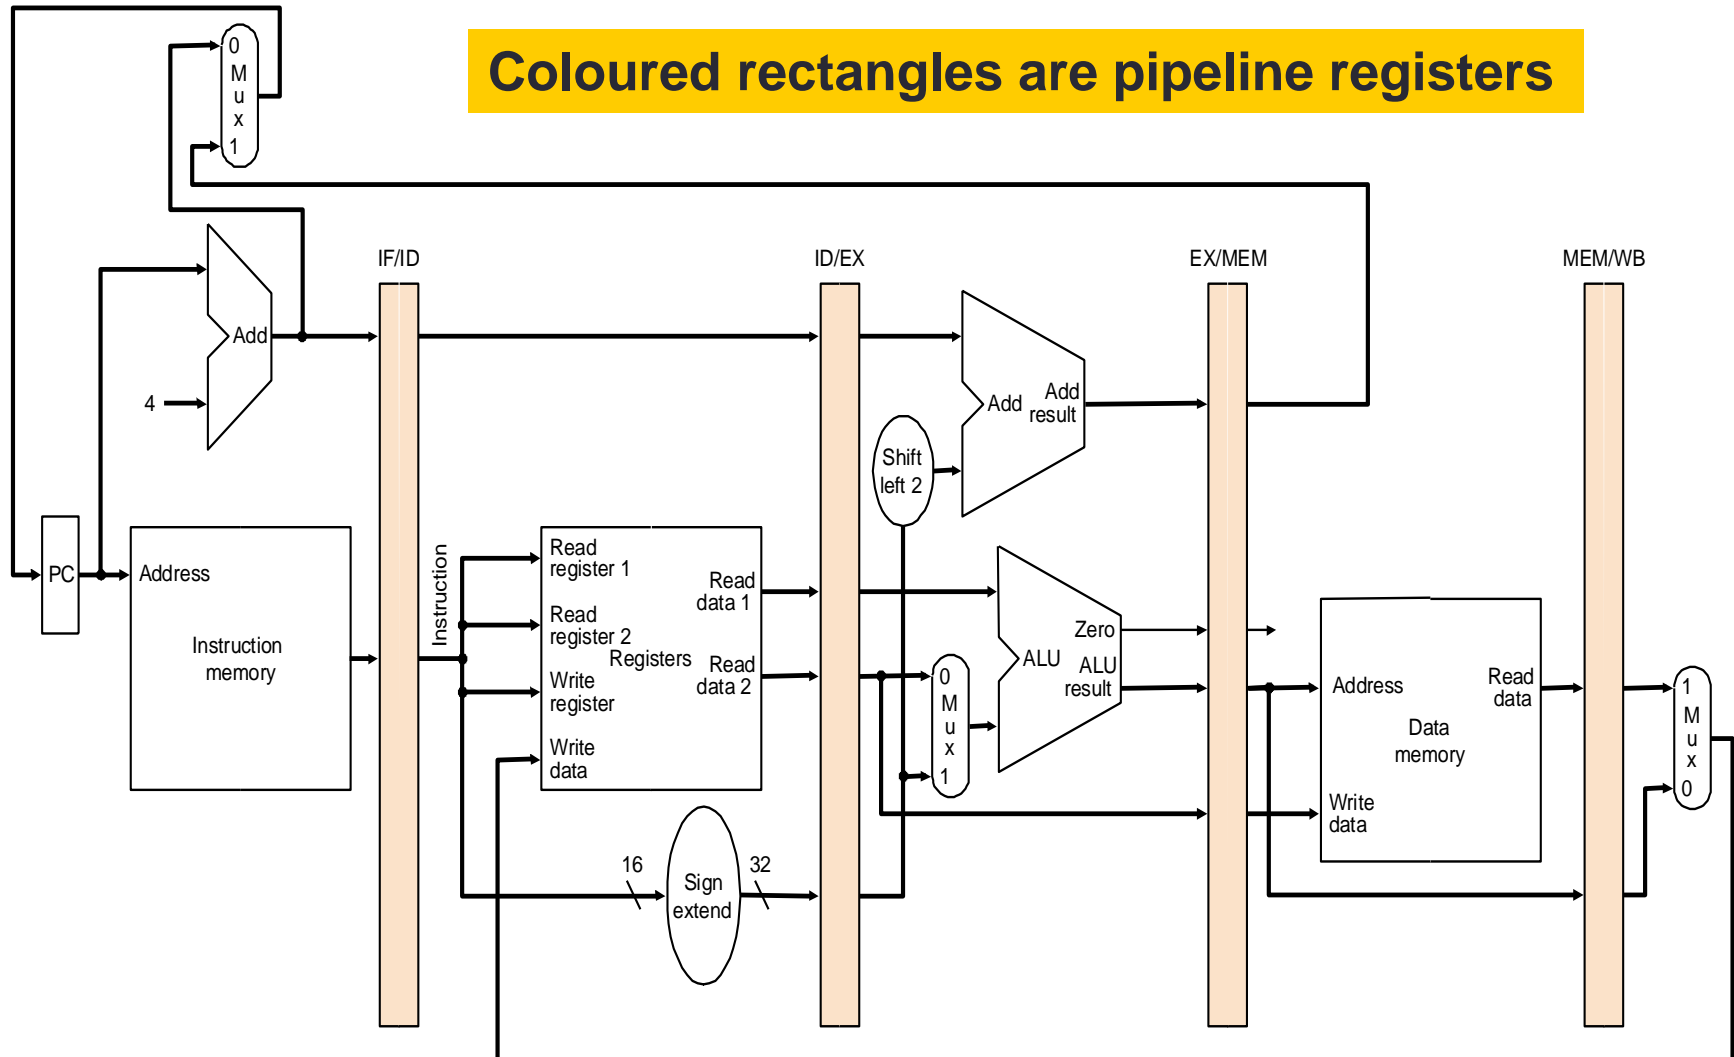
- **Single-cycle** implementation:
 - Update all state elements (PC, register file, data memory) at the end of a clock cycle
- **Pipelined** implementation:
 - One cycle per pipeline stage
 - Data required for each stage needs to be stored separately (why?)

3. MIPS Pipeline: Datapath (2/3)


- Data used by **subsequent instructions**:
 - Store in programmer-visible state elements: **PC**, register file and memory
- Data used by **same instruction** in later pipeline stages:
 - Additional registers in datapath called **pipeline registers**
 - **IF/ID**: register between **IF** and **ID**
 - **ID/EX**: register between **ID** and **EX**
 - **EX/MEM**: register between **EX** and **MEM**
 - **MEM/WB**: register between **MEM** and **WB**
- Why no register at the end of **WB** stage?

3. MIPS Pipeline: Datapath (3/3)

Coloured rectangles are pipeline registers

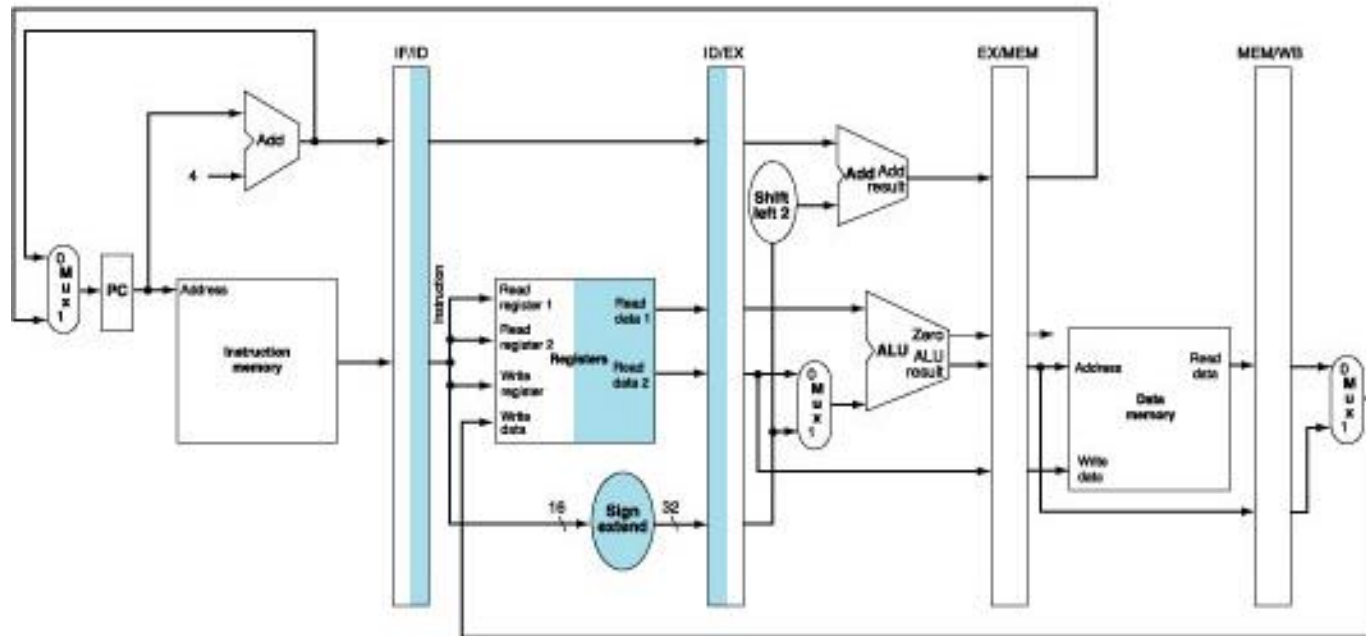


3. Pipeline Datapath: **IF** Stage



- At the end of a cycle, **IF/ID** receives (stores):
 - Instruction read from InstructionMemory[PC]
 - $PC + 4$
- $PC + 4$
 - Also connected to one of the MUX's inputs (another coming later)

3. Pipeline Datapath: ID Stage




At the beginning of a cycle
IF/ID register supplies:

- ❖ Register numbers for reading two registers
- ❖ 16-bit offset to be sign-extended to 32-bit

At the end of a cycle
ID/EX receives:

- ❖ Data values read from register file
- ❖ 32-bit immediate value
- ❖ $PC + 4$

3. Pipeline Datapath: **EX** Stage



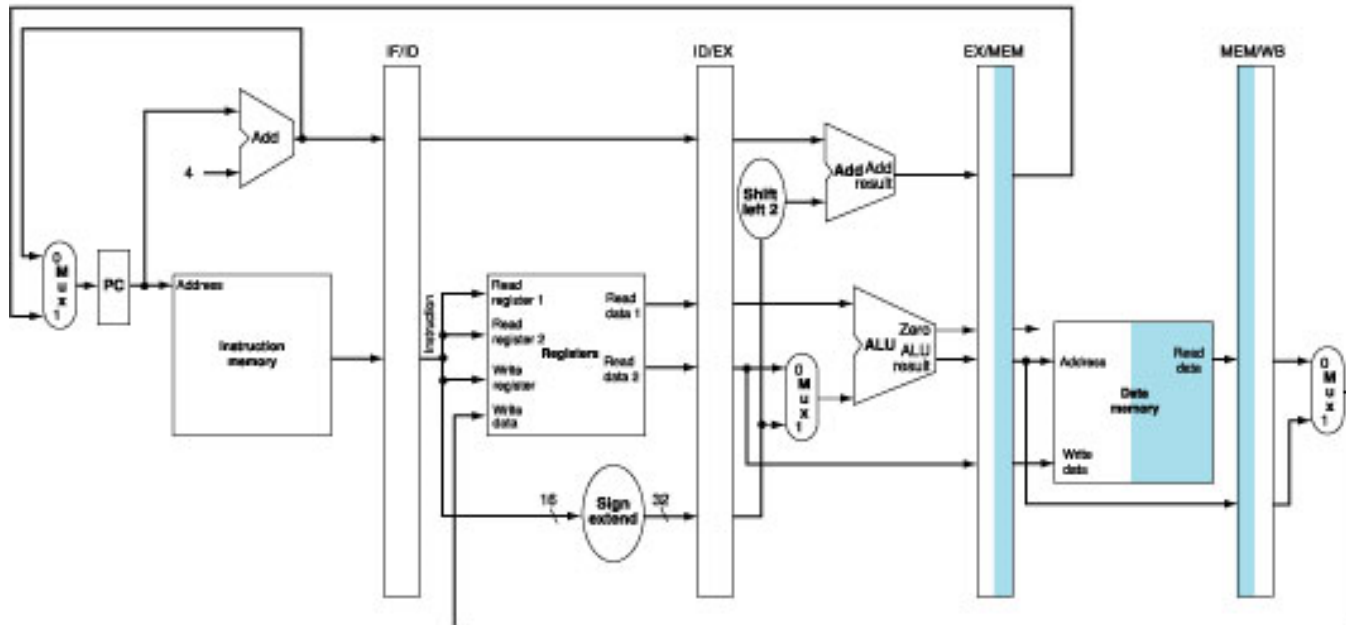
**At the beginning of a cycle
ID/EX register supplies:**

- ❖ Data values read from register file
- ❖ 32-bit immediate value
- ❖ PC + 4

**At the end of a cycle
EX/MEM receives:**

- ❖ (PC + 4) + (Immediate x 4)
- ❖ ALU result
- ❖ isZero? signal
- ❖ Data Read 2 from register file

3. Pipeline Datapath: **MEM** Stage



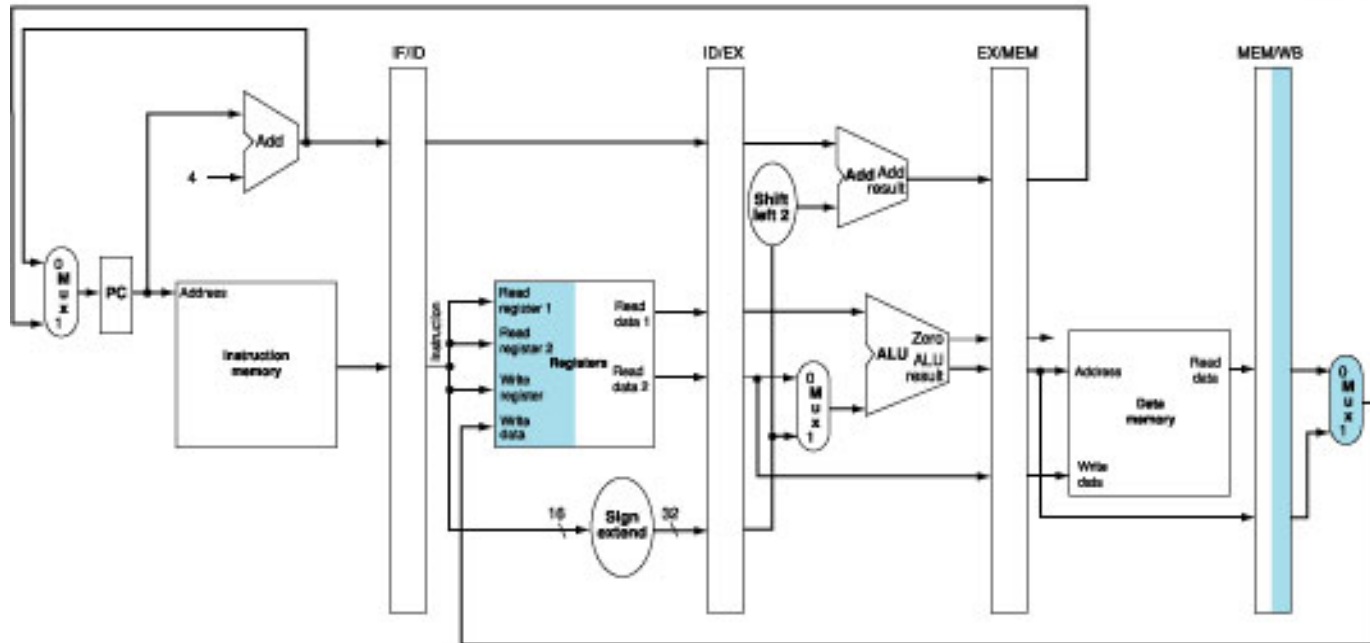
At the beginning of a cycle
EX/MEM register supplies:

- ❖ $(PC + 4) + (Immediate \times 4)$
- ❖ ALU result
- ❖ `isZero?` signal
- ❖ Data Read 2 from register file

At the end of a cycle
MEM/WB receives:

- ❖ ALU result
- ❖ Memory read data

3. Pipeline Datapath: **WB Stage**



At the beginning of a cycle
MEM/WB register supplies:

- ❖ ALU result
- ❖ Memory read data

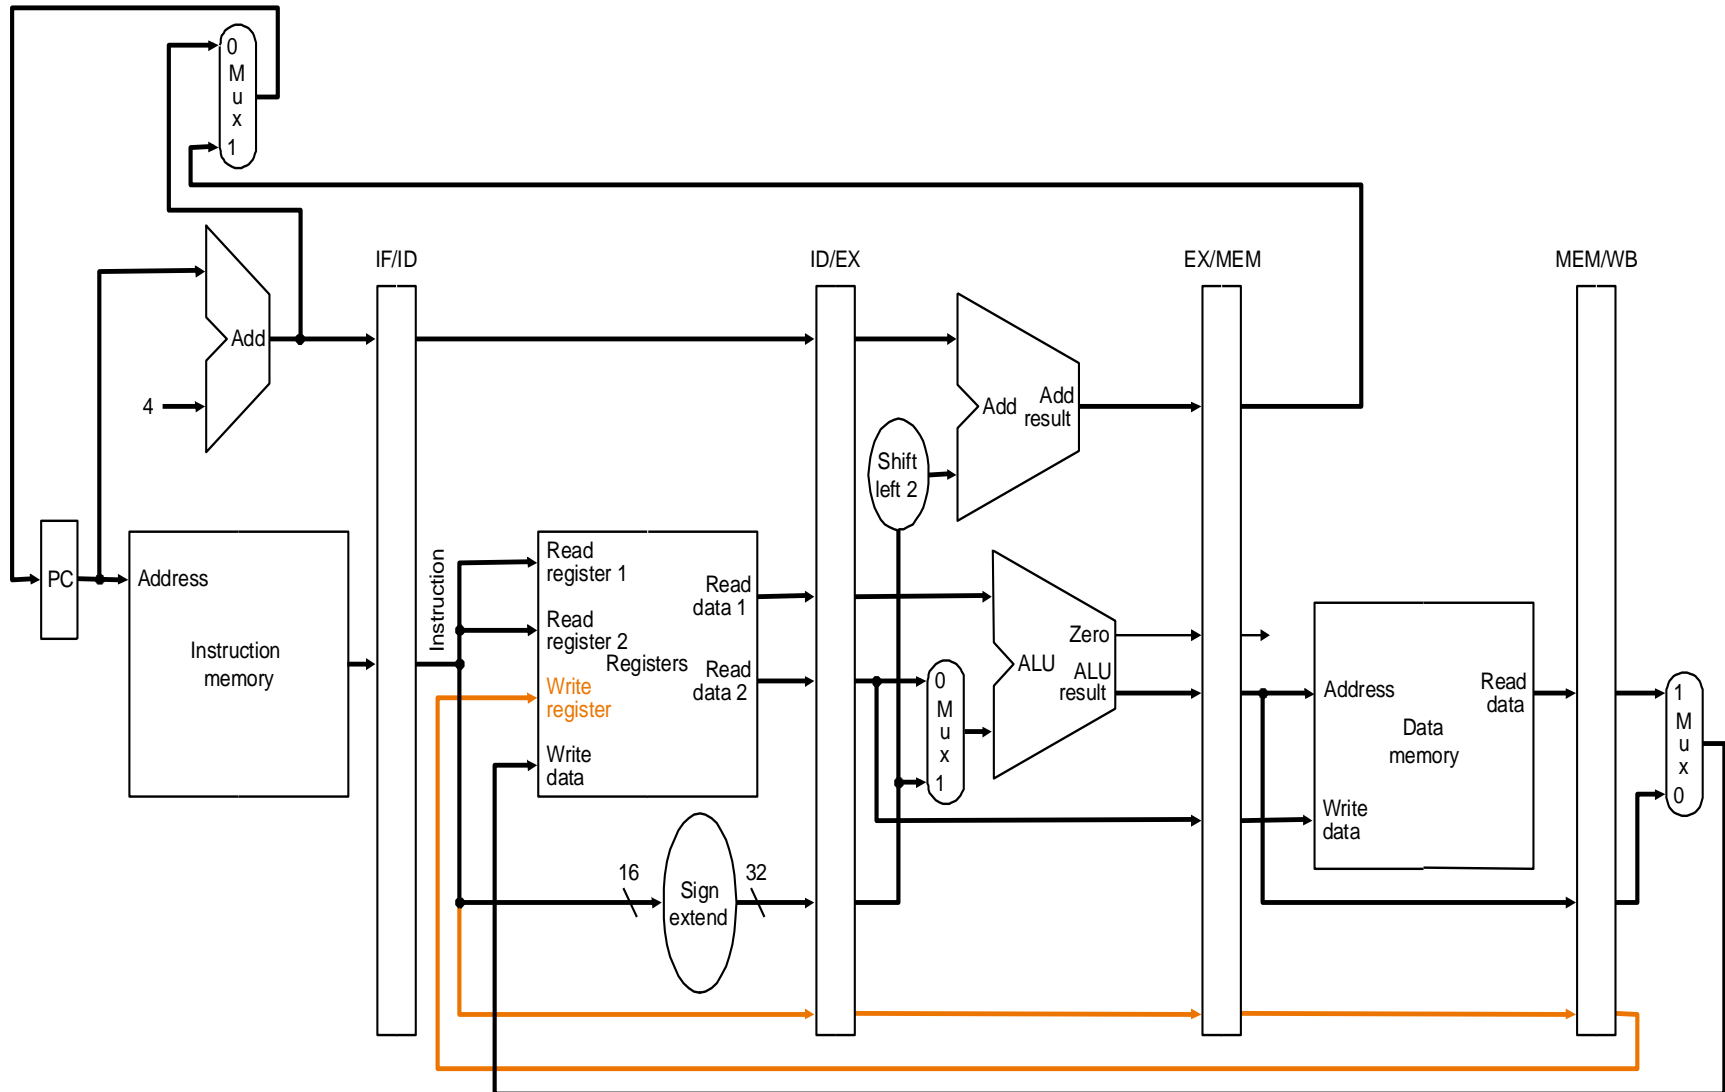
At the end of a cycle

- ❖ Result is written back to register file (if applicable)
- ❖ **There is a bug here.....**

3. Corrected Datapath (1/2)

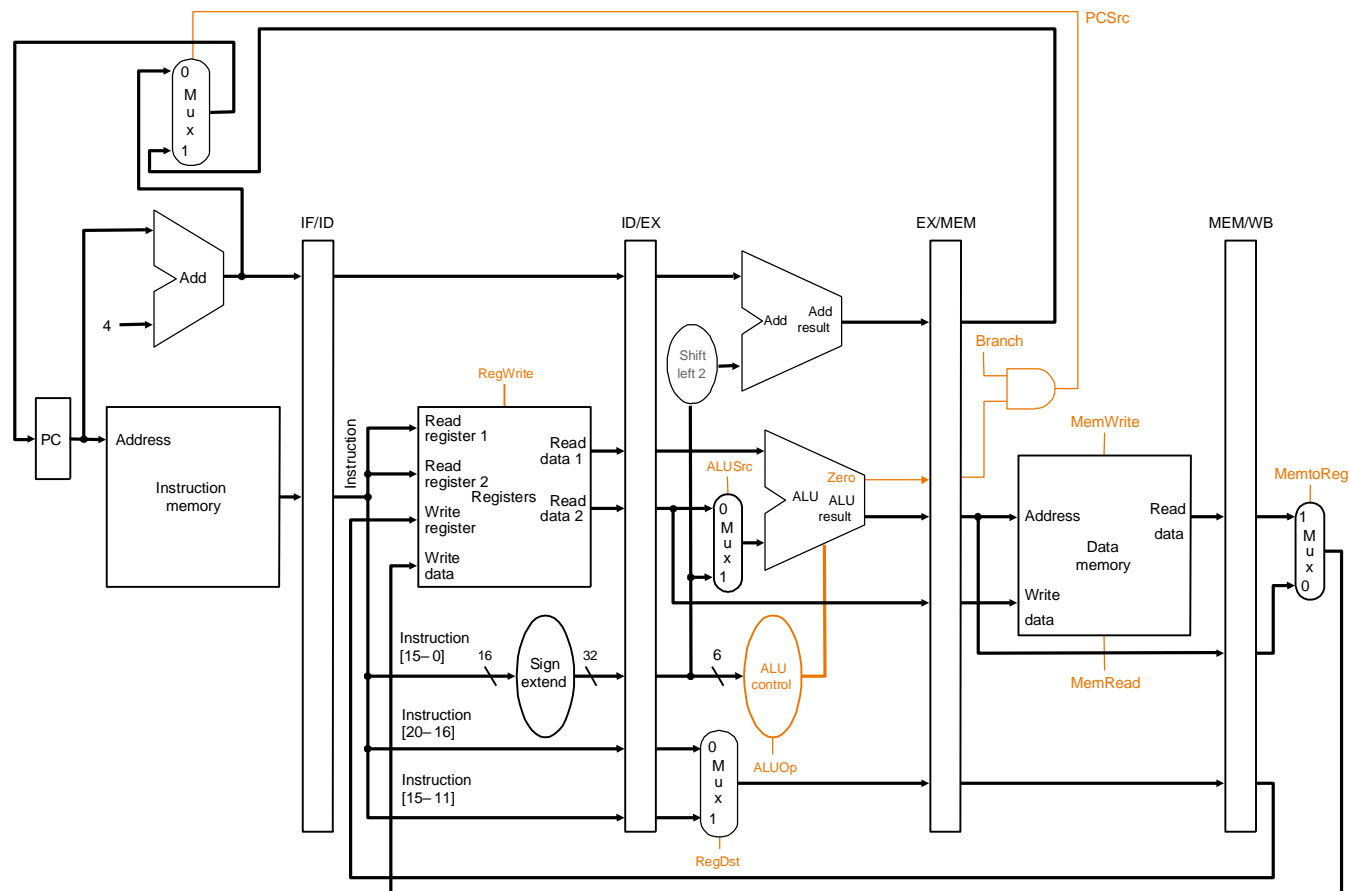
- Observe the “Write register” number
 - Supplied by the **IF/ID** pipeline register
 - It is NOT the correct write register for the instruction now in **WB** stage!
- **Solution:**
 - Pass “Write register” number from **ID/EX** through **EX/MEM** to **MEM/WB** pipeline register for use in **WB** stage
 - i.e. let the "Write register" number follows the instruction through the pipeline until it is needed in WB stage

3. Corrected Datapath (2/2)

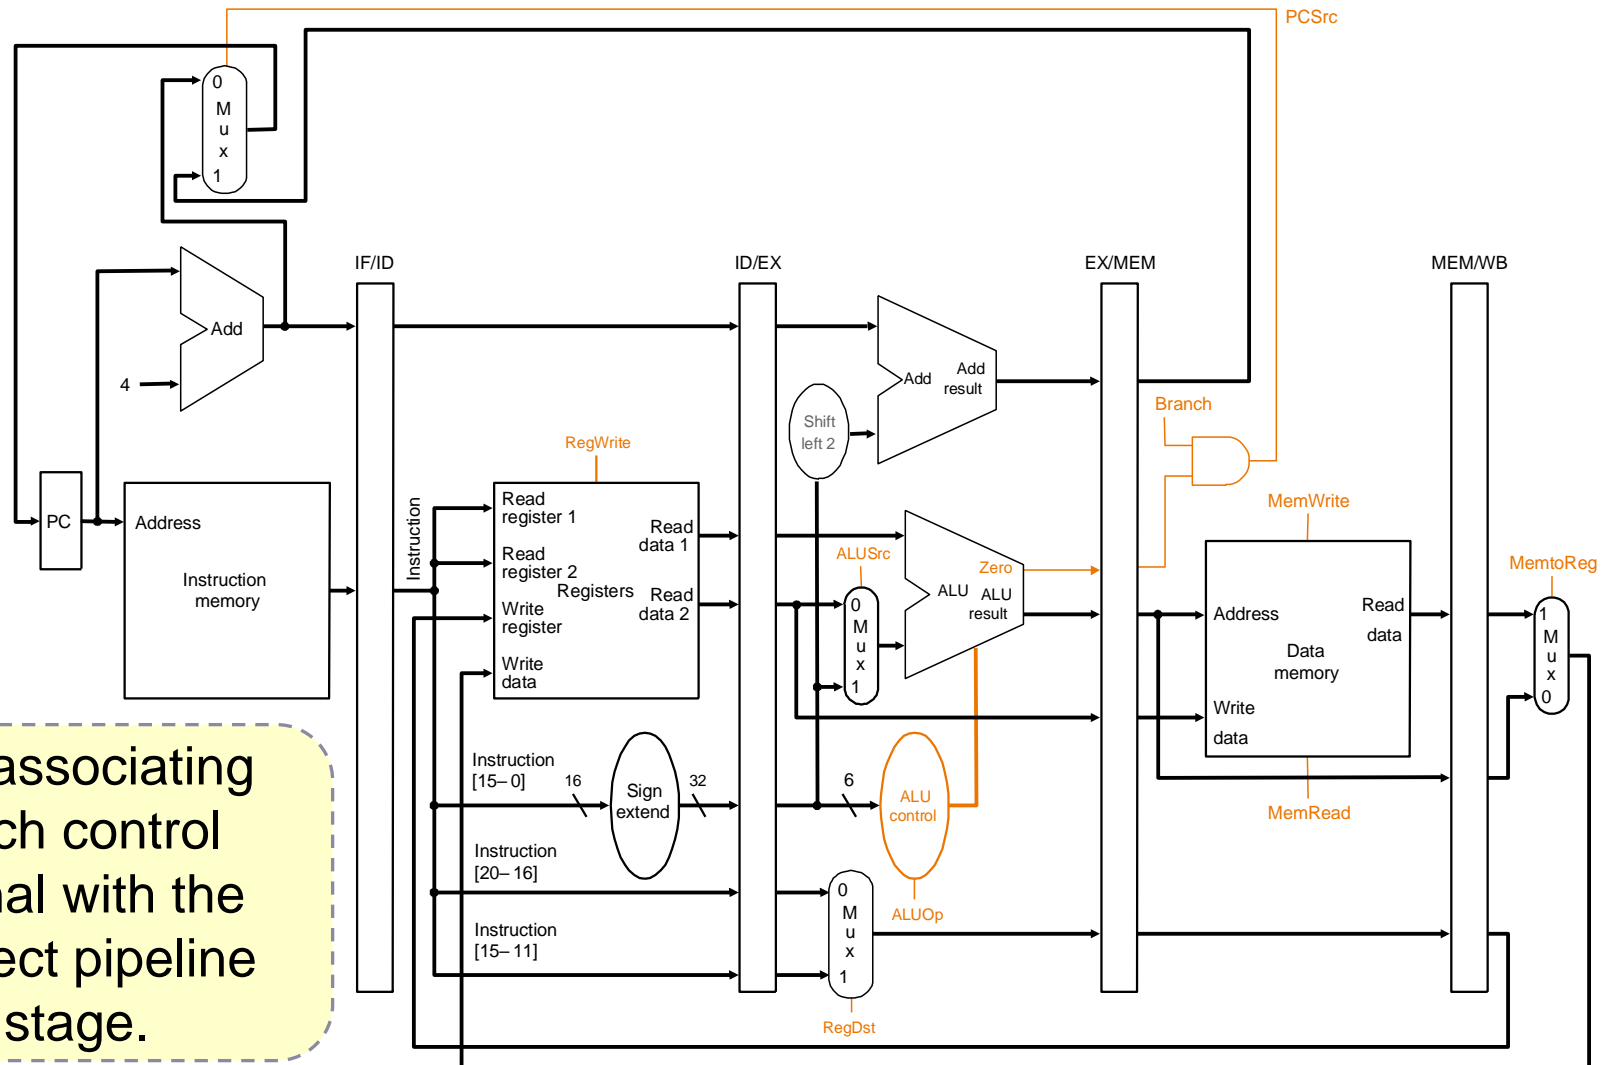


4. Pipeline Control: Main Idea

- Same control signals as single-cycle datapath
- Difference:** Each control signal belongs to a particular pipeline stage



4. Pipeline Control: Try it!



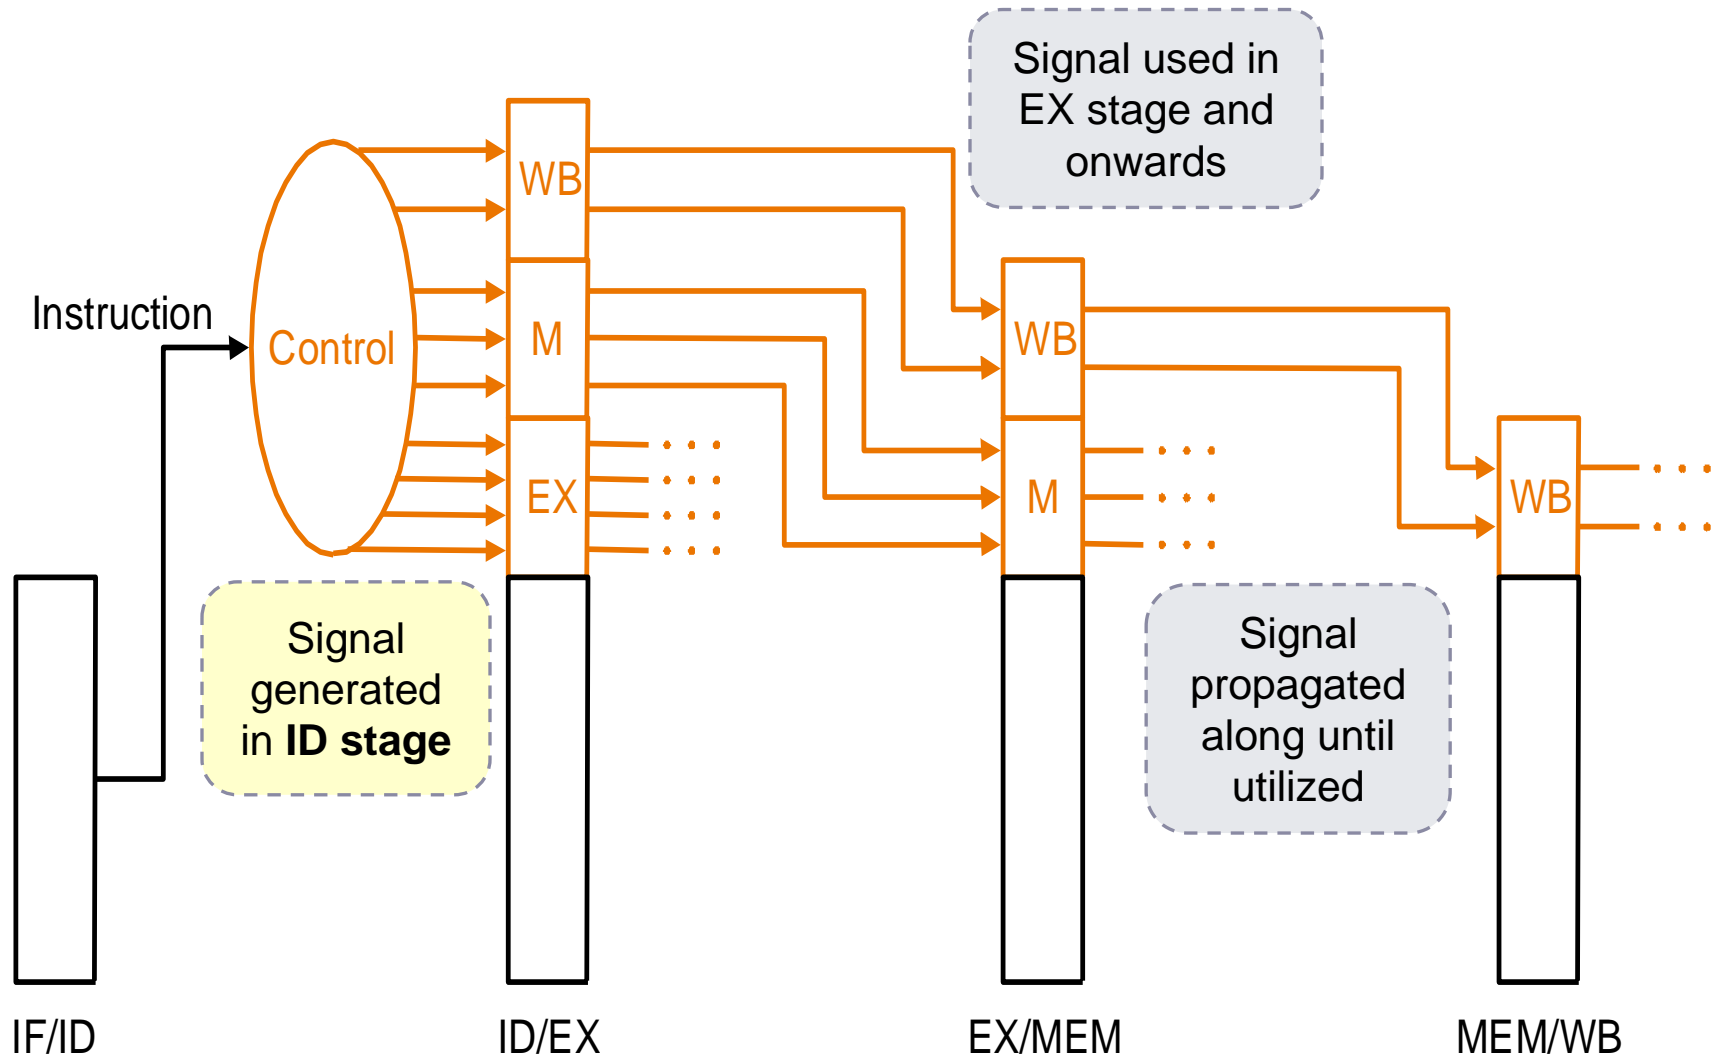
4. Pipeline Control: Grouping

- Group control signals according to pipeline stage

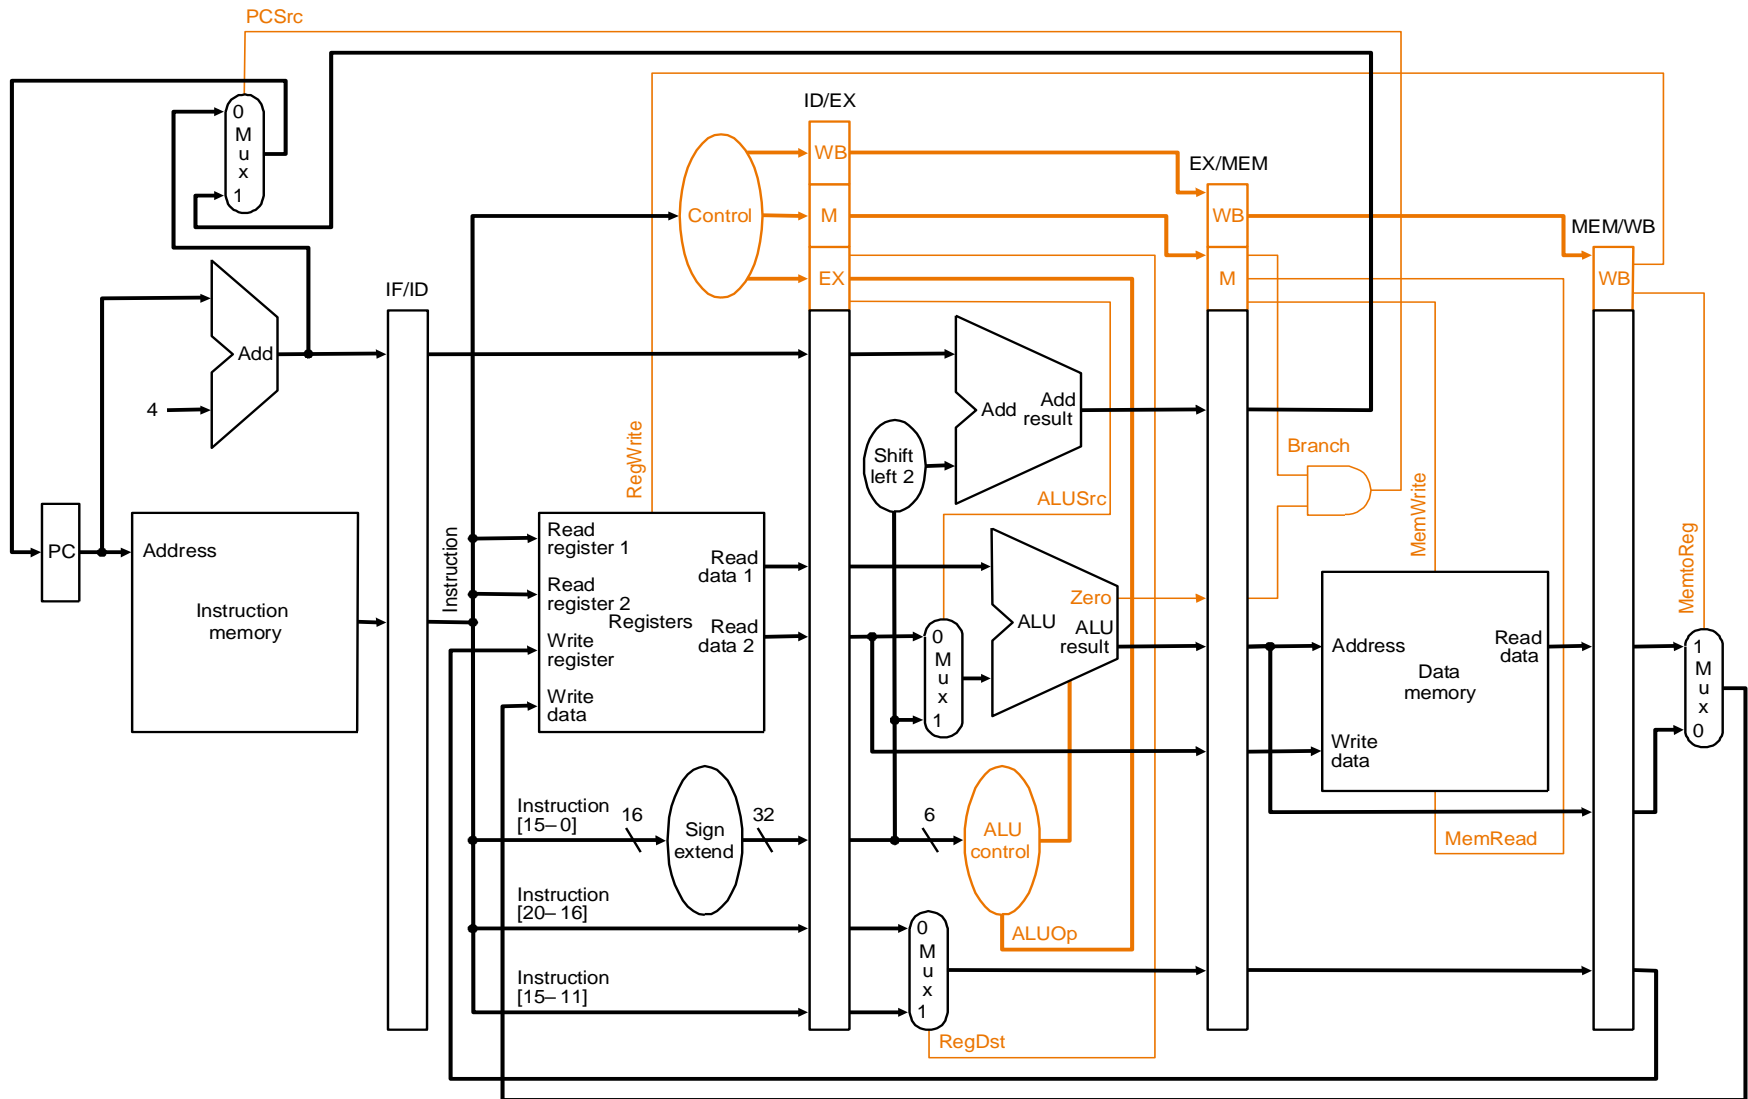
	RegDst	ALUSrc	MemTo Reg	Reg Write	Mem Read	Mem Write	Branch	ALUOp	
								op1	op0
R-type	1	0	0	1	0	0	0	1	0
lw	0	1	1	1	1	0	0	0	0
sw	X	1	X	0	0	1	0	0	0
beq	X	0	X	0	0	0	1	0	1

	EX Stage				MEM Stage			WB Stage	
	RegDst	ALUSrc	ALUOp		Mem Read	Mem Write	Branch	MemTo Reg	Reg Write
			op1	op0					
R-type	1	0	1	0	0	0	0	0	1
lw	0	1	0	0	1	0	0	1	1
sw	X	1	0	0	0	1	0	X	0
beq	X	0	0	1	0	0	1	X	0

4. Pipeline Control: Implementation




4. Pipeline Control: Datapath and Control




5. Different Implementations


Single-Cycle



Multi-Cycle



Pipeline



5. Comparison of Performance (1/2)

Single-Cycle Processor

- Cycle time: $CT_{seq} = \max(\sum_{k=1}^N T_k)$
 - T_k = Time for operation in stage k
 - N = Number of stages
- Execution Time for **I** instructions:
 - $Time_{seq} = I \times CT_{seq}$

Instruction	Inst Mem	Reg read	ALU	Data Mem	Reg write	Total
ALU (eg: add)	2	1	2		1	6
lw	2	1	2	2	1	8
sw	2	1	2	2		7
beq	2	1	2			5

Assume 100 instructions.

Cycle time = **8ns**

Time to execute 100 instructions = $100 \times 8\text{ns} = \mathbf{800\text{ns}}$

Multi-Cycle Processor

- Cycle time: $CT_{multi} = \max(T_k)$
- Execution Time for **I** instructions:
 - $Time_{multi} = I \times \text{Average CPI} \times CT_{multi}$
 - Average CPI needed as each instruction takes different number of cycles

Cycle time = **2ns**

Given average CPI = 4.6

Time to execute 100 instructions
= $100 \times 4.6 \times 2\text{ns} = \mathbf{920\text{ns}}$

5. Comparison of Performance (2/2)

Single-Cycle Processor

- Cycle time: $CT_{seq} = \max(\sum_{k=1}^N T_k)$
 - T_k = Time for operation in stage k
 - N = Number of stages
- Execution Time for **I** instructions:
 - $Time_{seq} = I \times CT_{seq}$

Instruction	Inst Mem	Reg read	ALU	Data Mem	Reg write	Total
ALU (eg: add)	2	1	2		1	6
lw	2	1	2	2	1	8
sw	2	1	2	2		7
beq	2	1	2			5

Assume 100 instructions.

Cycle time = **8ns**

Time to execute 100 instructions = $100 \times 8\text{ns} = \mathbf{800\text{ns}}$.

Pipelining Processor

- Cycle time: $CT_{pipeline} = \max(T_k) + T_d$
 - T_d = Overhead for pipelining, e.g. pipeline register
- Cycles needed for **I** instructions:
 - $I + N - 1$ (need $N - 1$ cycles to fill up the pipeline)
- Execution Time for **I** instructions:
 - $Time_{pipeline} = (I + N - 1) \times CT_{pipeline}$

Assume pipeline register latency = **0.5ns**

Cycle time = $2\text{ns} + 0.5\text{ns} = 2.5\text{ns}$

Time to execute 100 instructions
= $(100 + 5 - 1) \times 2.5\text{ns} = \mathbf{260\text{ns}}$.

5. Pipeline Processor: Ideal Speedup

- Assumptions for ideal case:
 - Every stage takes the same amount of time $\rightarrow \sum_{k=1}^N T_k = N \times T_1$
 - No pipeline overhead $\rightarrow T_d = 0$
 - $I \gg N$ (Number of instructions is much larger than number of stages)

$$\begin{aligned} \text{Speedup}_{\text{pipeline}} &= \frac{\text{Time}_{\text{seq}}}{\text{Time}_{\text{pipeline}}} \\ &= \frac{I \times \sum_{k=1}^N T_k}{(I+N-1) \times (\max(T_k) + T_d)} = \frac{I \times N \times T_1}{(I+N-1) \times T_1} \\ &\approx \frac{I \times N \times T_1}{I \times T_1} \\ &= N \end{aligned}$$

Conclusion:

Pipeline processor can gain **N** times speedup, where **N** is the number of pipeline stages.

Review Question

- Given this code:

```
add $t0, $s0, $s1  
sub $t1, $s0, $s1  
sll $t2, $s0, 2  
srl $t3, $s1, 2
```

a) 4 cycles

b) $4/(100 \times 10^6) = 40 \text{ ns}$

c) $4 + 4 = 8 \text{ cycles}$

d) $8/(500 \times 10^6) = 16 \text{ ns}$

- a) How many cycles will it take to execute the code on a single-cycle datapath?
- b) How long will it take to execute the code on a single-cycle datapath, assuming a 100 MHz clock?
- c) How many cycles will it take to execute the code on a 5-stage MIPS pipeline?
- d) How long will it take to execute the code on a 5-stage MIPS pipeline, assuming a 500 MHz clock?

End of File