Lecture #21

Pipelining
Part II: Hazards

NUS National University of Singapore | School of Computing

Lecture #21: Pipelining II

1. Pipeline Hazards

2. Structural Hazards

3. Instruction Dependencies

4. Data Hazards

   4.1 Forwarding

   4.2 Stall

   4.3 Exercises

Lecture #21: Pipelining II

5. Control Dependency

6. Control Hazards

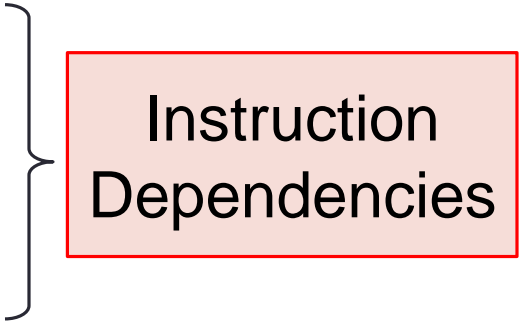   6.1 Early Branch

   6.2 Branch Prediction

   6.3 Delayed Branched

7. Multiple Issue Processors (reading)

1. Pipeline Hazards

- Speedup from pipeline implementation:
  - Based on the assumption that a new instructions can be "pumped" into pipeline every cycle

- However, there are **pipeline hazards**
  - Problems that prevent next instruction from immediately following previous instruction
  - **Structural hazards:**
    - Simultaneous use of a hardware resource
  - **Data hazards:**
    - Data dependencies between instructions
  - **Control hazards:**
    - Change in program flow
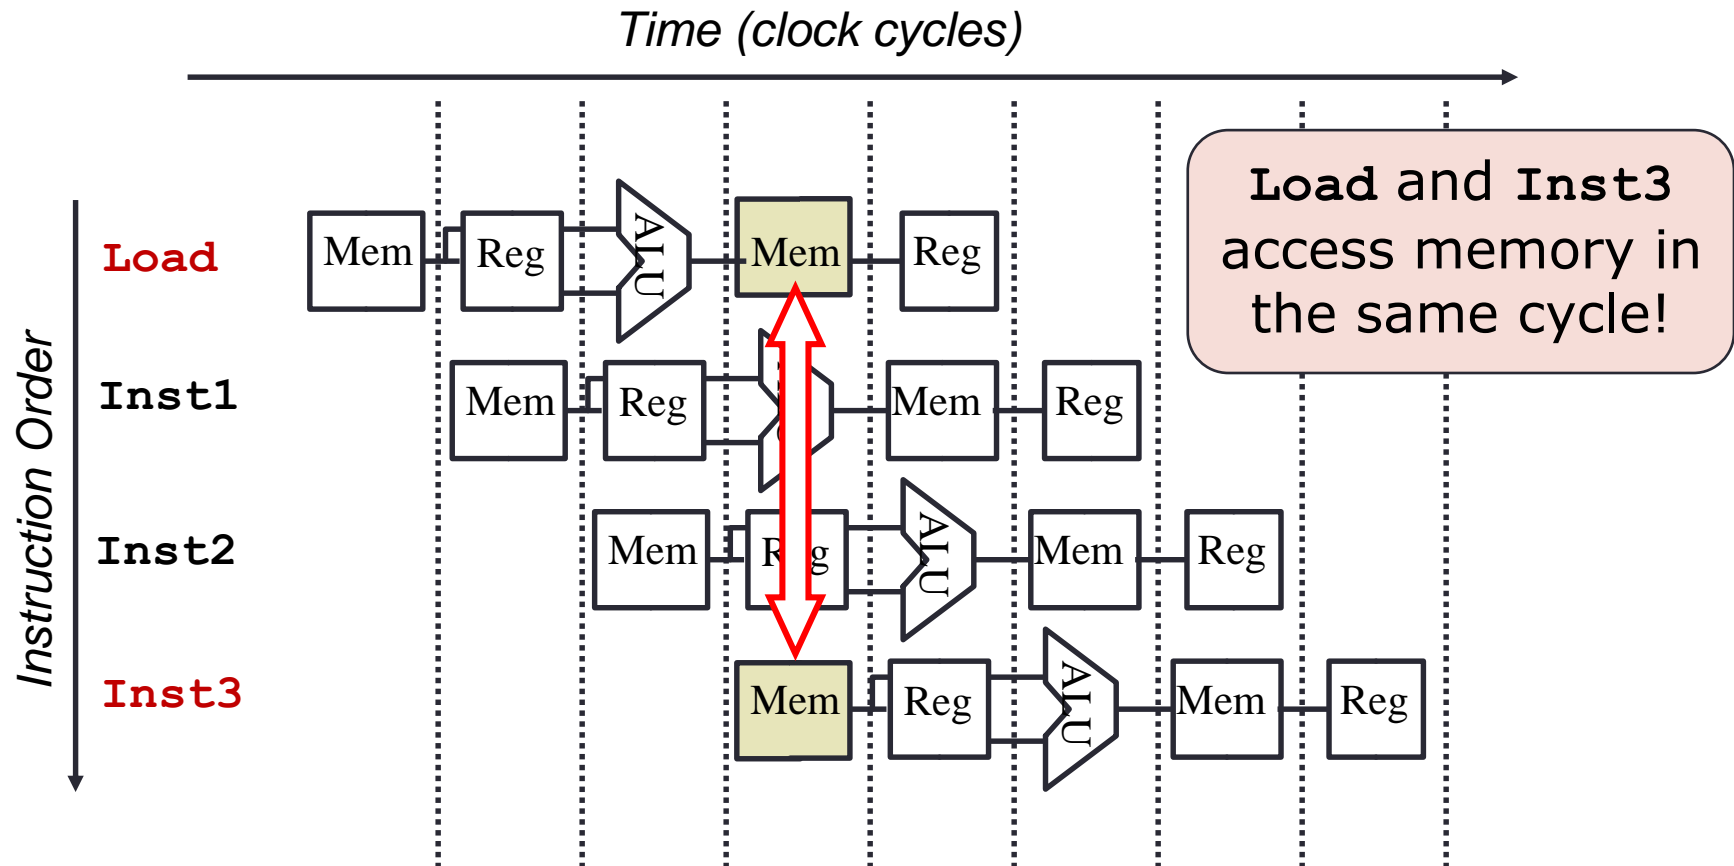
Instruction Dependencies
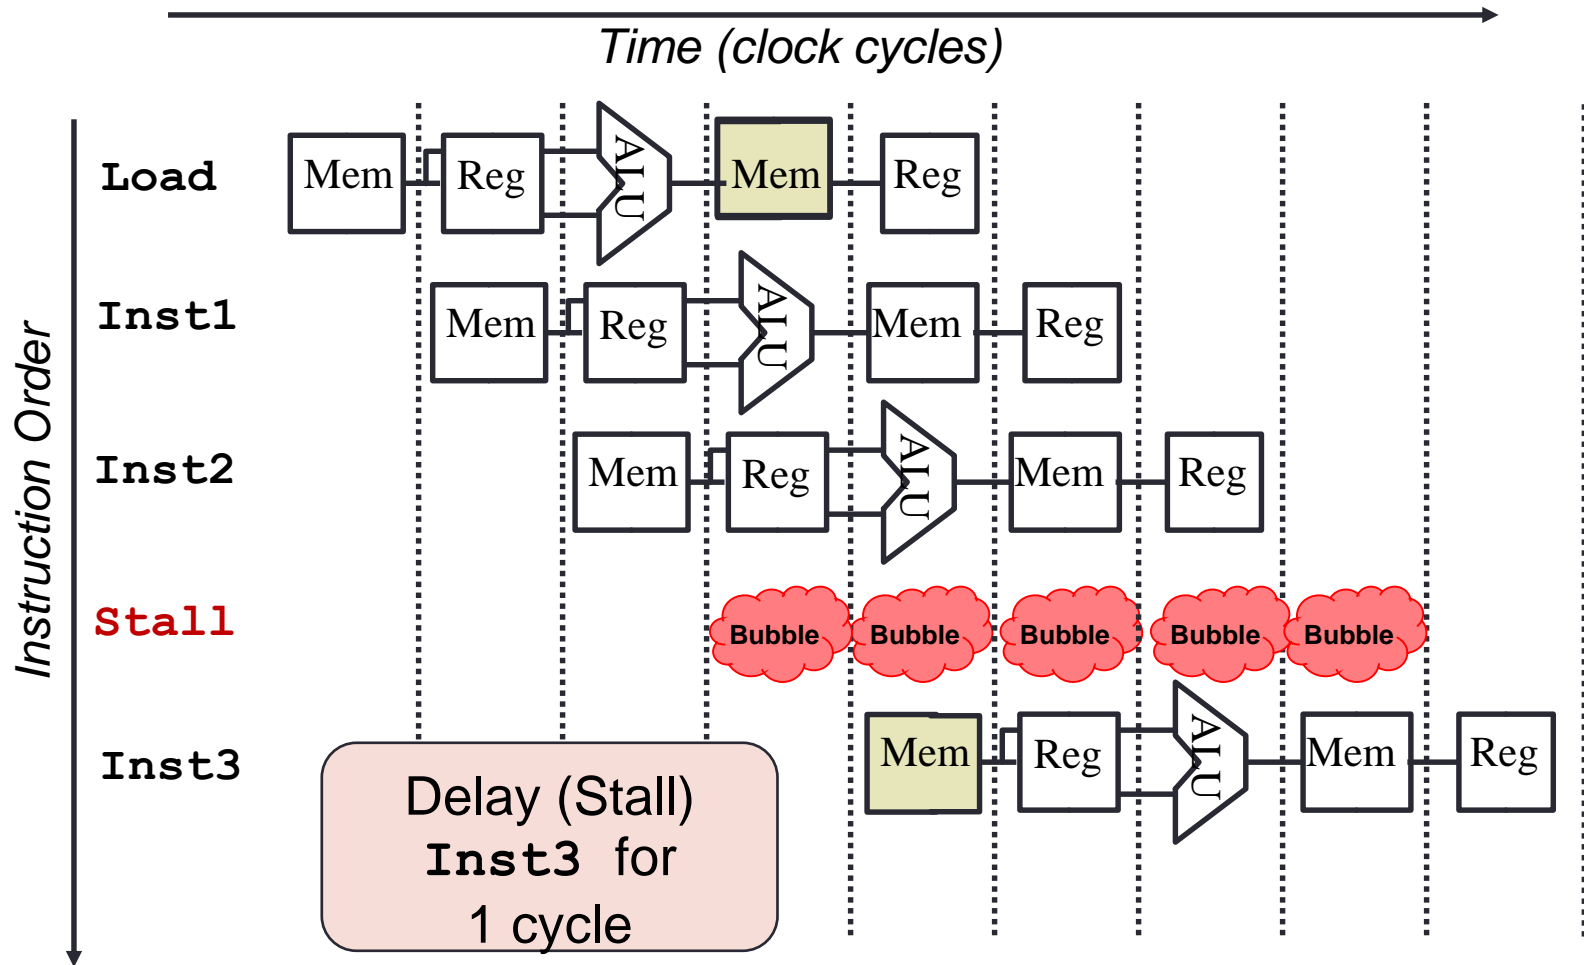
1. Graphical Notation for Pipeline
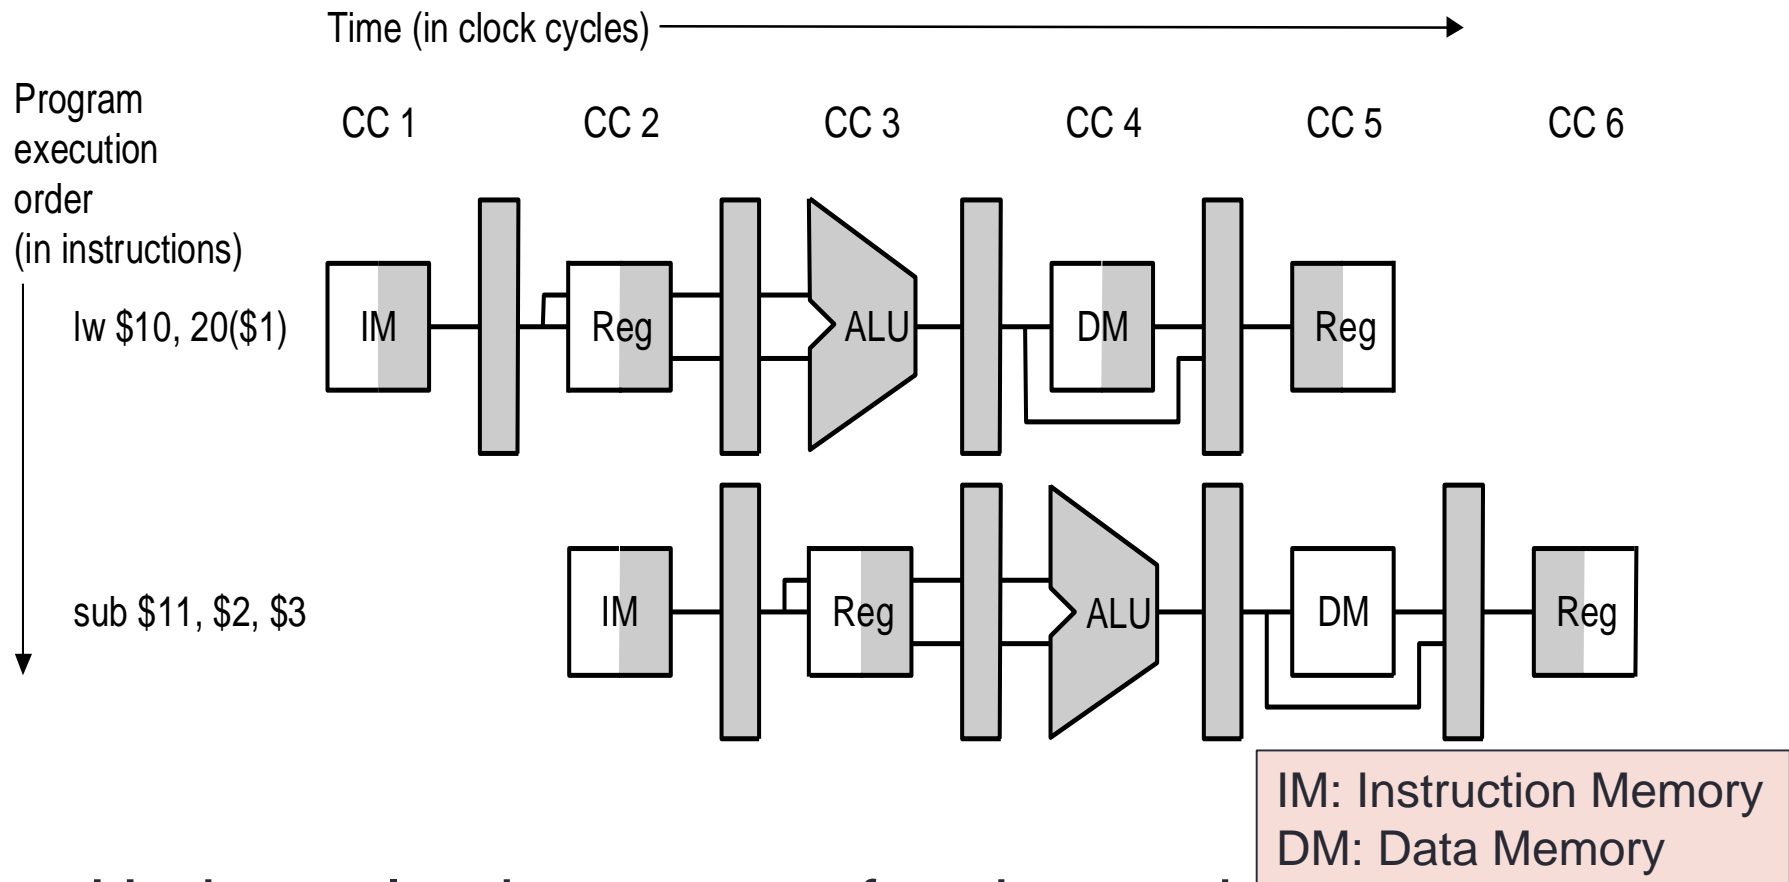
- Horizontal = the stages of an instruction
- Vertical = the instructions in different pipeline stages

2. Structural Hazard: Example

- If there is only a **single memory module:**

Time (clock cycles)

Load and Inst3 access memory in the same cycle!

2. Solution 1: Stall the Pipeline

Time (clock cycles)

Instruction Order

Load   Mem — Reg — ALU — Mem — Reg

Inst1   Mem — Reg — ALU — Mem — Reg

Inst2   Mem — Reg — ALU — Mem — Reg

Stall   Bubble Bubble Bubble Bubble Bubble

Inst3   Mem — Reg — ALU — Mem — Reg

Delay (Stall) **Inst3** for 1 cycle
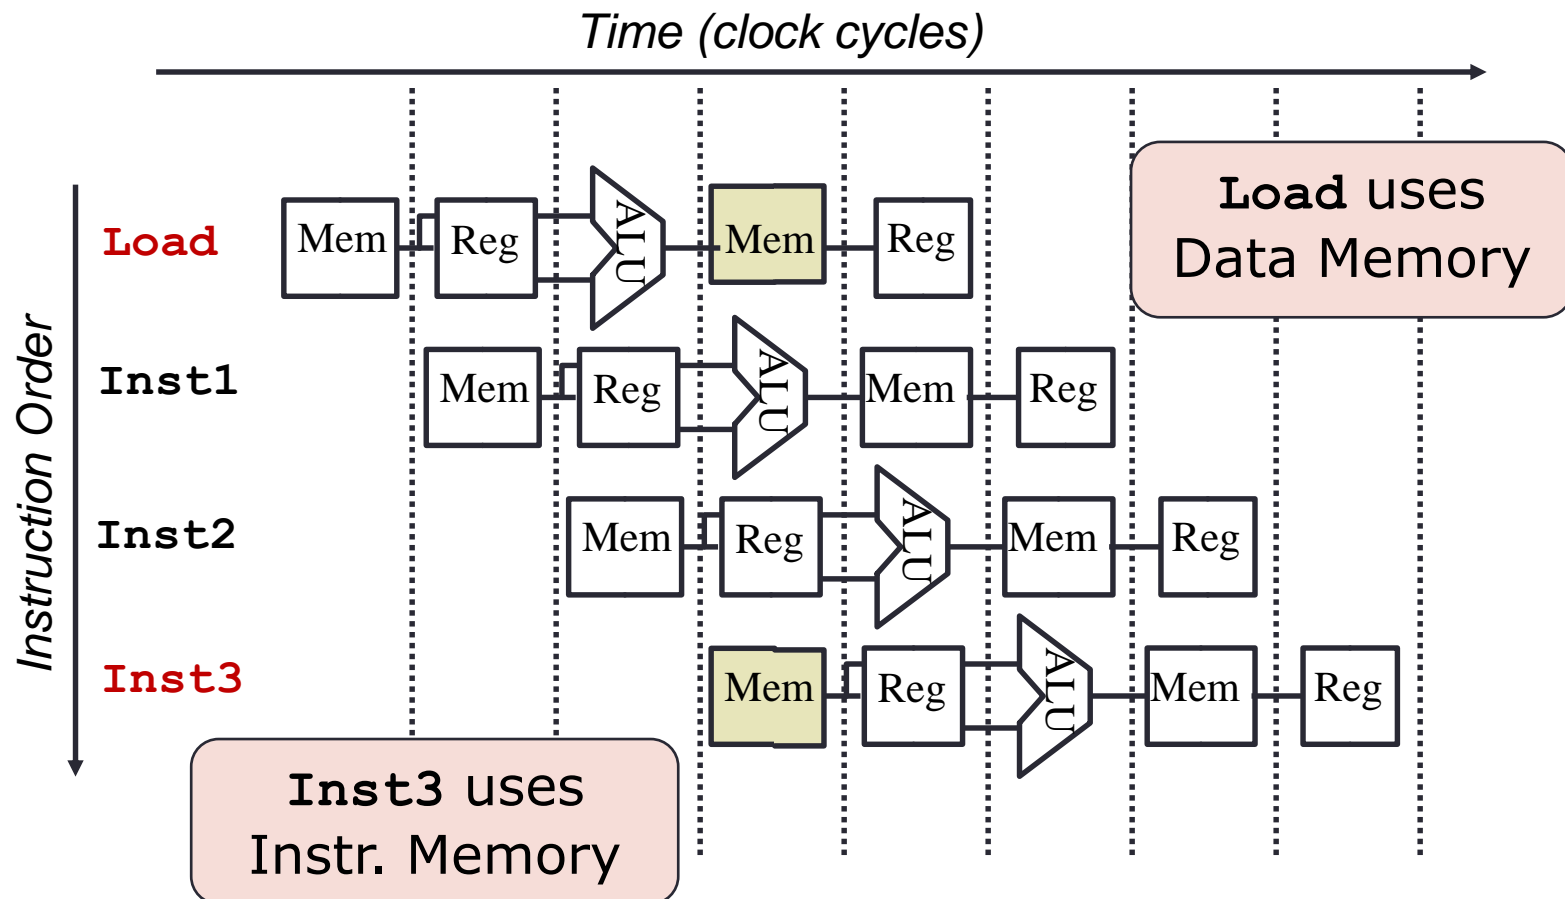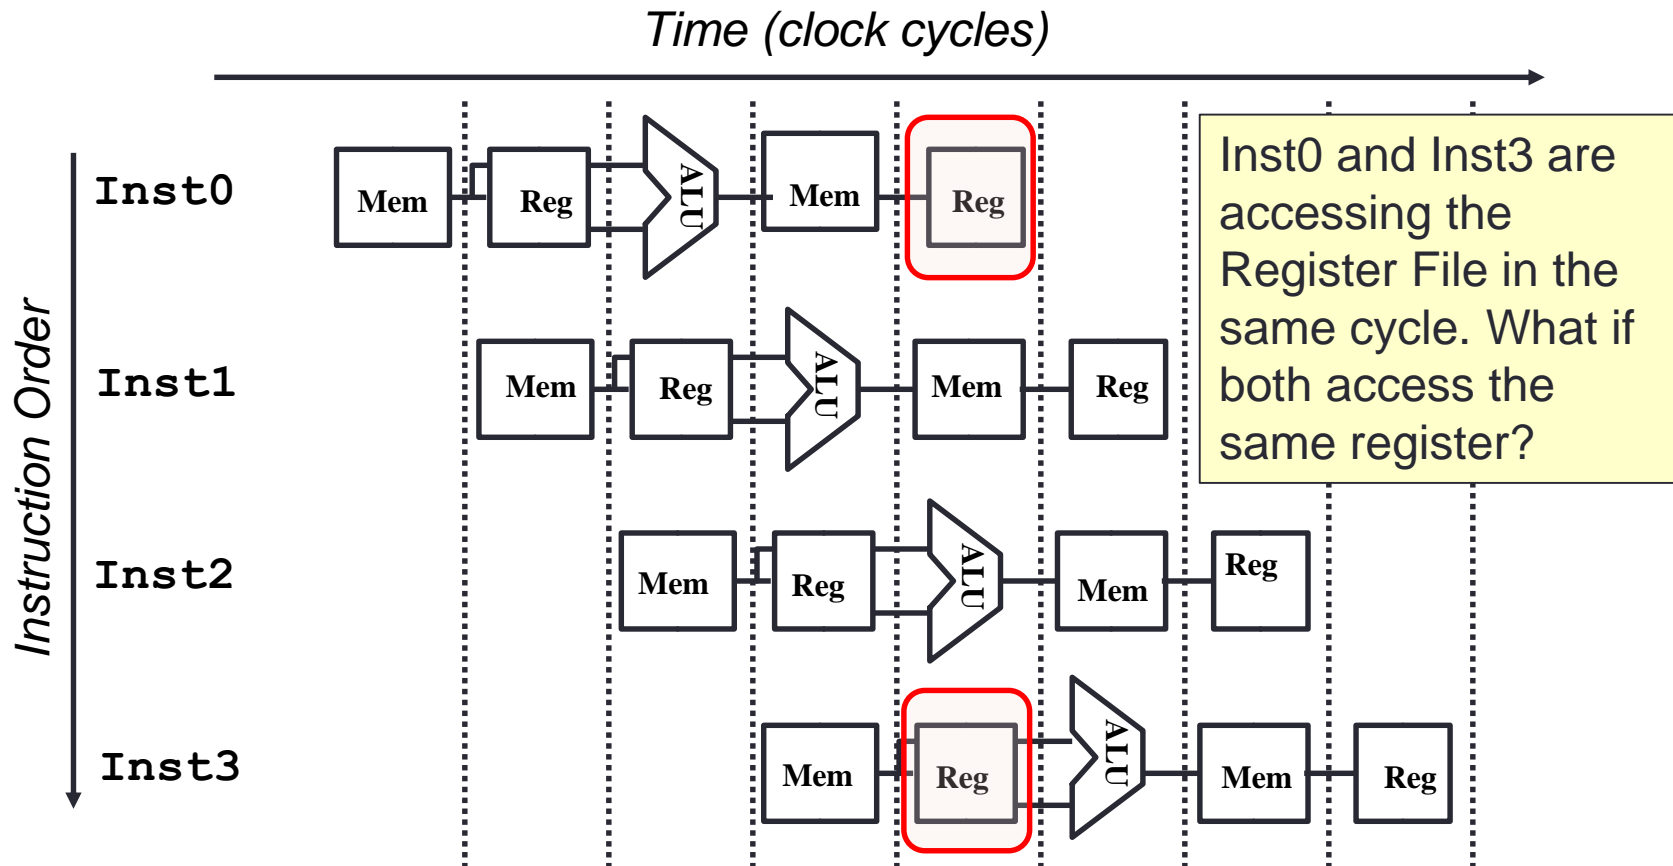
2. Solution 2: Separate Memory

- Split memory into **Data** and **Instruction** memory

2. Quiz (1/2)

Is there another conflict?

Time (clock cycles)

Instruction Order

Inst0

Inst1

Inst2

Inst3

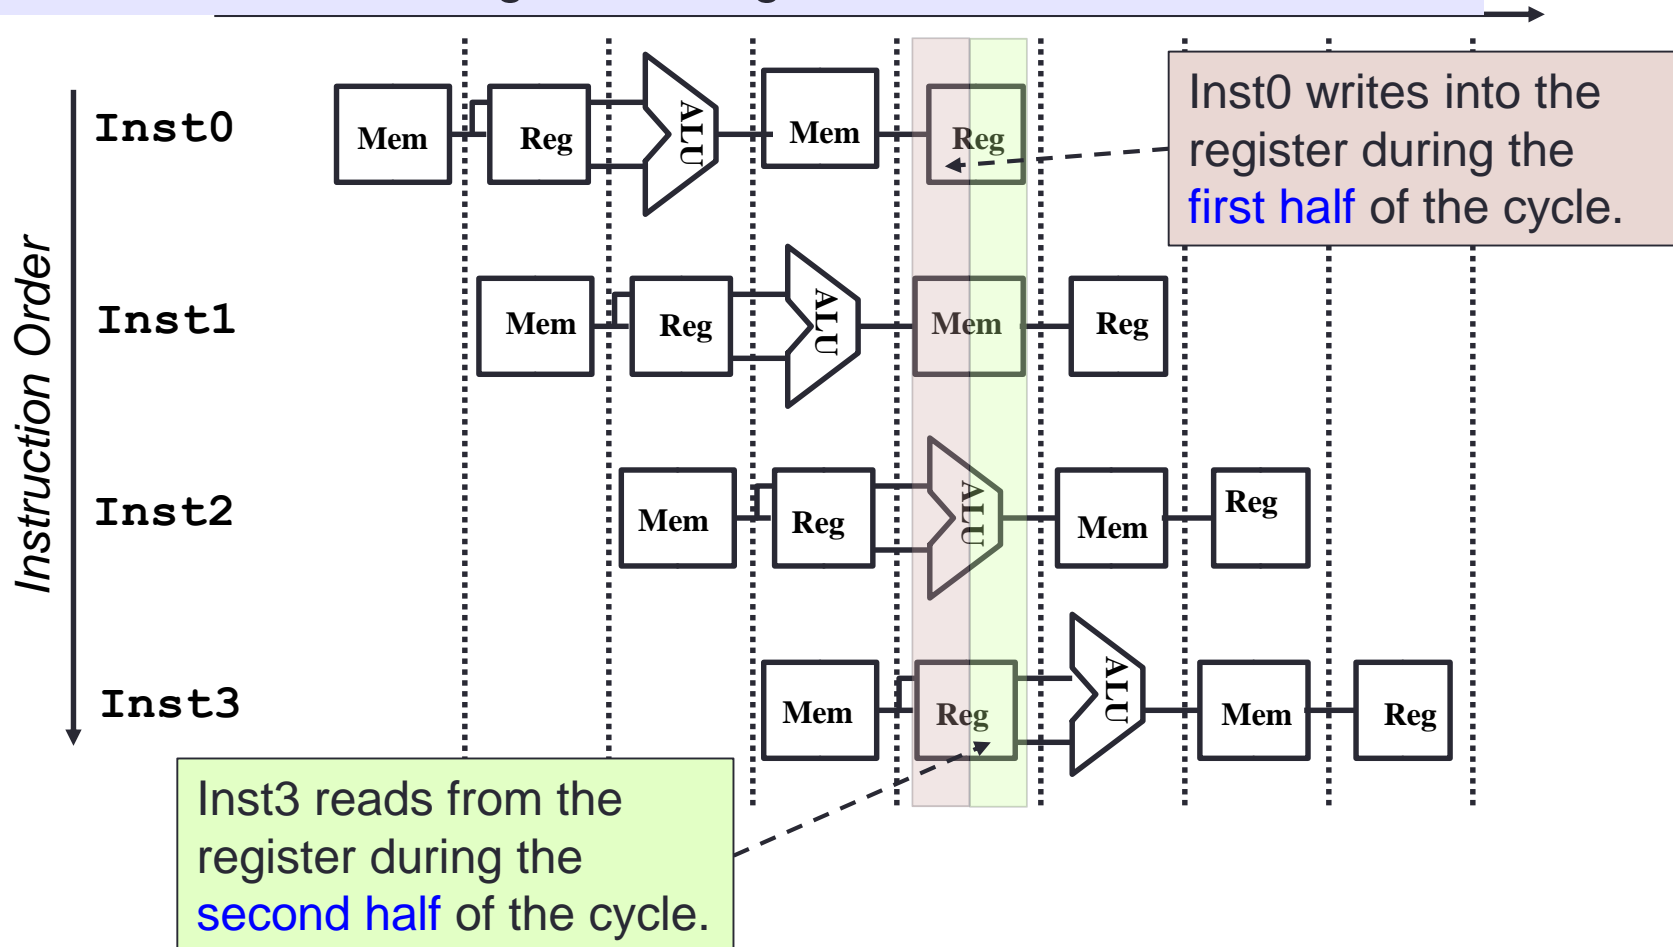Inst0 and Inst3 are accessing the Register File in the same cycle. What if both access the same register?

2. Quiz (2/2)

Recall that registers are very fast memory.
Solution: Split cycle into half; first half for writing into a register;
second half for reading from a register.

Inst0 writes into the register during the first half of the cycle.

Inst3 reads from the register during the second half of the cycle.

3. Instruction Dependencies

- Instructions can have relationship that prevent pipeline execution:
  - Although a partial overlap maybe possible in some cases

- When different instructions accesses (read/write) the same register
  - Register contention is the cause of dependency
  - Known as **data dependency**

- When the execution of an instruction depends on another instruction
  - Control flow is the cause of dependency
  - Known as **control dependency**

- Failure to handle dependencies can affect **program correctness**!

3. Data Dependency: RAW

- "**R**ead-**A**fter-**W**rite" **Definition:**
  - Occurs when a later instruction **reads** from the destination register **written** by an earlier instruction
  - Also known as **true data dependency**

```
i1: add $1, $2, $3 #writes to $1
i2: sub $4, $1, $5 #reads from $1
```

- Effect of incorrect execution:
  - If **i2** reads register **$1** before **i1** can write back the result, **i2** will get a *stale result (old result)*

3. Other Data Dependencies

- Similarly, we have:
  - **WAR**: **W**rite-**a**fter-**R**ead dependency
  - **WAW**: **W**rite-**a**fter-**W**rite dependency

- Fortunately, these dependencies **do not cause any pipeline hazards**

- They affect the processor only when instructions are executed out of program order:
  - i.e. in Modern SuperScalar Processor

4. RAW Dependency: Hazards?
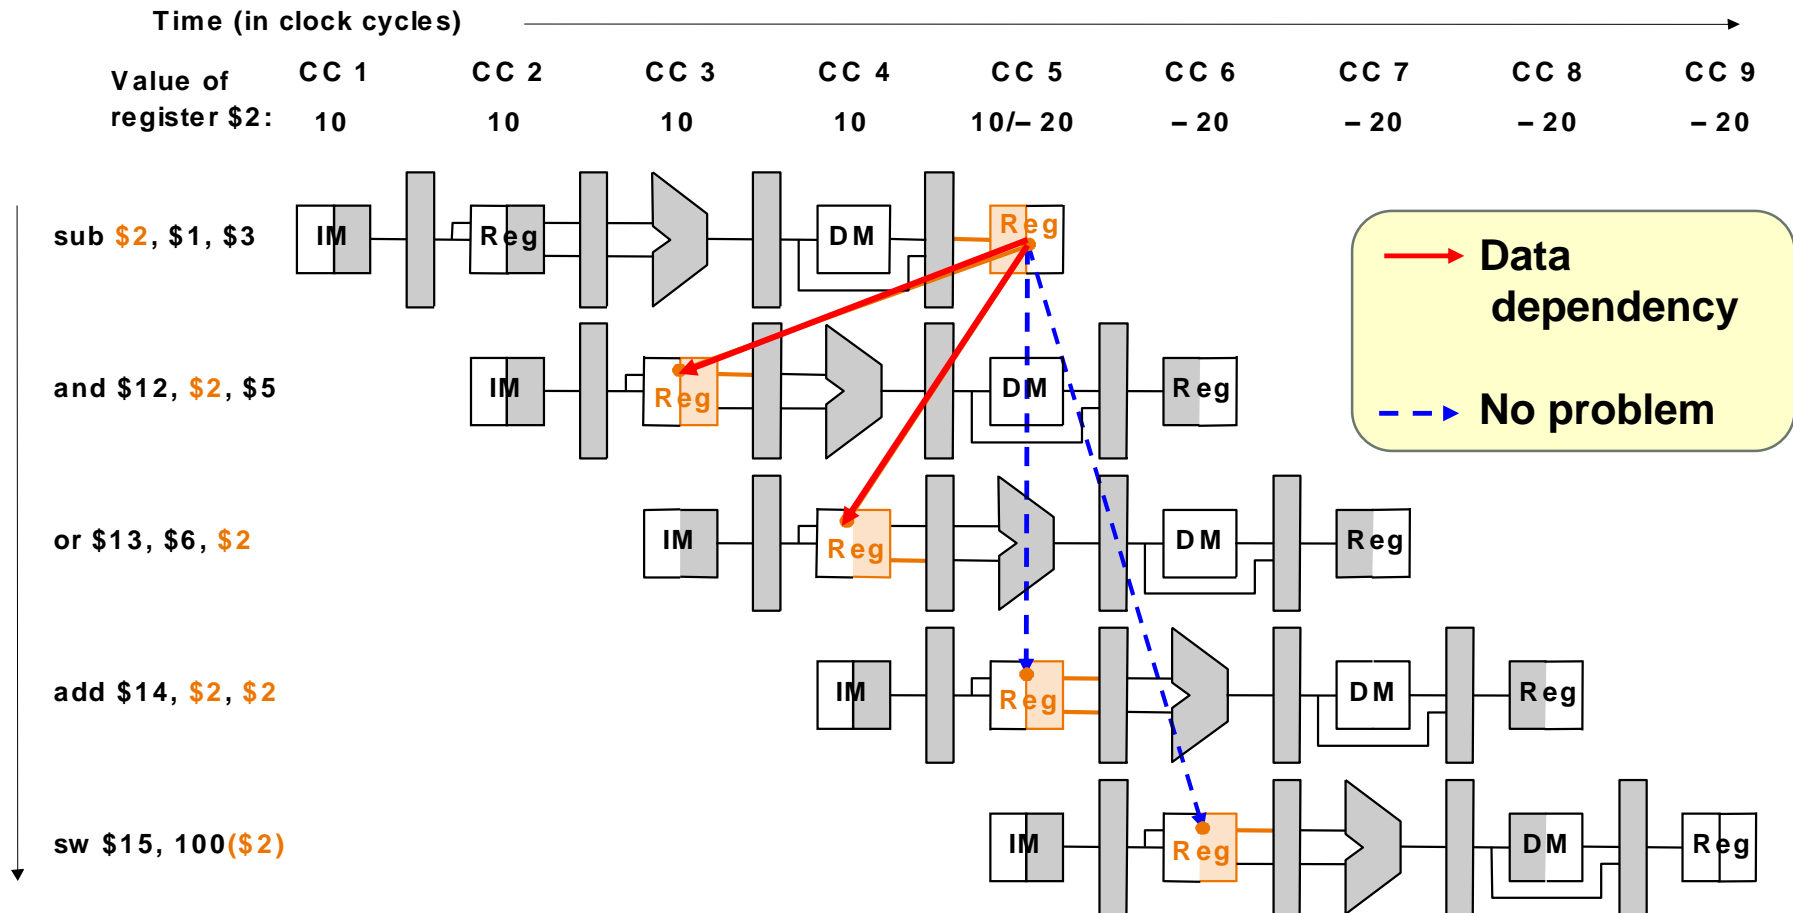
- Suppose we are executing the following code fragment:

```
sub $2, $1, $3      #i1
and $12, $2, $5     #i2
or  $13, $6, $2     #i3
add $14, $2, $2     #i4
sw  $15, 100($2)    #i5
```

- Note the multiple uses of register **$2**

- Question:
  - Which are the instructions require special handling?

4. RAW Data Hazards

- Value from prior instruction is needed before write back

4. RAW Data Hazards: Observations

- Questions:
  - When is the result from **sub** instruction actually produced?
    - End of EX stage for **sub** or clock cycle 3
  - When is the data actually needed by **and**?
    - Beginning of **and**'s EX stage or clock cycle 4
  - When is the data actually needed by **or**?
    - Beginning of **or**'s EX stage or clock cycle 5

- **Solution:**
  - **Forward** the result to any trailing (later) instructions before it is reflected in register file
  - ➔ **Bypass (replace)** the data read from register file

4.1 RAW Data Hazards: **Forwarding**

Time (in clock cycles)

	CC 1	CC 2	CC 3	CC 4	CC 5	CC 6	CC 7	CC 8	CC 9
Value of register $2 :	10	10	10	10	10/– 20	– 20	– 20	– 20	– 20
Value of EX/MEM :	X	X	X	– 20	X	X	X	X	X
Value of MEM/WB :	X	X	X	X	– 20	X	X	X	X

sub **$2**, $1, $3

and $12, **$2**, $5

or $13, $6, **$2**

add $14, **$2**, **$2**

sw $15, 100**($2)**

- **Forward** results from one stage to another
- **Bypass** data read from register file

4.2 Data Hazards: **LOAD** Instruction

Cannot solve with forwarding!

Data is needed **before** it is actually produced!
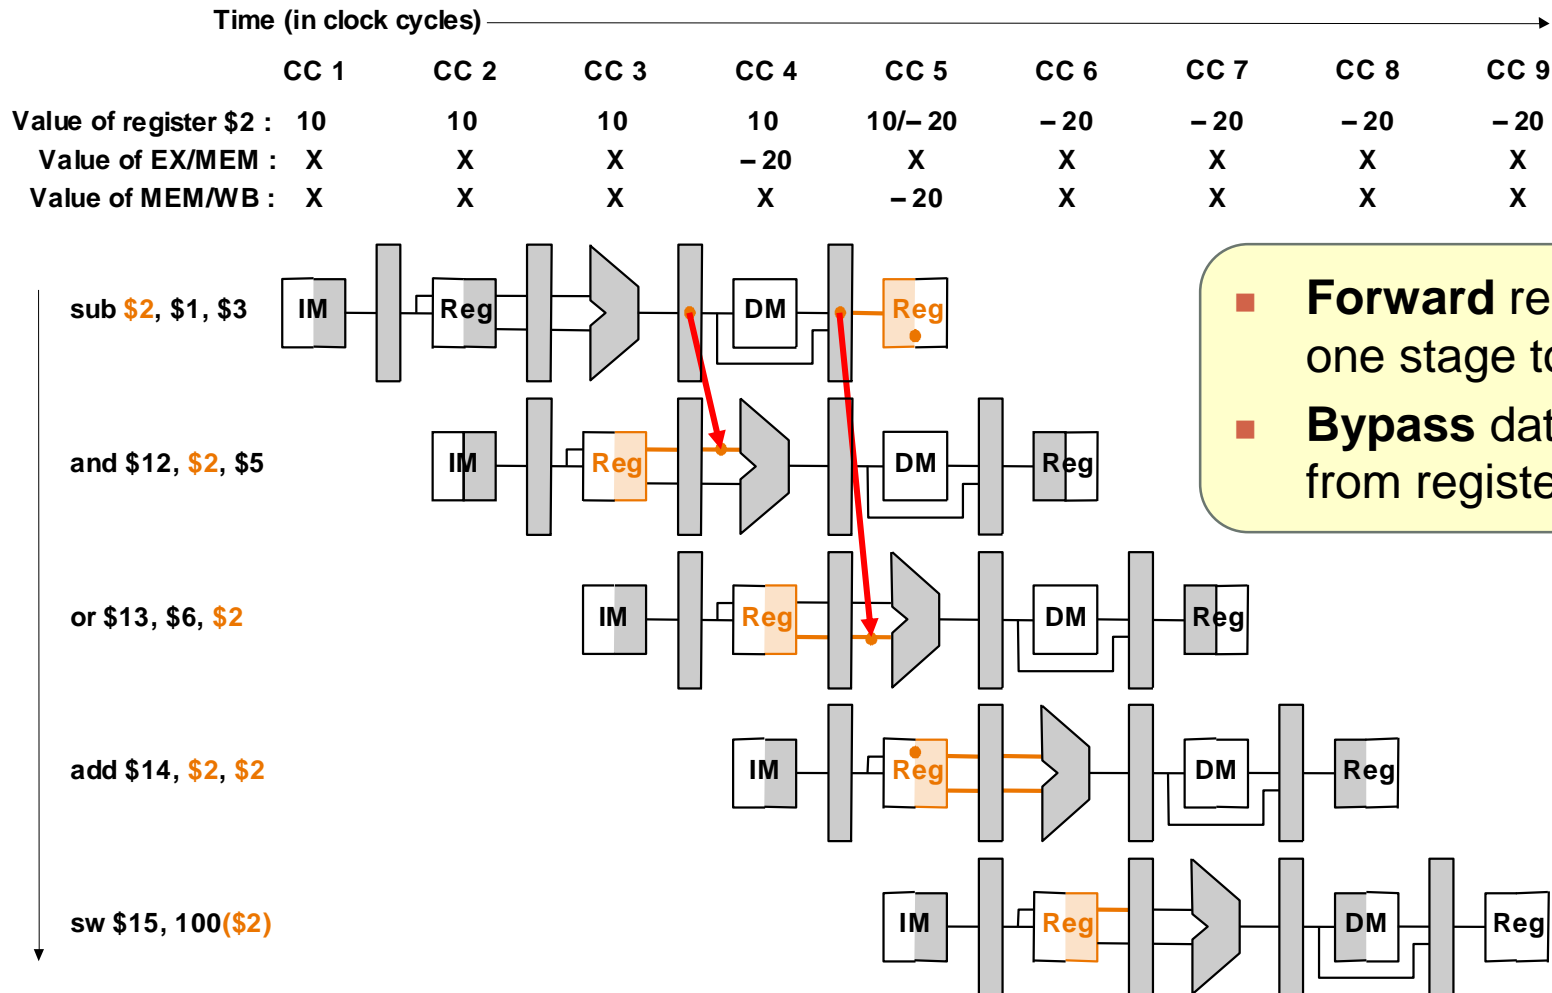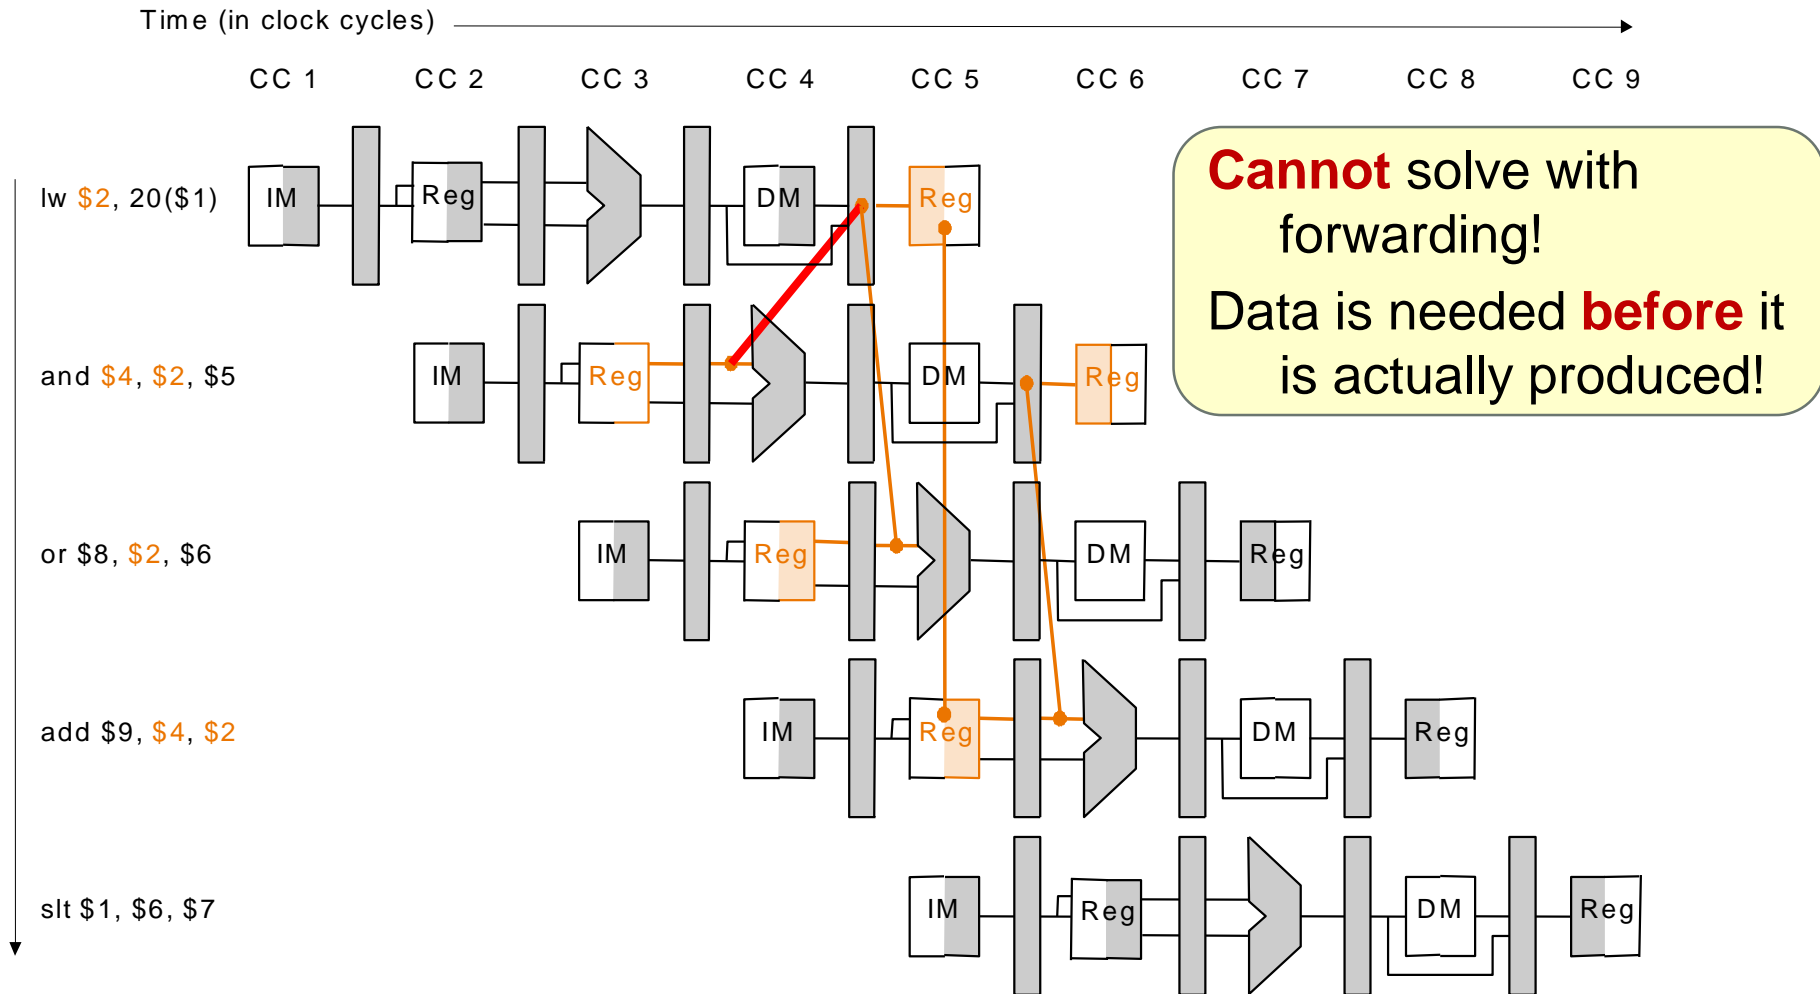
4.2 Data Hazards: **LOAD** Instruction Solution

Program execution order (in instructions)

Time (in clock cycles)

CC 1   CC 2   CC 3   CC 4   CC 5   CC 6   CC 7   CC 8   CC 9   CC 10

lw $2, 20($1)

and $4, $2, $5

or $8, $2, $6

add $9, $4, $2

slt $1, $6, $7

Stall the pipeline!
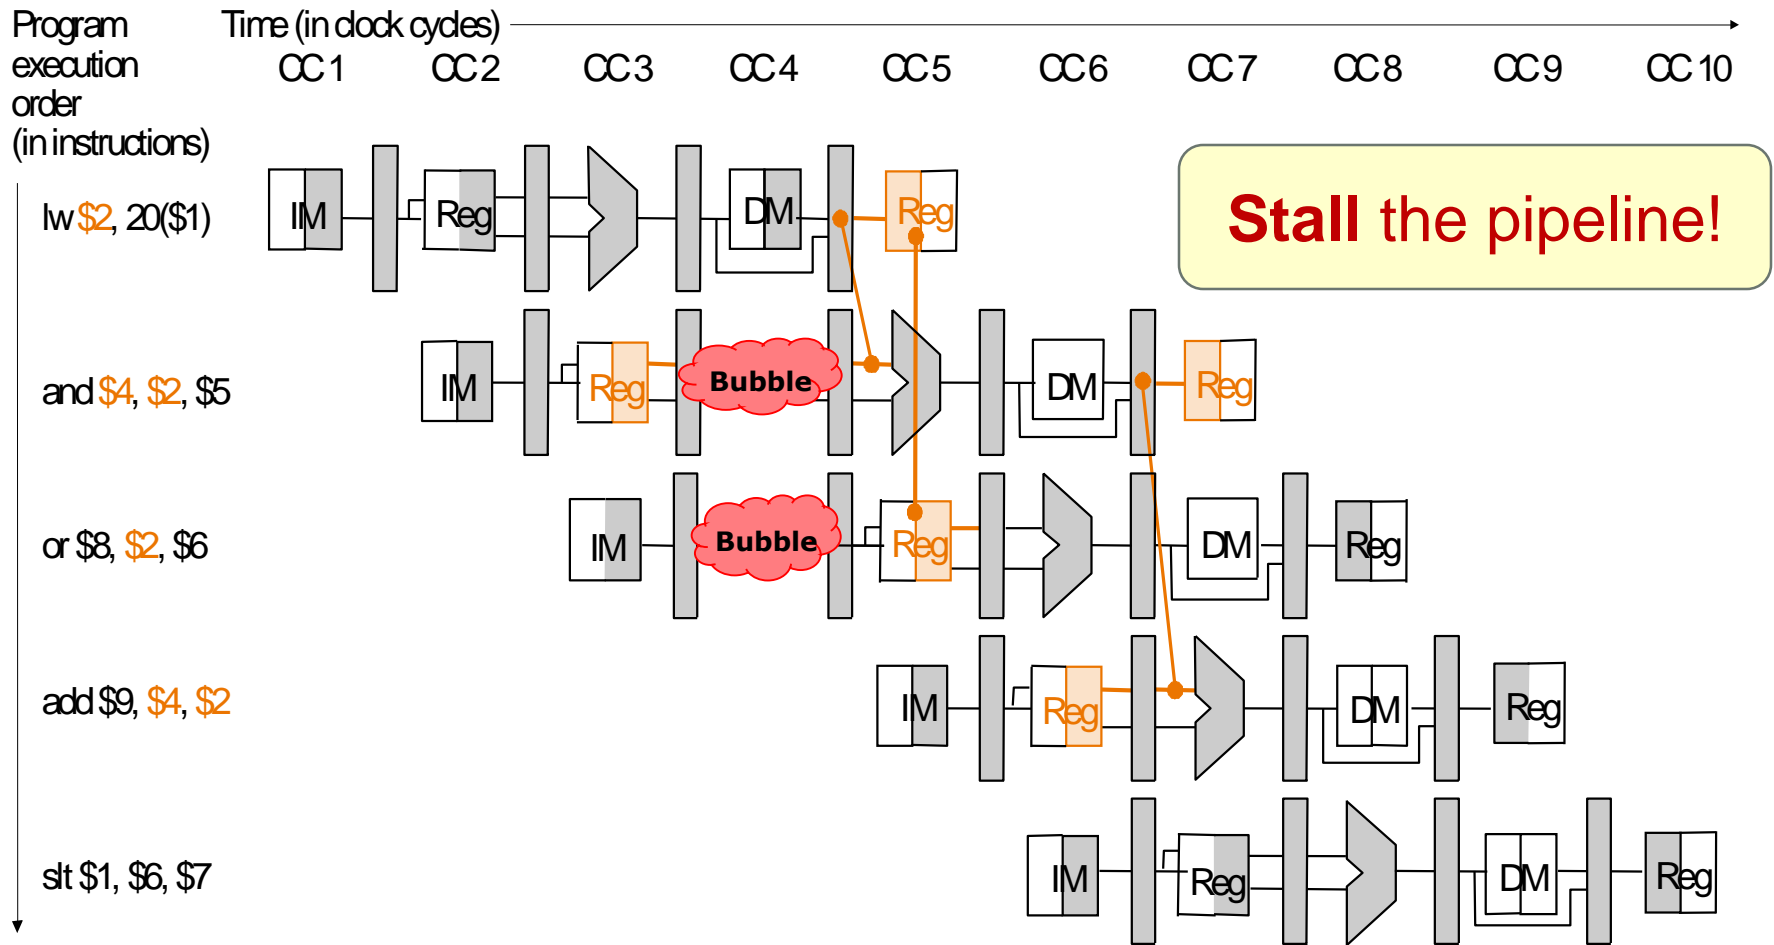
4.3 Exercise #1

- How many cycles will it take to execute the following code on a 5-stage pipeline
  - **without** forwarding?
  - **with** forwarding?

```
sub $2,  $1, $3
and $12, $2, $5
or  $13, $6, $2
add $14, $2, $2
sw  $15, 100($2)
```

4.3 Exercise #1: Without Forwarding

```
sub $2,  $1, $3
and $12, $2, $5
or  $13, $6, $2
add $14, $2, $2
sw  $15, 100($2)
```

	1	2	3	4	5	6	7	8	9	10	11
sub	IF	ID	EX	MEM	WB						
and		IF			ID	EX	MEM	WB			
or					IF	ID	EX	MEM	WB		
add						IF	ID	EX	MEM	WB	
sw							IF	ID	EX	MEM	WB

4.3 Exercise #1: With Forwarding

```
sub $2,  $1, $3
and $12, $2, $5
or  $13, $6, $2
add $14, $2, $2
sw  $15, 100($2)
```

	1	2	3	4	5	6	7	8	9	10	11
sub	IF	ID	EX	MEM	WB						
and		IF	ID	EX	MEM	WB					
or			IF	ID	EX	MEM	WB				
add				IF	ID	EX	MEM	WB			
sw					IF	ID	EX	MEM	WB		

4.3 Exercise #2

- How many cycles will it take to execute the following code on a 5-stage pipeline
  - **without** forwarding?
  - **with** forwarding?

```
lw  $2,  20($3)
and $12, $2, $5
or  $13, $6, $2
add $14, $2, $2
sw  $15, 100($2)
```

4.3 Exercise #2: Without Forwarding

```
lw   $2,  20($3)
and $12, $2, $5
or  $13, $6, $2
add $14, $2, $2
sw  $15, 100($2)
```

	1	2	3	4	5	6	7	8	9	10	11
lw	IF	ID	EX	MEM	WB						
and		IF			ID	EX	MEM	WB			
or					IF	ID	EX	MEM	WB		
add						IF	ID	EX	MEM	WB	
sw							IF	ID	EX	MEM	WB

4.3 Exercise #2: With Forwarding

```
lw   $2,  20($3)
and  $12, $2, $5
or   $13, $6, $2
add  $14, $2, $2
sw   $15, 100($2)
```

	1	2	3	4	5	6	7	8	9	10	11
lw	IF	ID	EX	MEM	WB						
and		IF	ID		EX	MEM	WB				
or			IF		ID	EX	MEM	WB			
add					IF	ID	EX	MEM	WB		
sw						IF	ID	EX	MEM	WB	

5. Control Dependency

- **Definition:**
  - An instruction `j` is control dependent on `i` if `i` controls whether or not `j` executes
  - Typically `i` would be a branch instruction
- Example:

```
i1: beq $3, $5, label    # branch
i2: add $1, $2, $4       # depends on i1
...    ...    ...
```

- Effect of incorrect execution:
  - If `i2` is allowed to execute before `i1` is determined, register `$1` maybe incorrectly changed!
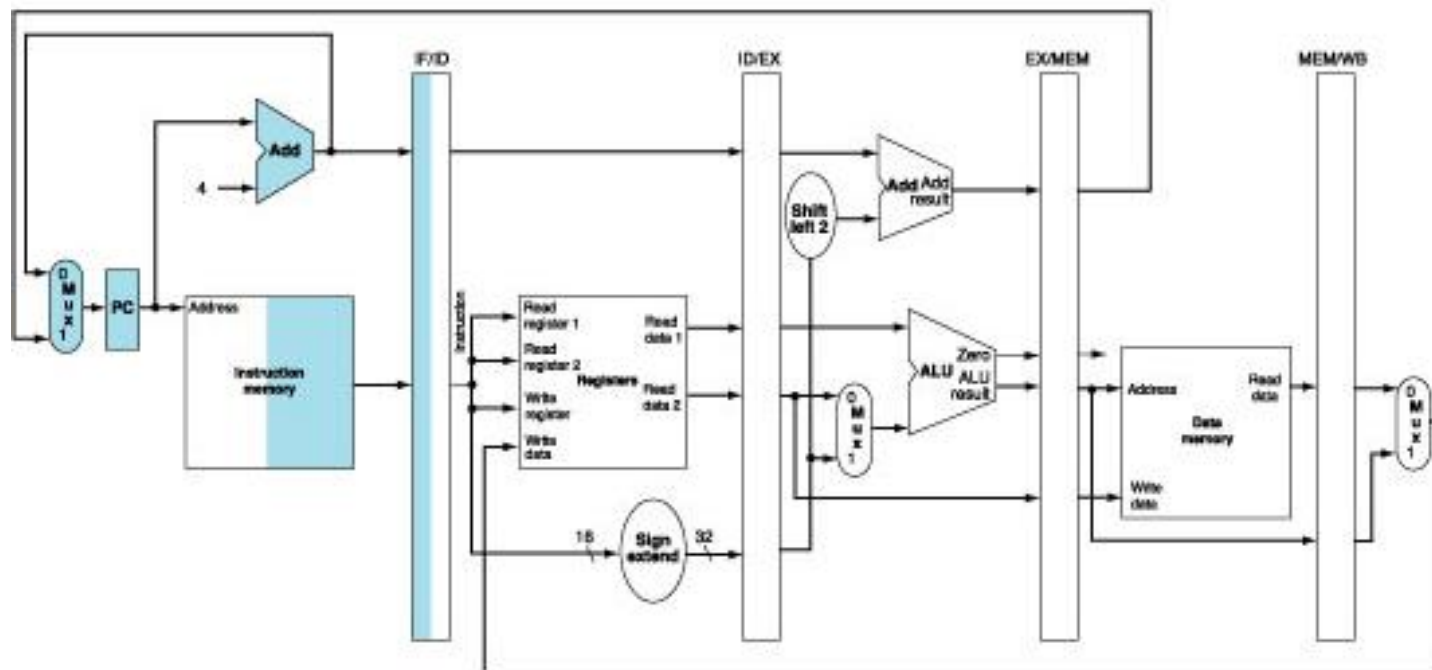
5. Control Dependency: **Example**

- Let us turn to a code fragment with a conditional branch:

$1 \neq $3

```
40 beq  $1, $3, 7
44 and $12, $2, $5
48 or   $13, $6, $2
52 add $14, $2, $2
.. ..........
72 lw  $4,  5($7)
```
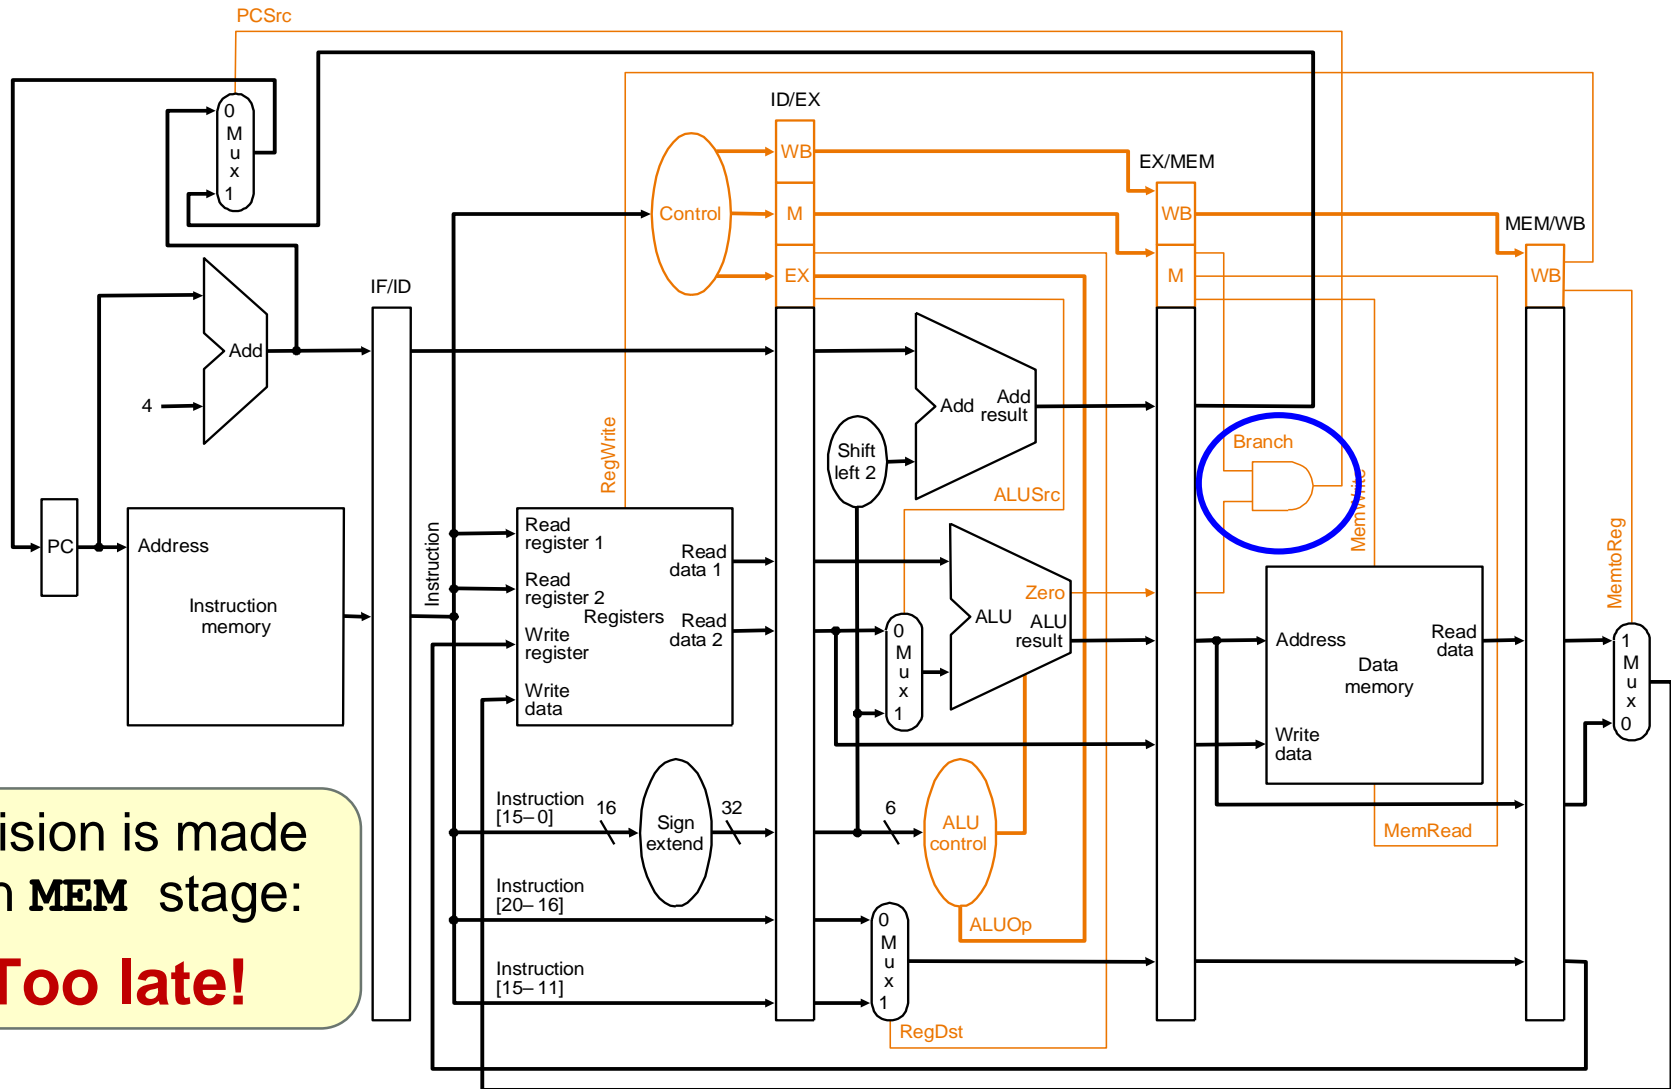
$1 = $3
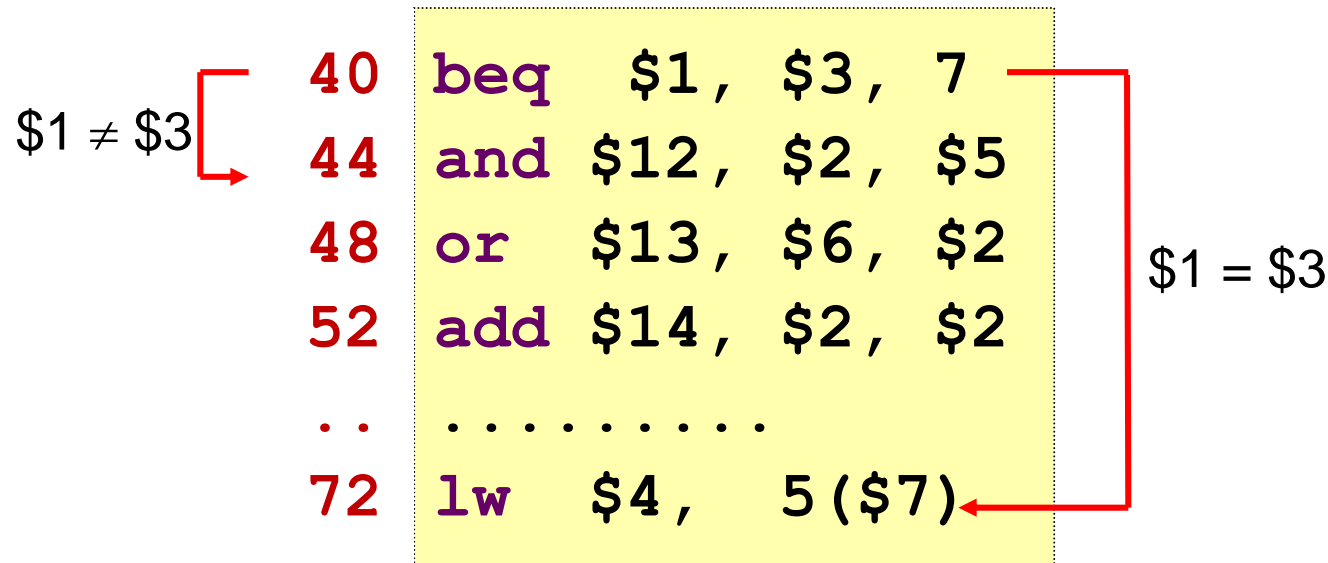
- How does the code affect a pipeline processor?
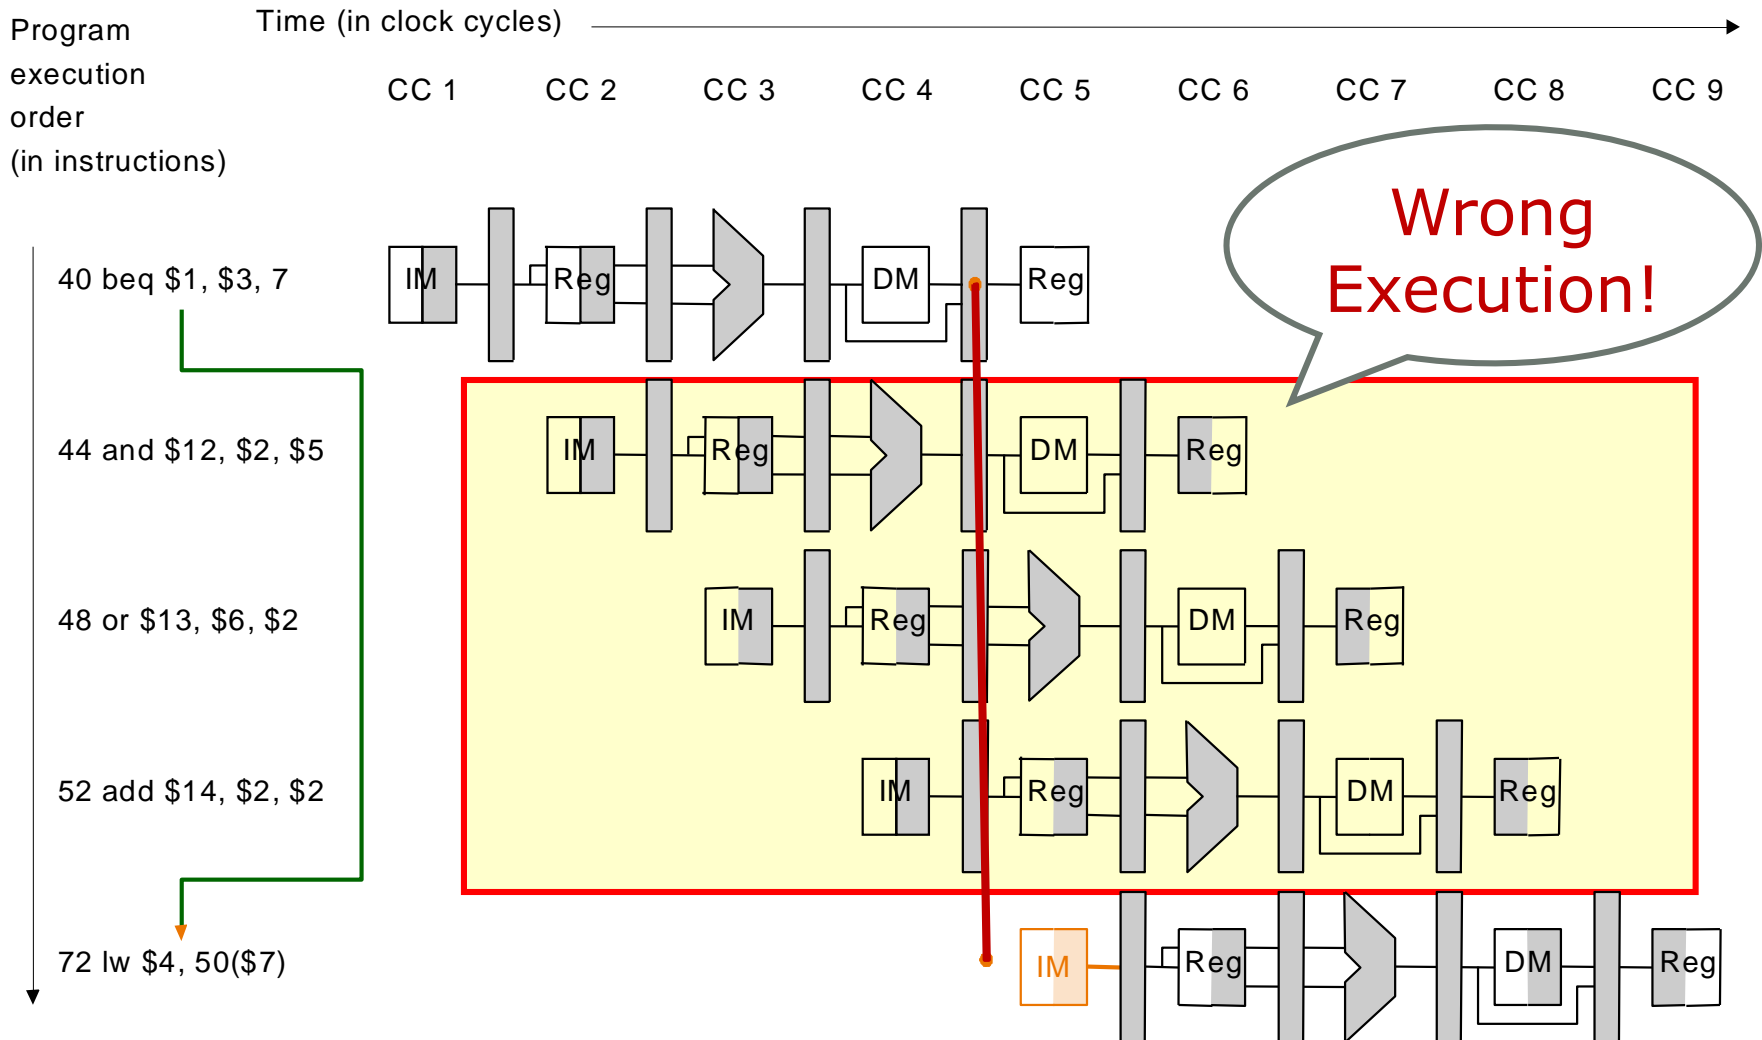
5. Pipeline Execution: **IF** Stage

- Read instruction from memory using the address in PC and put it in **IF/ID** register
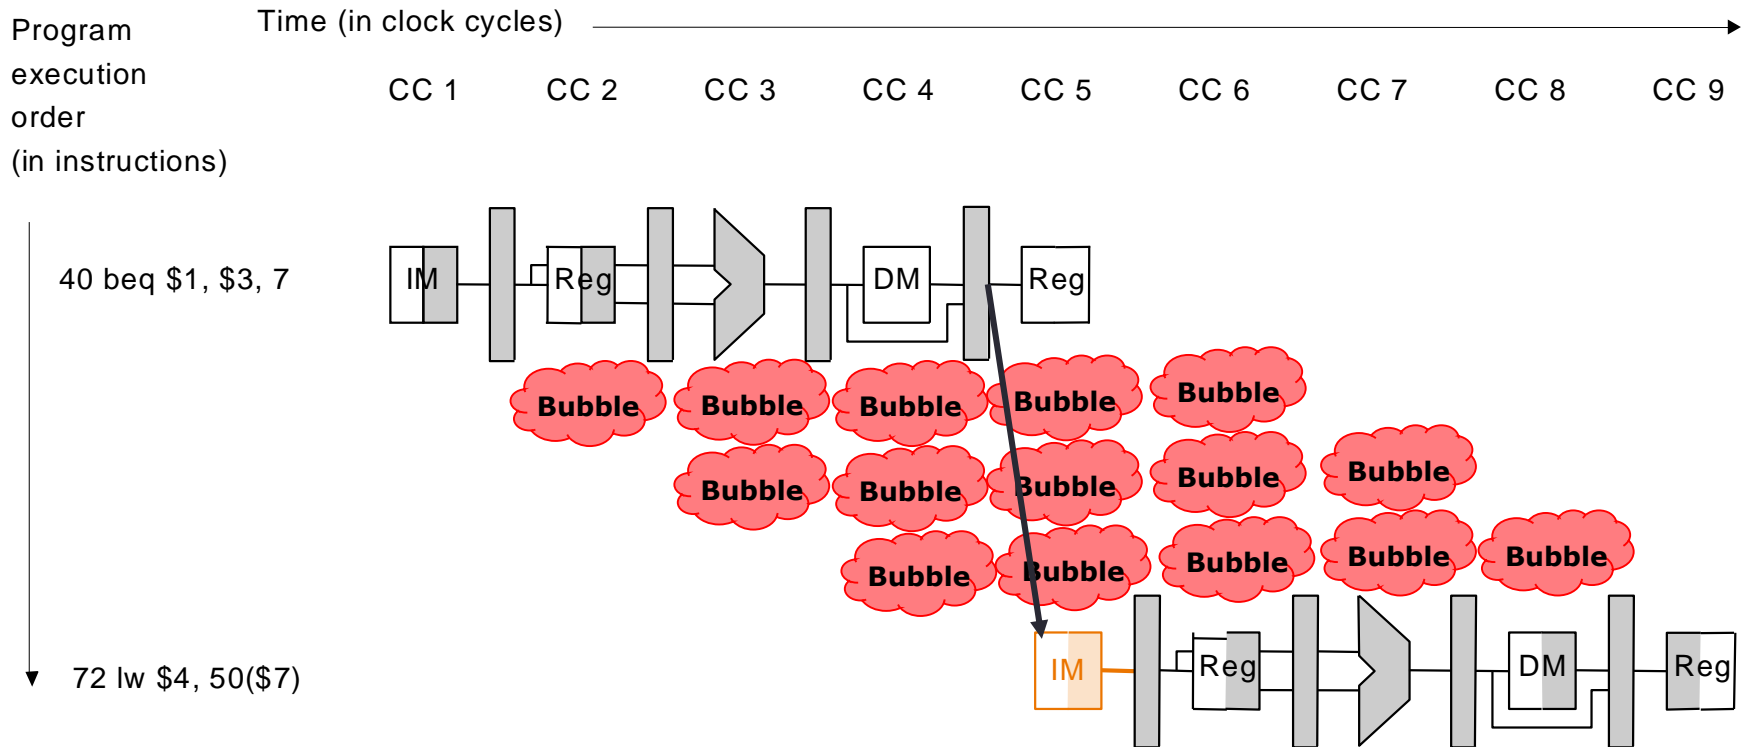- PC address is incremented by 4 and then written back to the PC for next instruction

5. Control Dependency: **Why?**

Decision is made in **MEM** stage:
Too late!

5. Control Dependency: **Example**

# 6. Control Hazards: **Stall Pipeline?**

- Wait until the branch outcome is known and then fetch the correct instructions
- → Introduces **3 clock cycles delay**

6. Control Hazards: **Reducing the Penalty**

- Branching is very common in code:
  - A 3-cycle stall penalty is too heavy!

- Many techniques invented to reduce the control hazard penalty:
  - Move branch decision calculation to earlier pipeline stage
    - **Early Branch Resolution**
  - Guess the outcome before it is produced
    - **Branch Prediction**
  - Do something useful while waiting for the outcome
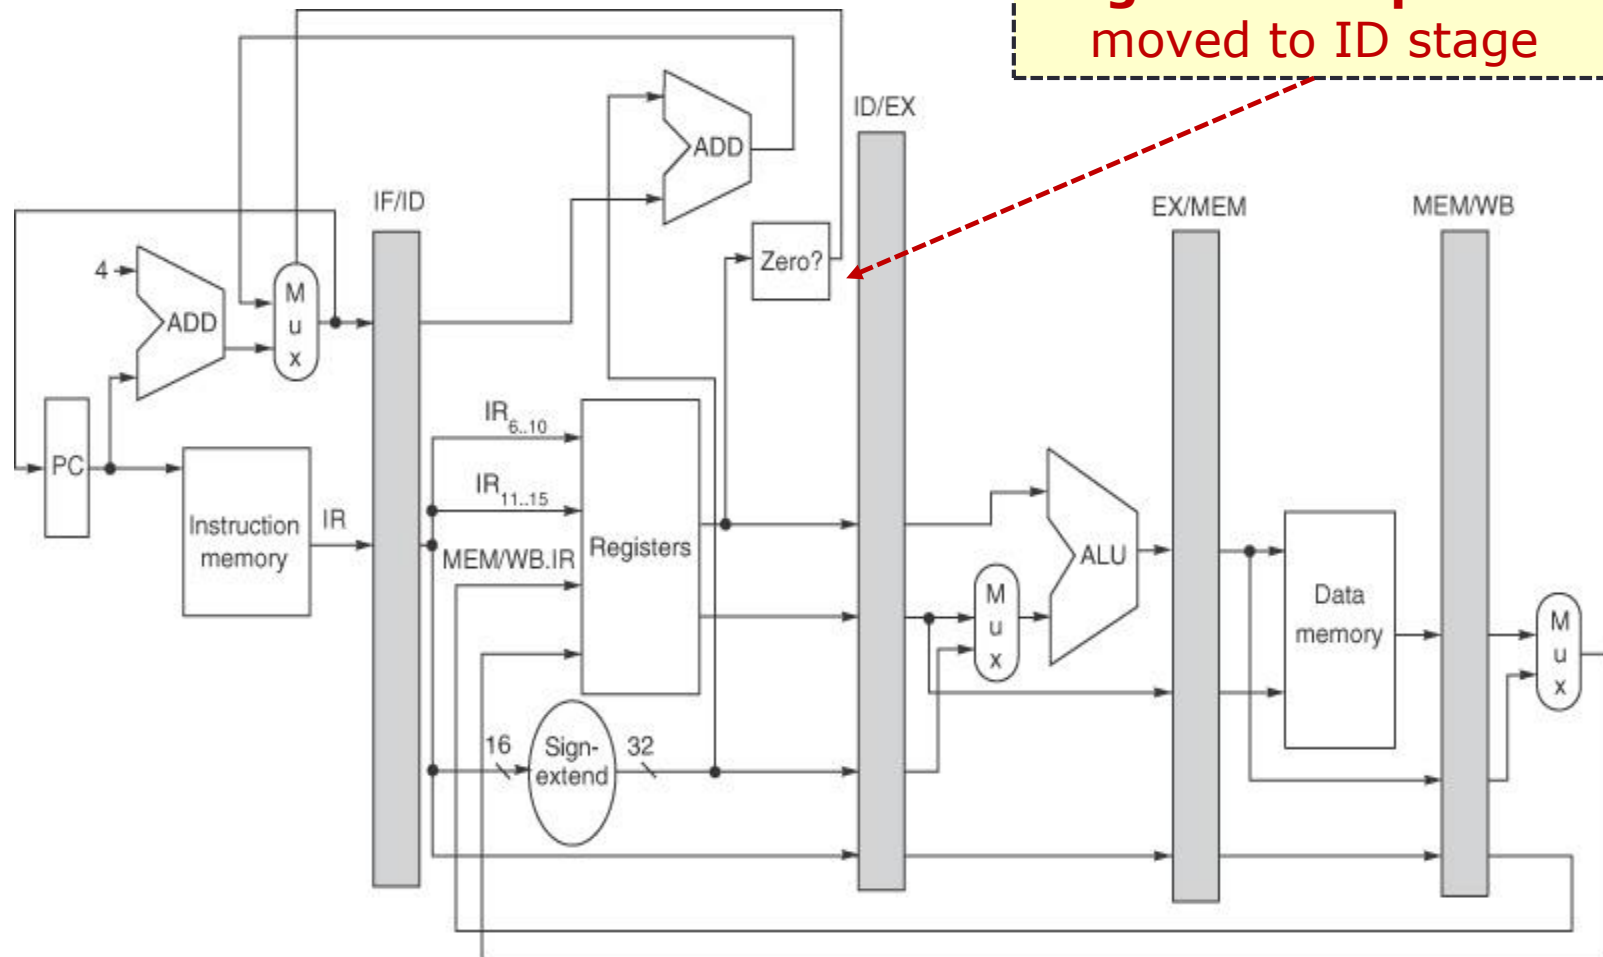   - **Delayed Branching**

6.1 Reduce Stalls: **Early Branch** (1/3)

- **Make decision in ID stage instead of MEM**
  - Move branch target address calculation
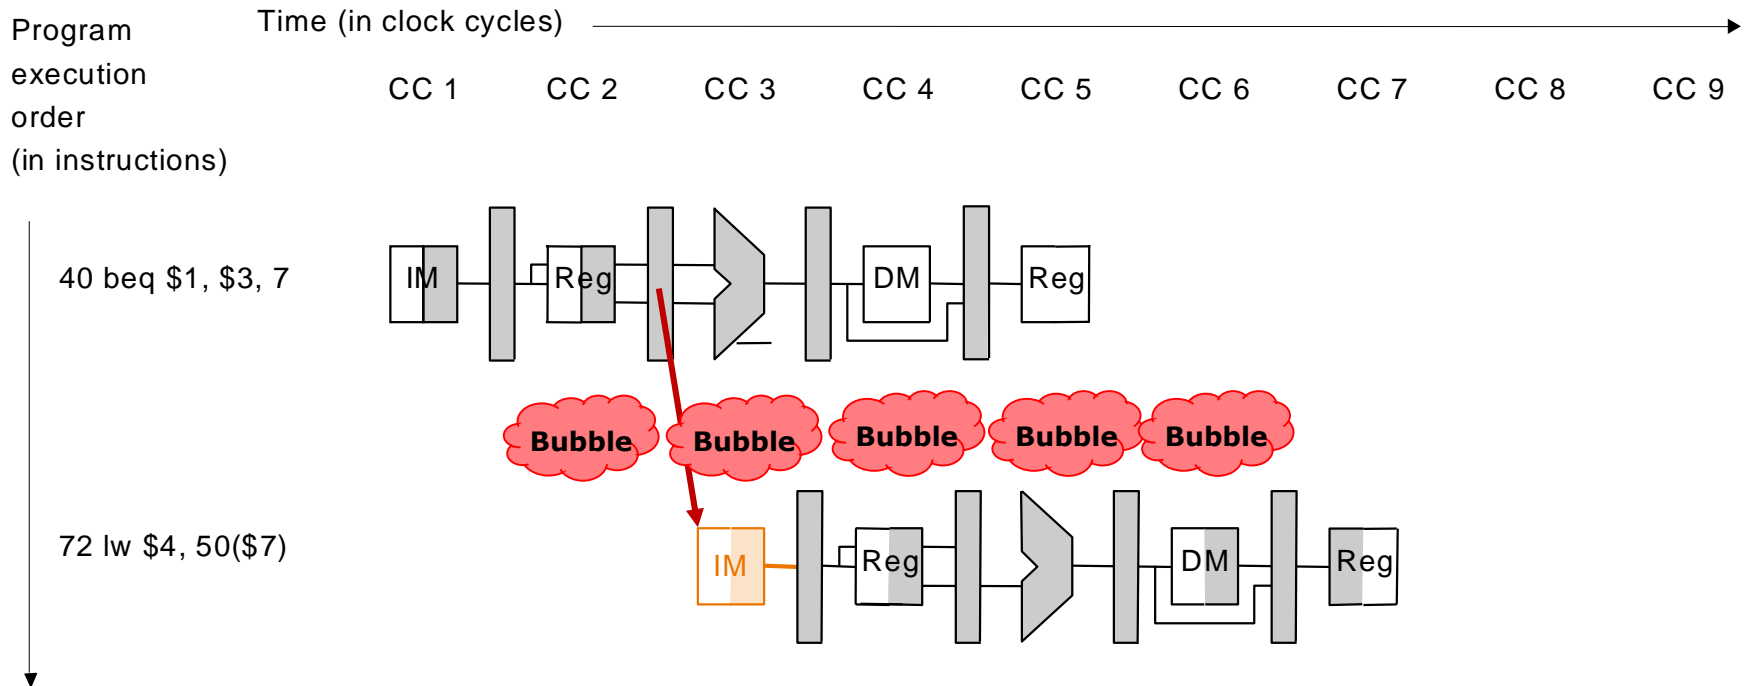 - Move register comparison ➔ cannot use ALU for register comparison any more
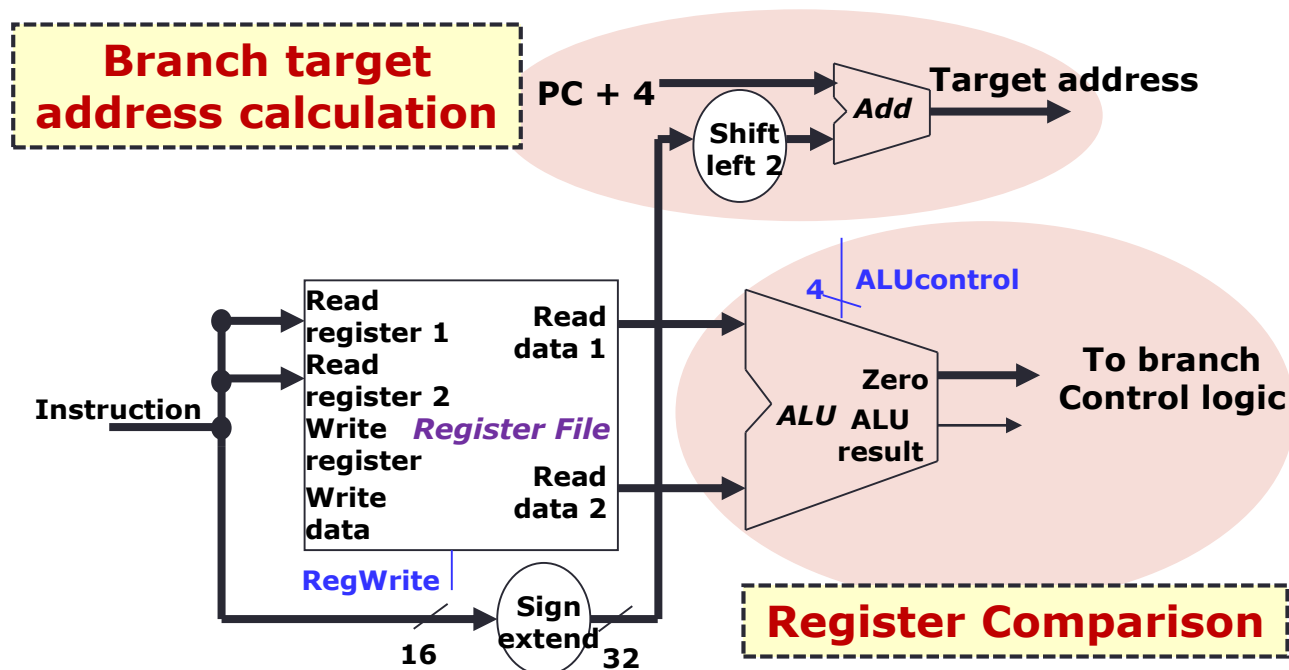
6.1 Reduce Stalls: **Early Branch** (2/3)
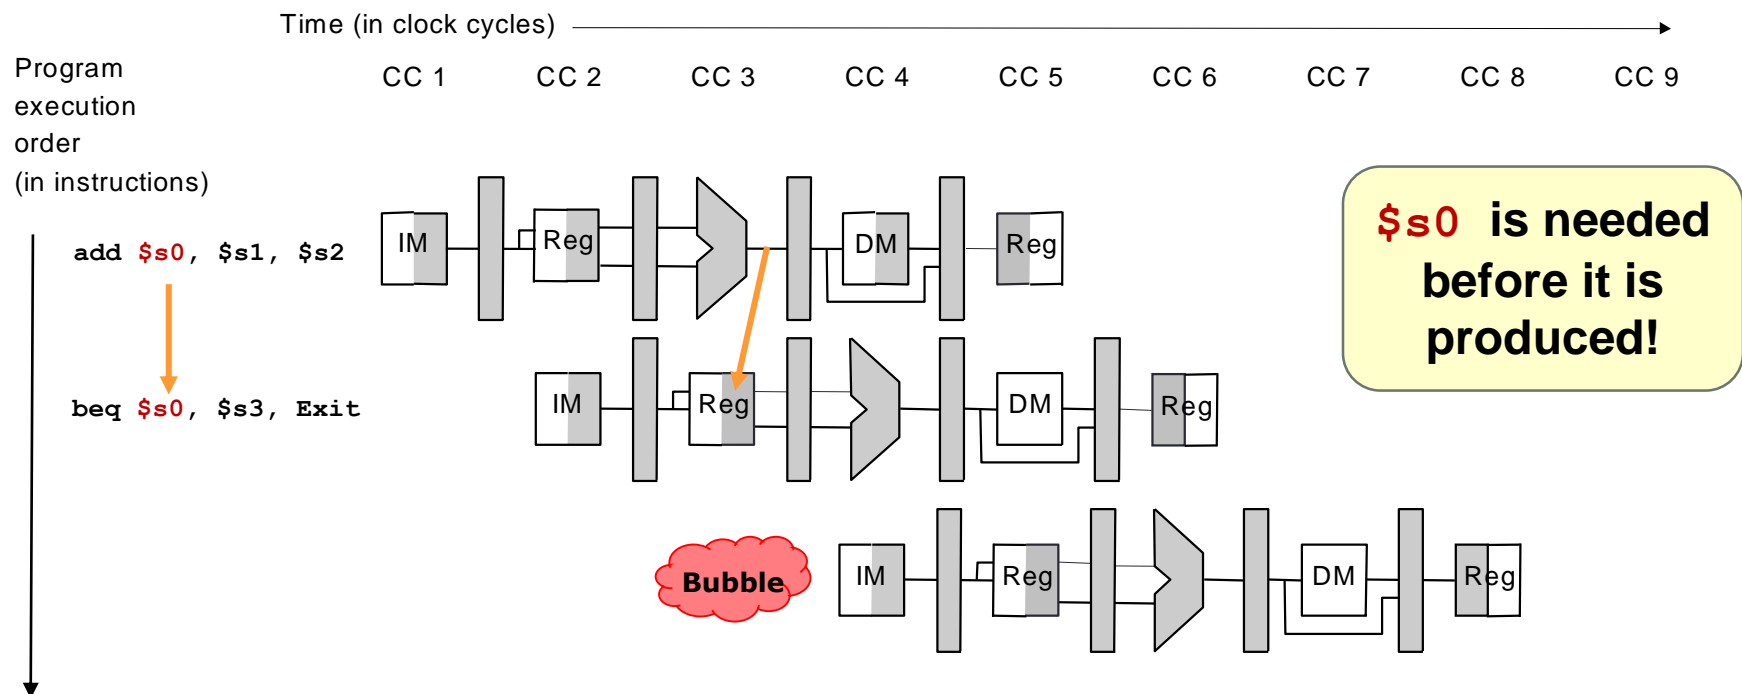
Register comparison moved to ID stage

6.1 Reduce Stalls: **Early Branch** (3/3)

- Wait until the branch decision is known:
  - Then fetch the correct instruction
- Reduced from 3 to **1 clock cycle delay**
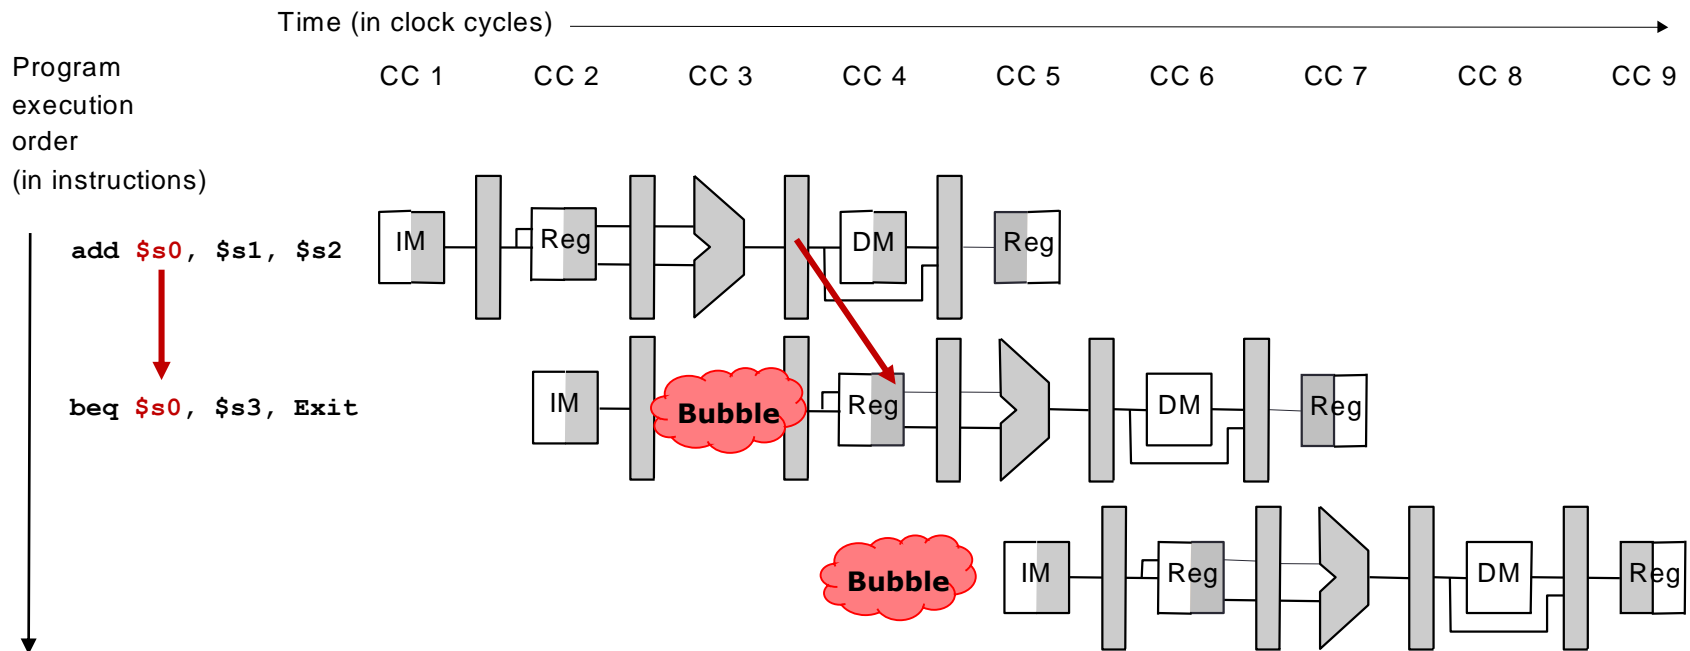
6.1 Early Branch: **Problems** (1/3)

- However, if the register(s) involved in the comparison is produced by preceding instruction:
  - Further stall is still needed!

Time (in clock cycles)

Program execution order (in instructions)

CC 1   CC 2   CC 3   CC 4   CC 5   CC 6   CC 7   CC 8   CC 9

add $s0, $s1, $s2

beq $s0, $s3, Exit

Bubble

$s0 is needed before it is produced!

# 6.1 Early Branch: **Problems** (2/3)

- ## **Solution:**
  - ### Add forwarding path from ALU to ID stage
  - ### **One clock cycle delay** is still needed

6.1 Early Branch: **Problems** (3/3)

- Problem is worse with **load** followed by **branch**
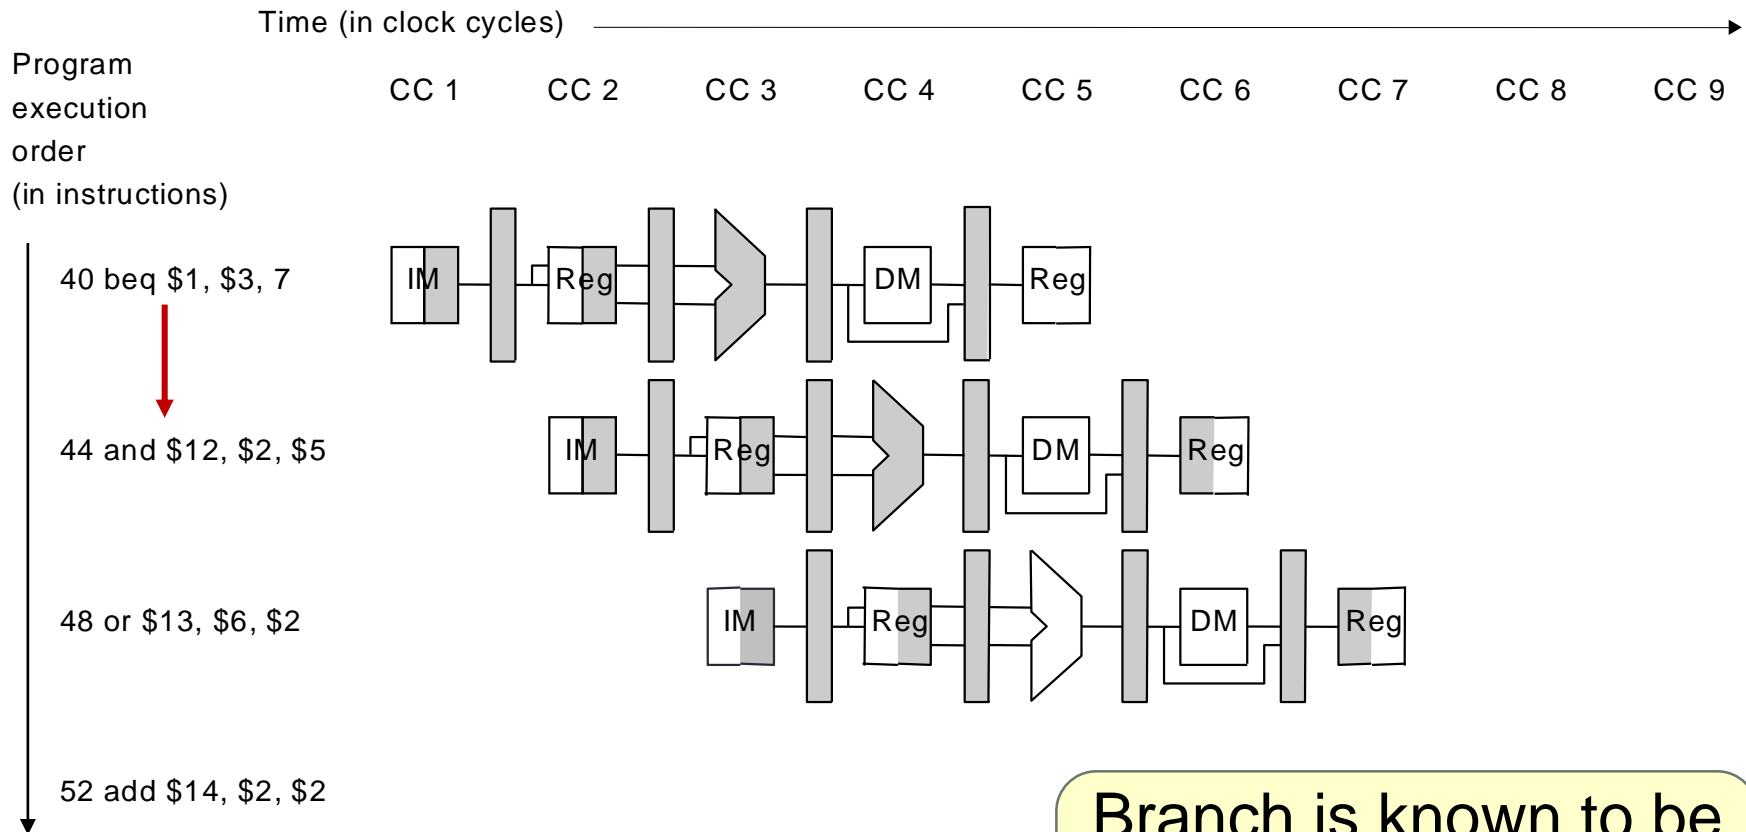
- **Solution:**
  - MEM to ID forwarding and 2 more stall cycles!
  - In this case, we ended up with 3 total stall cycles ➔ no improvement!

Time (in clock cycles)

CC 1   CC 2   CC 3   CC 4   CC 5   CC 6   CC 7   CC 8   CC 9

Program execution order (in instructions)

lw $s0, 0($s1)

beq $s0, $s3, Exit

ALU ➔ ID forwarding cannot help!

Bubble   Bubble   Bubble

6.2 Reduce Stalls: **Branch Prediction**

- There are many branch prediction schemes
  - We only cover the simplest in this course ☺

- Simple prediction:
  - All branches are assumed to be **not taken**
  - → Fetch the successor instruction and start pumping it through the pipeline stages

- When the actual branch outcome is known:
  - **Not taken**: Guessed correctly ➔ No pipeline stall
  - **Taken**: Guessed wrongly ➔ Wrong instructions in the pipeline ➔ **Flush** successor instruction from the pipeline

6.2 Branch Prediction: **Correct Prediction**

Time (in clock cycles)

Program
execution
order
(in instructions)

| CC 1 | CC 2 | CC 3 | CC 4 | CC 5 | CC 6 | CC 7 | CC 8 | CC 9 |

40 beq $1, $3, 7

      IM   Reg   DM   Reg

44 and $12, $2, $5

          IM   Reg   DM   Reg

48 or $13, $6, $2

                IM   Reg   DM   Reg

52 add $14, $2, $2

Branch is known to be
not taken in cycle 3
➔ no stall needed!

6.2 Branch Prediction: **Wrong Prediction**

Time (in clock cycles)

Program execution order (in instructions)

CC 1   CC 2   CC 3   CC 4   CC 5   CC 6   CC 7   CC 8   CC 9

40 beq $1, $3, 7

44 and $12, $2, $5

72 lw $4, 50($7)

Bubble   Bubble   Bubble   Bubble
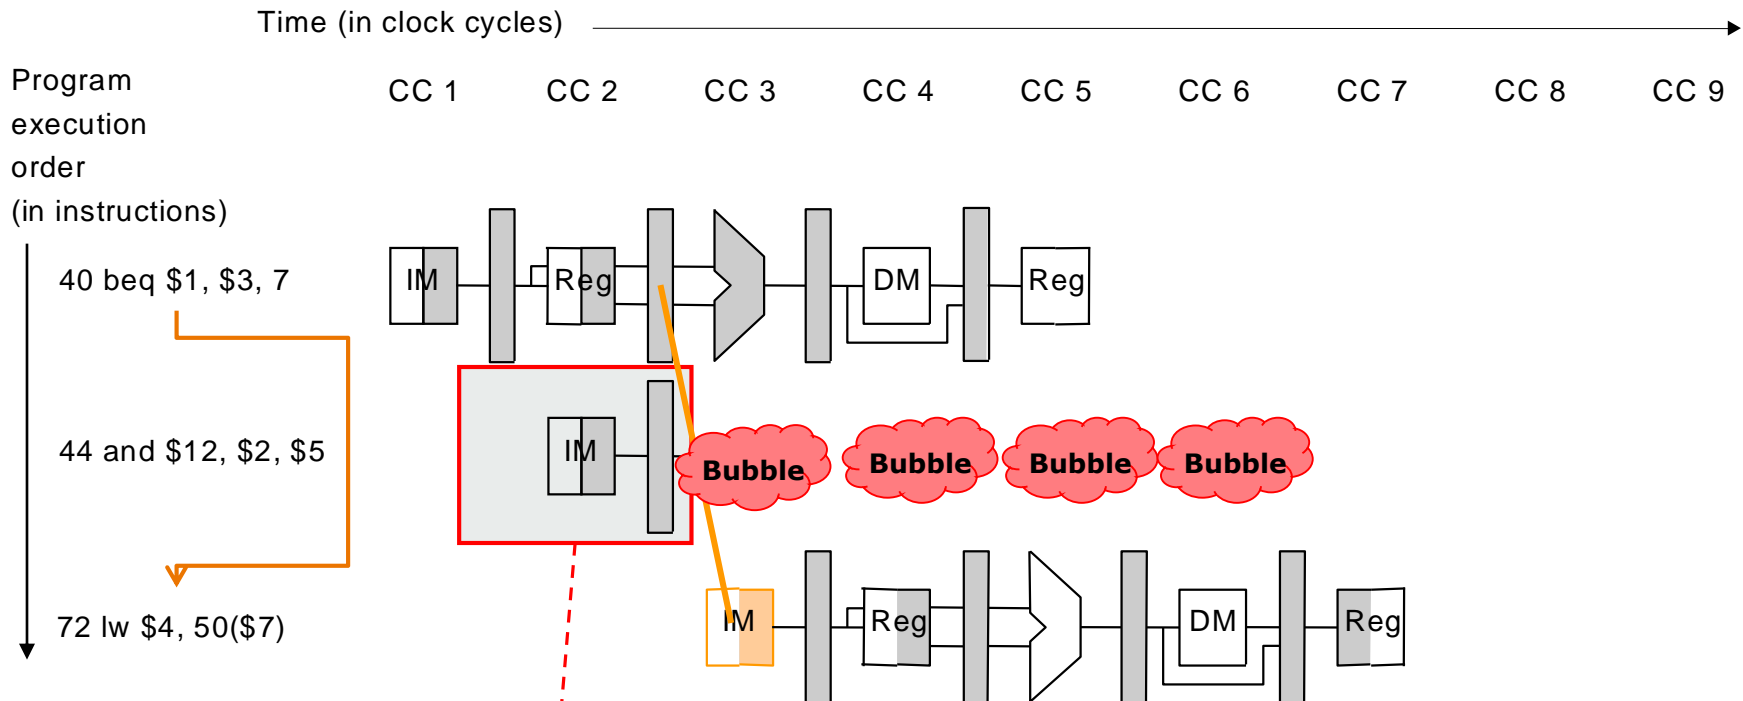
Branch is known to be **taken** in cycle 3
➔ "**and**" instruction should not be executed
➔ Flushed from pipeline

6.2 Exercise #3: **Branch Prediction**

- How many cycles will it take to execute the following code on a 5-stage pipeline **with** forwarding and …
  - **without** branch prediction?
  - **with** branch prediction (predict not taken)?

```
        addi $s0, $zero, 10
Loop:   addi $s0, $s0, -1
        bne  $s0, $zero, Loop
        sub  $t0, $t1, $t2
```

- Decision making moved to **ID** stage
- Total instructions = $1 + 10 \times 2 + 1 = 22$
- **Ideal** pipeline = $4 + 22 = 26$ cycles

6.2 Exercise #3: Without Branch Prediction

	1	2	3	4	5	6	7	8	9	10	11
addi¹	IF	ID	EX	MEM	WB						
addi²		IF	ID	EX	MEM	WB					
bne			IF		ID	EX	MEM	WB			
addi²						IF	ID	EX	MEM	WB	

- Data dependency between (`addi $s0, $s0, -1`) and `bne` incurs 1 cycle of delay. There are 10 iterations, hence 10 cycles of delay.
- Every `bne` incurs a cycle of delay to execute the next instruction. There are 10 iterations, hence 10 cycles of delay.
- Total number of cycles of delay = 20.
- Total execution cycles = 26 + 20 = **46 cycles.**

6.2 Exercise #3: With Branch Prediction

	1	2	3	4	5	6	7	8	9	10	11
$addi^1$	IF	ID	EX	MEM	WB						
$addi^2$		IF	ID	EX	MEM	WB					
bne			IF		ID	EX	MEM	WB			
sub					IF						
$addi^2$						IF	ID	EX	MEM	WB	

Predict not taken.
- The data dependency remains, hence 10 cycles of delay for 10 iterations.
- In the first 9 iterations, the branch prediction is wrong, hence 1 cycle of delay.
- In the last iteration, the branch prediction is correct, hence saving 1 cycle of delay.
- Total number of cycles of delay = 19.
- Total execution cycles = 26 + 19 = **45 cycles.**

6.3 Reduce Stalls: **Delayed Branch**

- **Observation:**
  - Branch outcome takes **X** number of cycles to be known
  - ➔ **X** cycles stall

- **Idea:**
  - Move **non-control dependent instructions** into the X slots following a branch
    - Known as the **branch-delay slot**
  - ➔ These instructions are executed **regardless of the branch outcome**
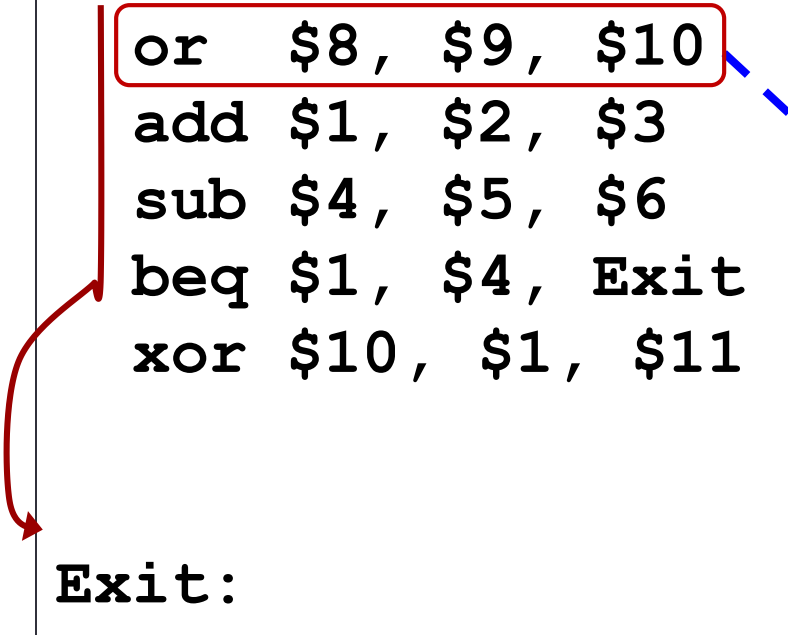
- In our MIPS processor:
  - Branch-Delay slot = **1** (with the early branch)
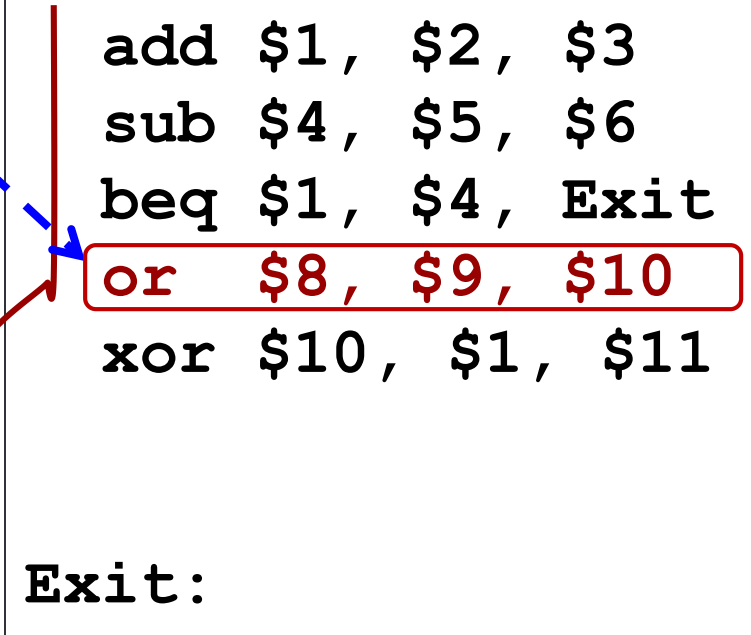
6.3 Delayed Branch: **Example**

Non-delayed branch	*Delayed branch*
| ```
or   $8, $9, $10
add $1, $2, $3
sub $4, $5, $6
beq $1, $4, Exit
xor $10, $1, $11


Exit:
``` | ```
add $1, $2, $3
sub $4, $5, $6
beq $1, $4, Exit
or   $8, $9, $10
xor $10, $1, $11


Exit:
``` |

- The "**or**" instruction is moved into the delayed slot:
  - Get executed regardless of the branch outcome
  - ➔ Same behavior as the original code!

6.3 Delayed Branch: **Observation**

- **Best case scenario**
  - There is an instruction **preceding the branch** which **can be moved** into the delayed slot
    - Program correctness must be preserved!

- **Worst case scenario**
  - Such instruction cannot be found
  - ➔ Add a no-op (**nop**) instruction in the branch-delay slot

- Re-ordering instructions is a common method of program optimization
  - Compiler must be smart enough to do this
  - Usually can find such an instruction at least 50% of the time
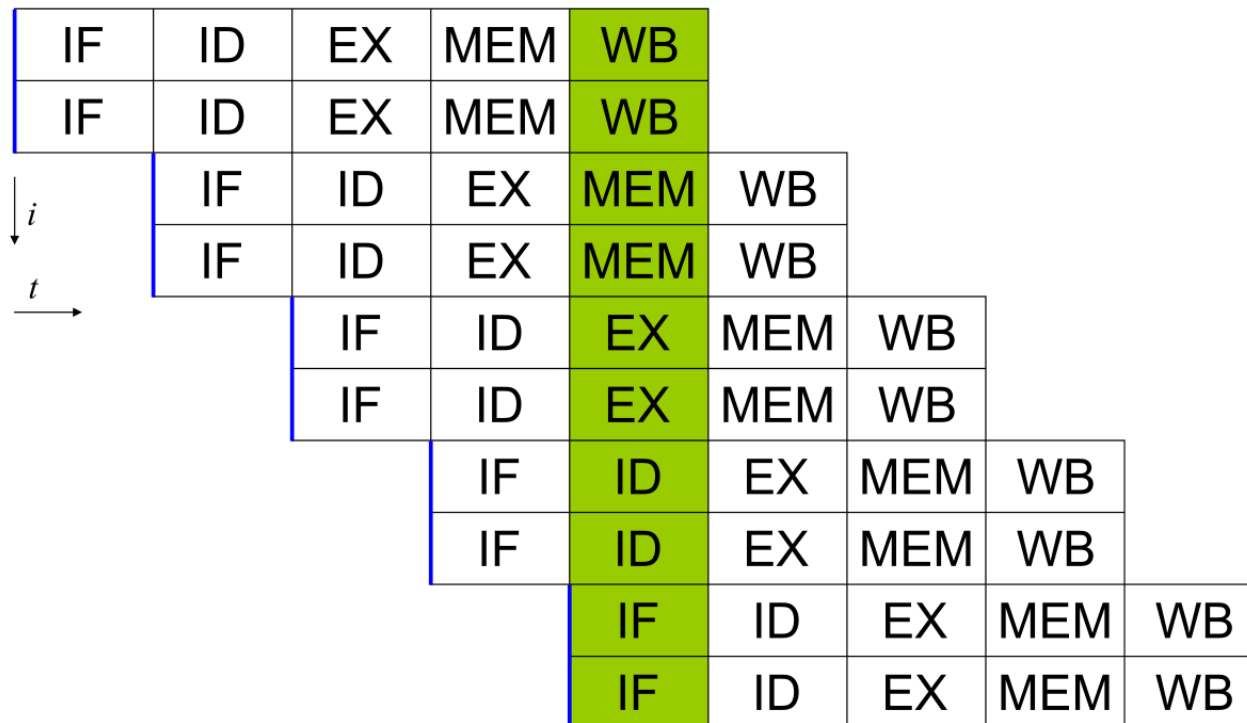
7. Multiple Issue Processors (1/2) For reading only

- Multiple Issue processors
  - **Multiple instructions** in every pipeline stage
  - 4 washer, 4 dryer…

- **Static multiple issue**:
  - EPIC (Explicitly Parallel Instruction Computer) or VLIW (Very Long Instruction Word), e.g. IA64
  - Compiler specifies the set of instructions that execute together in a given clock cycle
  - Simple hardware, complex compiler

- **Dynamic multiple issue:**
  - Superscalar processor: Dominant design of modern processors
  - Hardware decides which instructions to execute together
  - Complex hardware, simpler compiler

7. Multiple Issue Processors (2/2)

For reading only

- A 2-wide superscalar pipeline:
  - By fetching and dispatching two instructions at a time, a maximum of two instructions per cycle can be completed.

IF	ID	EX	MEM	WB				
IF	ID	EX	MEM	WB				
	IF	ID	EX	MEM	WB			
	IF	ID	EX	MEM	WB			
		IF	ID	EX	MEM	WB		
		IF	ID	EX	MEM	WB		
			IF	ID	EX	MEM	WB	
			IF	ID	EX	MEM	WB	
				IF	ID	EX	MEM	WB
				IF	ID	EX	MEM	WB

i ↓

t →

Summary

- Pipelining is a fundamental concept in computer systems
  - Multiple instructions in flight
  - Limited by length of the longest stage
  - Hazards create trouble by stalling pipeline

- Pentium 4 has 22 pipeline stages!

Reading

- ## 3rd edition
  - Sections 6.1 – 6.3
  - Sections 6.4 – 6.6 (data hazards and control hazards in details; read for interest; not in syllabus)

- ## 4th edition
  - Sections 4.5 – 4.6
  - Sections 4.7 – 4.8 (data hazards and control hazards in details; read for interest; not in syllabus)

End of File