



## 6.1 Flip-flop Characteristic Tables

- Each type of flip-flop has its own behaviour, shown by its **characteristic table**.

$J$	$K$	$Q(t+1)$	Comments
0	0	$Q(t)$	No change
0	1	0	Reset
1	0	1	Set
1	1	$Q(t)'$	Toggle

$S$	$R$	$Q(t+1)$	Comments
0	0	$Q(t)$	No change
0	1	0	Reset
1	0	1	Set
1	1	?	Unpredictable

$D$	$Q(t+1)$
0	0      Reset
1	1      Set

$T$	$Q(t+1)$
0	$Q(t)$ No change
1	$Q(t)'$ Toggle

## 6.3 Flip-flop Excitation Tables (1/2)

- **Excitation tables**: given the required transition from present state to next state, determine the flip-flop input(s).

$Q$	$Q^+$	$J$	$K$
0	0	0	X
0	1	1	X
1	0	X	1
1	1	X	0

$JK$  Flip-flop

$Q$	$Q^+$	$S$	$R$
0	0	0	X
0	1	1	0
1	0	0	1
1	1	X	0

$SR$  Flip-flop

$Q$	$Q^+$	$D$
0	0	0
0	1	1
1	0	0
1	1	1

$D$  Flip-flop

$Q$	$Q^+$	$T$
0	0	0
0	1	1
1	0	1
1	1	0

$T$  Flip-flop

# MIPS Reference Data

①



## ARITHMETIC CORE INSTRUCTION SET

② OPCODE

### CORE INSTRUCTION SET

NAME, MNEMONIC	FOR-MAT	OPERATION (in Verilog)	OPCODE / FUNCT (Hex)
Add	add R	R[rd] = R[rs] + R[rt]	(1) 0 / 20 <sub>hex</sub>
Add Immediate	addi I	R[rt] = R[rs] + SignExtImm	(1,2) 8 <sub>hex</sub>
Add Imm. Unsigned	addiu I	R[rt] = R[rs] + SignExtImm	(2) 9 <sub>hex</sub>
Add Unsigned	addu R	R[rd] = R[rs] + R[rt]	(0) 21 <sub>hex</sub>
And	and R	R[rd] = R[rs] & R[rt]	0 / 24 <sub>hex</sub>
And Immediate	andi I	R[rt] = R[rs] & ZeroExtImm	(3) c <sub>hex</sub>
Branch On Equal	beq I	if(R[rs]==R[rt]) PC=PC+4+BranchAddr	(4) 4 <sub>hex</sub>
Branch On Not Equal	bne I	if(R[rs]!=R[rt]) PC=PC+4+BranchAddr	(4) 5 <sub>hex</sub>
Jump	j J	PC=JumpAddr	(5) 2 <sub>hex</sub>
Jump And Link	jal J	R[31]=PC+8; PC=JumpAddr	(5) 3 <sub>hex</sub>
Jump Register	jr R	PC=R[rs]	0 / 08 <sub>hex</sub>
Load Byte Unsigned	lbu I	R[rt] = {24'b0, M[R[rs] +SignExtImm](7:0)}	(2) 24 <sub>hex</sub>
Load Halfword Unsigned	lhu I	R[rt] = {16'b0, M[R[rs] +SignExtImm](15:0)}	(2) 25 <sub>hex</sub>
Load Linked	ll I	R[rt] = M[R[rs]+SignExtImm]	(2,7) 30 <sub>hex</sub>
Load Upper Imm.	lui I	R[rt] = {imm, 16'b0}	f <sub>hex</sub>
Load Word	lw I	R[rt] = M[R[rs]+SignExtImm]	(2) 23 <sub>hex</sub>
Nor	nor R	R[rd] = ~ (R[rs]   R[rt])	0 / 27 <sub>hex</sub>
Or	or R	R[rd] = R[rs]   R[rt]	0 / 25 <sub>hex</sub>
Or Immediate	ori I	R[rt] = R[rs]   ZeroExtImm	(3) d <sub>hex</sub>
Set Less Than	slt R	R[rd] = (R[rs] < R[rt]) ? 1 : 0	0 / 2a <sub>hex</sub>
Set Less Than Imm.	slti I	R[rt] = (R[rs] < SignExtImm) ? 1 : 0	a <sub>hex</sub>
Set Less Than Imm. Unsigned	sltiu I	R[rt] = (R[rs] < SignExtImm) ? 1 : 0	(2,6) b <sub>hex</sub>
Set Less Than Unsig.	sltu R	R[rd] = (R[rs] < R[rt]) ? 1 : 0	(6) 0 / 2b <sub>hex</sub>
Shift Left Logical	sll R	R[rd] = R[rt] << shamt	0 / 00 <sub>hex</sub>
Shift Right Logical	srl R	R[rd] = R[rt] >> shamt	0 / 02 <sub>hex</sub>
Store Byte	sb I	M[R[rs]+SignExtImm](7:0) = R[rt](7:0)	(2) 28 <sub>hex</sub>
Store Conditional	sc I	M[R[rs]+SignExtImm] = R[rt]; R[rt] = (atomic) ? 1 : 0	(2,7) 38 <sub>hex</sub>
Store Halfword	sh I	M[R[rs]+SignExtImm](15:0) = R[rt](15:0)	(2) 29 <sub>hex</sub>
Store Word	sw I	M[R[rs]+SignExtImm] = R[rt]	(2) 2b <sub>hex</sub>
Subtract	sub R	R[rd] = R[rs] - R[rt]	(1) 0 / 22 <sub>hex</sub>
Subtract Unsigned	subu R	R[rd] = R[rs] - R[rt]	0 / 23 <sub>hex</sub>

- (1) May cause overflow exception
- (2) SignExtImm = { 16{immediate[15]}, immediate }
- (3) ZeroExtImm = { 16{1b'0}, immediate }
- (4) BranchAddr = { 14{immediate[15]}, immediate, 2'b0 }
- (5) JumpAddr = { PC+4[31:28], address, 2'b0 }
- (6) Operands considered unsigned numbers (vs. 2's comp.)
- (7) Atomic test&set pair; R[rt] = 1 if pair atomic, 0 if not atomic

### BASIC INSTRUCTION FORMATS

R	opcode	rs	rt	rd	shamt	funct
	31	26 25	21 20	16 15	11 10	6 5
	0					
I	opcode	rs	rt	immediate		
	31	26 25	21 20	16 15		
	0					
J	opcode	address				
	31	26 25				
	0					

NAME, MNEMONIC	FOR-MAT	OPERATION	OPCODE / FUNCT (Hex)
Branch On FP True	bclt FI	if(FPcond)PC=PC+4+BranchAddr	(4) 11/8/1/--
Branch On FP False	bclif FI	if(!FPcond)PC=PC+4+BranchAddr	(4) 11/8/0/--
Divide	div R	Lo=R[rs]/R[rt]; Hi=R[rs]%R[rt]	0/--/--/1a
Divide Unsigned	divu R	Lo=R[rs]/R[rt]; Hi=R[rs]%R[rt]	(6) 0/--/--/1b
FP Add Single	add.s FR	F[fd] = F[fs] + F[ft]	11/10/--/0
FP Add Double	add.d FR	{F[fd],F[fd+1]} = {F[fs],F[fs+1]} + {F[ft],F[ft+1]}	11/11/--/0
FP Compare Single	c.x.s* FR	FPcond = (F[fs] op F[ft]) ? 1 : 0	11/10/--/y
FP Compare Double	c.x.d* FR	FPcond = ({F[fs],F[fs+1]} op {F[ft],F[ft+1]}) ? 1 : 0	11/11/--/y
* (x is eq, lt, or le) (op is ==, <, or <=) (y is 32, 3c, or 3e)			
FP Divide Single	div.s FR	F[fd] = F[fs] / F[ft]	11/10/--/3
FP Divide Double	div.d FR	{F[fd],F[fd+1]} = {F[fs],F[fs+1]} / {F[ft],F[ft+1]}	11/11/--/3
FP Multiply Single	mul.s FR	F[fd] = F[fs] * F[ft]	11/10/--/2
FP Multiply Double	mul.d FR	{F[fd],F[fd+1]} = {F[fs],F[fs+1]} * {F[ft],F[ft+1]}	11/11/--/2
FP Subtract Single	sub.s FR	F[fd] = F[fs] - F[ft]	11/10/--/1
FP Subtract Double	sub.d FR	{F[fd],F[fd+1]} = {F[fs],F[fs+1]} - {F[ft],F[ft+1]}	11/11/--/1
Load FP Single	lwc1 I	F[rt] = M[R[rs]+SignExtImm]	(2) 31/--/--/0
Load FP Double	ldc1 I	F[rt] = M[R[rs]+SignExtImm]; F[rt+1] = M[R[rs]+SignExtImm+4]	(2) 35/--/--/0
Move From Hi	mfhi R	R[rd] = Hi	0/--/--/10
Move From Lo	mfl0 R	R[rd] = Lo	0/--/--/12
Move From Control	mfc0 R	R[rd] = CR[rs]	10/0/--/0
Multiply	mult R	{Hi,Lo} = R[rs] * R[rt]	0/--/--/18
Multiply Unsigned	multu R	{Hi,Lo} = R[rs] * R[rt]	(6) 0/--/--/19
Shift Right Arith.	sra R	R[rd] = R[rt] >>> shamt	0/--/--/3
Store FP Single	swc1 I	M[R[rs]+SignExtImm] = F[rt]	(2) 39/--/--/0
Store FP Double	sdc1 I	M[R[rs]+SignExtImm] = F[rt]; M[R[rs]+SignExtImm+4] = F[rt+1]	(2) 3d/--/--/0

### FLOATING-POINT INSTRUCTION FORMATS

FR	opcode	fmt	ft	fs	fd	funct
	31	26 25	21 20	16 15	11 10	6 5
	0					
FI	opcode	fmt	ft	immediate		
	31	26 25	21 20	16 15		
	0					

### PSEUDOINSTRUCTION SET

NAME	MNEMONIC	OPERATION
Branch Less Than	blt	if(R[rs]<R[rt]) PC = Label
Branch Greater Than	bgt	if(R[rs]>R[rt]) PC = Label
Branch Less Than or Equal	b1e	if(R[rs]<=R[rt]) PC = Label
Branch Greater Than or Equal	bge	if(R[rs]>=R[rt]) PC = Label
Load Immediate	li	R[rd] = immediate
Move	move	R[rd] = R[rs]

### REGISTER NAME, NUMBER, USE, CALL CONVENTION

NAME	NUMBER	USE	PRESERVEDACROSS A CALL?
\$zero	0	The Constant Value 0	N.A.
\$at	1	Assembler Temporary	No
\$v0-\$v1	2-3	Values for Function Results and Expression Evaluation	No
\$a0-\$a3	4-7	Arguments	No
\$t0-\$t7	8-15	Temporaries	No
\$s0-\$s7	16-23	Saved Temporaries	Yes
\$t8-\$t9	24-25	Temporaries	No
\$k0-\$k1	26-27	Reserved for OS Kernel	No
\$gp	28	Global Pointer	Yes
\$sp	29	Stack Pointer	Yes
\$fp	30	Frame Pointer	Yes
\$ra	31	Return Address	Yes

**CS2100 Computer Organisation**  
**Tutorial #10: Pipelining**  
(Week 12: 1 – 5 November 2021)

**LumiNUS Discussion Questions**

- D1. Suppose the four stages in some 4-stage pipeline take the following timing: 2ns, 3ns, 4ns, and 2ns. Given 1000 instructions, what is the speedup (in two decimal places) of the pipelined processor compared to the non-pipelined single-cycle processor?
- D2. Let's try to understand pipeline processor by doing a detailed trace. Suppose the pipeline registers (also known as pipeline latches) store the following information:

IF / ID		ID / EX		EX / MEM		MEM / WB	
No Control Signal		MToR		MToR		MToR	
		RegWr		RegWr			
		MemRd		MemRd			
		MemWr					
		Branch		MemWr			
		RegDst		Branch		RegWr	
		ALUsrc					
		ALUop					
PC+4		PC+4		BrcTgt		MemRes	
OpCode				isZero?			
Rs		ALUOpr1		ALURes		ALURes	
Rt		ALUOpr2		ALUOpr2			
Rd		Rt					
Funct		Rd					
Imm (16)		Imm (32)		DstRNum		DstRNum	

Show the progress of the following instructions through the pipeline stages by filling in the content of pipeline registers. Note that these are the same instructions from Tutorial #5 Question 1 so that you can reuse some of the answers here.

- i.    **0x8df80000   # lw \$24, 0(\$15)       #Inst.Addr = 0x100**
- ii.   **0x1023000C   # beq \$1, \$3, 12       #Inst.Addr = 0x100**
- iii.   **0x0285c822   # sub \$25, \$20, \$5       #Inst.Addr = 0x100**

Assume that registers 1 to 31 have been initialized to a value that is equal to 101 + its register number. i.e. [\$1] = 102, [\$31] = 132 etc. You can put "X" in fields that are irrelevant for that instruction. Do note that in reality, these fields are actually generated but not utilized.

Part (i) has been worked out for you.

i. **0x8df80000 # lw \$24, 0(\$15) #Inst.Addr = 0x100**

IF / ID		ID / EX		EX / MEM		MEM / WB	
No Control Signal		MToR	1	MToR	1	MToR	1
		RegWr	1	RegWr	1		
		MemRd	1	MemRd	1		
		MemWr	0				
		Branch	0	MemWr	0	RegWr	1
		RegDst	0	Branch	0		
		ALUsrc	1			MemRes	Mem(116)
		ALUop	00	BrcTgt	X	ALURes	X
PC+4	0x104	PC+4	0x104	isZero?	X		
OpCode	0x23	ALUOpr1	116	ALURes	116	DstRNum	\$24
Rs	\$15	ALUOpr2	X	ALUOpr2	X		
Rt	\$24	Rt	\$24				
Rd	X	Rd	X	DstRNum	\$24		
Funct	X						
Imm (16)	0	Imm (32)	0				

D2. Given the following three formulas (See Lecture #20, Section 5 Performance):

$$CT_{seq} = \sum_{k=1}^N T_k$$

$$CT_{pipeline} = \max(T_k) + T_d$$

$$Speedup_{pipeline} = \frac{CT_{seq} \times InstNum}{CT_{pipeline} \times (N + InstNum - 1)}$$

For each of the following processor parameters, calculate  $CT_{seq}$ ,  $CT_{pipeline}$  and  $Speedup_{pipeline}$  (to two decimal places) for 10 instructions and for 10 million instructions.

	Stages Timing (for 5 stages, in ps)	Latency of pipeline register (in ps)
a.	300, 100, 200, 300, 100 (slow memory)	0
b.	200, 200, 200, 200, 200	40
c.	200, 200, 200, 200, 200 (ideal)	0

## Tutorial Questions

### 1. [AY2014/5 Semester 2 Exam]

Refer to the following MIPS program:

```
# register $s0 contains a 32-bit value
# register $s1 contains a non-zero 8-bit value
#           at the right most (least significant) byte
add  $t0, $s0, $zero    #inst A
add  $s2, $zero, $zero  #inst B
lp:  bne  $s2, $zero, done #inst C
     beq  $t0, $zero, done #inst D
     andi $t1, $t0, 0xFF  #inst E
     bne  $s1, $t1, nt     #inst F
     addi $s2, $s2, 1      #inst G
nt:  srl  $t0, $t0, 8      #inst H
     j    lp              #inst J
done:
```

We assume that the register **\$s0** contains **0xAFAFFAFA** and **\$s1** contains **0xFF**.

Given a 5-stage MIPS pipeline processor, for each of the parts below, give the total number of cycles needed for the first iteration of the execution from instructions **A** to **H** (i.e. excluding the “**j lp**” instruction). Remember to include the cycles needed for instruction **H** to finish the WB stage. Note that the questions are independent from each other.

- With only data forwarding mechanisms and no control hazard mechanism.
- With data forwarding and “assume not taken” branch prediction. Note that there is no early branching.  
[Recall that early branching means branch decision is made at stage 2 (Decode stage); no early branch means branch decision is made at stage 4 (Memory stage).]
- By swapping two instructions (from Instructions **A** to **H**), we can improve the performance of **early branching (with all additional forwarding paths)**. Give the two instructions that can be swapped. You only need to indicate the instruction letters in your answer.

Give the total number of cycles needed for the execution of the whole code in the worst case for each of the following assumptions. You may assume that the jump instruction (**j**) computes the address of the instruction to jump to in the MEM stage.

- With only data forwarding mechanisms and no control hazard mechanism.
- With data forwarding and “assume not taken” branch prediction. Note that there is no early branching.

2. [AY2017/8 Semester 2 Exam]

Refer to the MIPS code below. *A* and *B* are integer arrays whose base addresses are in **\$s0** and **\$s1** respectively. The arrays are of the same size *n* (number of elements). **\$s2** contains the value *n*. For this question, we will focus on the code from Instruction 1 onwards.

```
.data
A: .word 11, 9, 31, 2, 9, 1, 6, 10
B: .word 3, 7, 2, 12, 11, 41, 19, 35
n: .word 8

.text
main: la    $s0, A      # $s0 is the base address of array A
      la    $s1, B      # $s1 is the base address of array B
      la    $t0, n      # $t0 is the addr of n (size of array)
             # $s2 is the content of n
      beq   $s2, $zero, End # Inst1
      addi  $t8, $s2, -1    # Inst2
      sll   $t8, $t8, 2     # Inst3
Loop: add   $t0, $s0, $t8    # Inst4
      add   $t1, $s1, $t8    # Inst5
      lw    $t2, 0($t0)      # Inst6
      lw    $t3, 0($t1)      # Inst7
      andi  $t4, $t3, 3      # Inst8
      addi  $t4, $t4, -3     # Inst9
      beq   $t4, $zero, A1   # Inst10
      add   $t2, $t2, $t3    # Inst11
      j     A2              # Inst12
A1:   addi  $t2, $t2, 1      # Inst13
A2:   sw    $t2, 0($t0)      # Inst14
      addi  $t8, $t8, -8     # Inst15
      slt   $t7, $t8, $zero  # Inst16
      beq   $t7, $zero, Loop # Inst17

End:
```

Assuming a 5-stage MIPS pipeline system with forwarding and early branching, that is, the branch decision is made at the ID stage. No branch prediction is made and no delayed branching is used. For the jump (**j**) instruction, the computation of the target address to jump to is done at the ID stage as well.

Assume also that the first **beq** instruction begins at cycle 1.



- a. Suppose arrays *A* and *B* now each contains 200 positive integers. What is the minimum number and maximum number of instructions executed? (Consider only the above code segment from Inst1 to Inst17.)
- b. List out the instructions where some stall cycle(s) are inserted in executing that instruction in the pipeline. These include delay caused by data dependency and control hazard. You may write the instruction number InstX instead of writing out the instruction in full.
- c. How many cycles does one iteration of the loop (from Inst1 to Inst17) take if the **beq** instruction at Inst10 branches to *A1*? You have to count until the WB stage of Inst17.
- d. How many cycles does one iteration of the loop (from Inst1 to Inst17) take if the **beq** instruction at Inst10 does not branch to *A1*? You have to count until the WB stage of Inst17.

A blank pipeline chart is shown in the next page for your use. The Microsoft word version of it is also available on LumiNUS > Files and the CS2100 website.



3. [AY2020/21 Semester 2 Exam]

Study the following MIPS code on integer arrays *A* and *B* which contain the same number of elements. *\$s0* and *\$s1* contain the base addresses of *A* and *B* respectively; *\$s2* is the number of elements in array *A*; *\$s5* is *count*.

```

      add  $s5, $0, $0      # I1
      add  $t0, $0, $0      # I2
loop: slt  $t8, $t0, $s2    # I3
      beq  $t8, $0, end     # I4
      sll  $t1, $t0, 2      # I5
      add  $t3, $t1, $s0    # I6
      lw   $s3, 0($t3)      # I7
      andi $t9, $s3, 1      # I8
      beq  $t9, $0, skip    # I9
      add  $t4, $t1, $s1    # I10
      lw   $s4, 0($t4)      # I11
      sub  $s3, $s3, $s4    # I12
      sw   $s3, 0($t3)      # I13
      addi $s5, $s5, 1      # I14
skip: addi $t0, $t0, 1      # I15
      j    loop             # I16
end:

```

Assuming a 5-stage MIPS pipeline and all elements in array *A* are positive odd integers, answer the following questions. You need to count until the last stage of instruction I16.

- (a) How many cycles does this code segment take to complete its execution in the first iteration (I1 to I16) in an ideal pipeline, that is, one with no delays?

For parts (b) to (d) below, given the assumption for each part, how many additional cycles does this code segment (I1 to I16) take to complete its execution in the first iteration as compared to an ideal pipeline? (For example, if part (a) takes 12 cycles and part (b) takes 20 cycles, you are to answer part (b) with the value 8 and not 20.)

- (b) Assuming without forwarding and branch decision is made at MEM stage (stage 4). No branch prediction is made and no delayed branching is used.
- (c) Assuming without forwarding and branch decision is made at ID stage (stage 2). No branch prediction is made and no delayed branching is used.
- (d) Assuming with forwarding and branch decision is made at ID stage (stage 2). Branch prediction is made where the branch is predicted not taken, and no delayed branching is used.
- (e) Assuming the setting in part (d) above and you are not allowed to modify any of the instructions, is it possible to reduce the additional delay cycles in part (d) by rearranging some instructions, and if possible, by how many cycles? Explain your answer. (Answer with no explanation will not be awarded any mark.)

**CS2100 Computer Organisation**  
**Tutorial #10: Pipelining**  
**Answers to Selected Questions**

---

**Tutorial Questions**

1. [AY2014/5 Semester 2 Exam]

Refer to the following MIPS program:

```
# register $s0 contains a 32-bit value
# register $s1 contains a non-zero 8-bit value
#           at the right most (least significant) byte
add  $t0, $s0, $zero      #inst A
add  $s2, $zero, $zero    #inst B
lp:  bne $s2, $zero, done  #inst C
     beq $t0, $zero, done  #inst D
     andi $t1, $t0, 0xFF   #inst E
     bne $s1, $t1, nt      #inst F
     addi $s2, $s2, 1      #inst G
nt:  srl $t0, $t0, 8       #inst H
     j    lp              #inst J
done:
```

We assume that the register **\$s0** contains **0xAFAFFAFA** and **\$s1** contains **0xFF**.

Given a 5-stage MIPS pipeline processor, for each of the parts below, give the total number of cycles needed for the first iteration of the execution from instructions **A** to **H** (i.e. excluding the “**j lp**” instruction). Remember to include the cycles needed for instruction **H** to finish the WB stage. Note that the questions are independent from each other.

- a. With only data forwarding mechanisms and no control hazard mechanism.
- b. With data forwarding and “assume not taken” branch prediction. Note that there is no early branching.  
[Recall that early branching means branch decision is made at stage 2 (Decode stage); no early branch means branch decision is made at stage 4 (Memory stage).]
- c. By swapping two instructions (from Instructions **A** to **H**), we can improve the performance of **early branching (with all additional forwarding paths)**. Give the two instructions that can be swapped. You only need to indicate the instruction letters in your answer.

Give the total number of cycles needed for the execution of the whole code in the worst case for each of the following assumptions. You may assume that the jump instruction (**j**) computes the address of the instruction to jump to in the MEM stage.

- d. With only data forwarding mechanisms and no control hazard mechanism.
- e. With data forwarding and “assume not taken” branch prediction. Note that there is no early branching.

## Answers:

### (a) 20 cycles

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
add	F	D	E	M	W															
add		F	D	E	M	W														
bne			F	D	E	M	W													
beq							F	D	E	M	W									
andi											F	D	E	M	W					
bne												F	D	E	M	W				
addi		The addi instruction is not executed.																		
srl																F	D	E	M	W

### (b) 14 cycles

	1	2	3	4	5	6	7	8	9	10	11	12	13	14
add	F	D	E	M	W									
add		F	D	E	M	W								
bne			F	D	E	M	W							
beq				F	D	E	M	W						
andi					F	D	E	M	W					
bne						F	D	E	M	W				
addi							F	D	E	*	*			
srl								F	D	*	*	*		
j									F	*	*	*	*	
srl										F	D	E	M	W

Cost of wrong prediction

### (c) Swap instructions A and B to reduce the delay between instructions B and C.

```

    add    $s2, $zero, $zero    #inst B
    add    $t0, $s0, $zero      #inst A
lp:   bne   $s2, $zero, done     #inst C
      beq   $t0, $zero, done     #inst D
      andi  $t1, $t0, 0xFF       #inst E
      bne   $s1, $t1, nt         #inst F
      addi  $s2, $s2, 1          #inst G
nt:   srl   $t0, $t0, 8          #inst H
      j     lp                  #inst J
done:

```

(d)

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21
add	F	D	E	M	W																
add		F	D	E	M	W															
bne			F	D	E	M	W														
beq							F	D	E	M	W										
andi											F	D	E	M	W						
bne												F	D	E	M	W					
<del>addi</del>																					
srl															F	D	E	M	W		
j																F	D	E	M	W	
bne																					F

In the worst case, 4 bytes of the data are examined → 4 iterations. The loop from Instructions C to J (one iteration) takes 18 cycles (not counting the WB stage of the j instruction which overlaps with the bne instruction). There are 2 cycles before the first iteration, and 9 cycles for Instructions C and D in the fifth iteration.

Therefore, total = 2 + (4×18) + 9 = **83 cycles**.

(e)

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
add	F	D	E	M	W										
add		F	D	E	M	W									
bne			F	D	E	M	W								
beq				F	D	E	M	W							
andi					F	D	E	M	W						
bne						F	D	E	M	W					
addi							F	D	E	*	*				
srl								F	D	*	*	*			
j									F	*	*	*	*		
srl										F	D	E	M	W	
j											F	D	E	M	W
bne															F

In the worst case, 4 bytes of the data are examined → 4 iterations. The loop from Instructions C to J (one iteration) takes 12 cycles (not counting the WB stage of the j instruction which overlaps with the bne instruction). There are 2 cycles before the first iteration, and 6 cycles for Instructions C and D in the fifth iteration.

Therefore, total = 2 + (4×12) + 6 = **56 cycles**.

2. [AY2017/8 Semester 2 Exam]

Refer to the MIPS code below. *A* and *B* are integer arrays whose base addresses are in *\$s0* and *\$s1* respectively. The arrays are of the same size *n* (number of elements). *\$s2* contains the value *n*. For this question, we will focus on the code from Instruction 1 onwards.

```
.data
A: .word 11, 9, 31, 2, 9, 1, 6, 10
B: .word 3, 7, 2, 12, 11, 41, 19, 35
n: .word 8
.text
main: la    $s0, A      # $s0 is the base address of array A
      la    $s1, B      # $s1 is the base address of array B
      la    $t0, n      # $t0 is the addr of n (size of array)
             # $s2 is the content of n
      beq   $s2, $zero, End # Inst1
      addi  $t8, $s2, -1    # Inst2
      sll   $t8, $t8, 2     # Inst3
Loop: add   $t0, $s0, $t8   # Inst4
      add   $t1, $s1, $t8   # Inst5
      lw    $t2, 0($t0)     # Inst6
      lw    $t3, 0($t1)     # Inst7
      andi  $t4, $t3, 3     # Inst8
      addi  $t4, $t4, -3    # Inst9
      beq   $t4, $zero, A1  # Inst10
      add   $t2, $t2, $t3   # Inst11
      j     A2              # Inst12
A1:   addi  $t2, $t2, 1     # Inst13
A2:   sw    $t2, 0($t0)     # Inst14
      addi  $t8, $t8, -8    # Inst15
      slt   $t7, $t8, $zero # Inst16
      beq   $t7, $zero, Loop # Inst17
End:
```

Assuming a 5-stage MIPS pipeline system with forwarding and early branching, that is, the branch decision is made at the ID stage. No branch prediction is made and no delayed branching is used. For the jump (j) instruction, the computation of the target address to jump to is done at the ID stage as well.

Assume also that the first **beq** instruction begins at cycle 1.

- Suppose arrays *A* and *B* now each contains 200 positive integers. What is the minimum number and maximum number of instructions executed? (Consider only the above code segment from Inst1 to Inst17.)
- List out the instructions where some stall cycle(s) are inserted in executing that instruction in the pipeline. These include delay caused by data dependency and control hazard. You may write the instruction number InstX instead of writing out the instruction in full.
- How many cycles does one iteration of the loop (from Inst1 to Inst17) take if the **beq** instruction at Inst10 branches to A1? You have to count until the WB stage of Inst17.
- How many cycles does one iteration of the loop (from Inst1 to Inst17) take if the **beq** instruction at Inst10 does not branch to A1? You have to count until the WB stage of Inst17.

**Answers:**

**The code does this:**

```
int i;
for (i=n-1; i>=0; i-=2) {
    if (B[i]%4 == 3)
        A[i] = A[i] + 1;
    else
        A[i] = A[i] + B[i];
}
```

**(a) Minimum =  $3 + 100 \times 12 = 1203$**

**Maximum =  $3 + 100 \times 13 = 1303$**

In the loop (Inst4 to Inst17), there are two paths after Inst10: one that skips Inst11 and Inst12, and the other skips Inst13.

**(b) Due to control: Inst2, Inst4, Inst11, Inst13**

**Due to data: Inst8, Inst10, Inst17**

**(c) 24 cycles**

**(d) 26 cycles**



3. [AY2020/21 Semester 2 Exam]

Study the following MIPS code on integer arrays *A* and *B* which contain the same number of elements. *\$s0* and *\$s1* contain the base addresses of *A* and *B* respectively; *\$s2* is the number of elements in array *A*; *\$s5* is *count*.

```

        add    $s5, $0, $0    # I1
        add    $t0, $0, $0    # I2
loop:   slt    $t8, $t0, $s2  # I3
        beq    $t8, $0, end    # I4
        sll    $t1, $t0, 2     # I5
        add    $t3, $t1, $s0   # I6
        lw     $s3, 0($t3)     # I7
        andi   $t9, $s3, 1     # I8
        beq    $t9, $0, skip   # I9
        add    $t4, $t1, $s1   # I10
        lw     $s4, 0($t4)     # I11
        sub    $s3, $s3, $s4   # I12
        sw     $s3, 0($t3)     # I13
        addi   $s5, $s5, 1     # I14
skip:   addi   $t0, $t0, 1     # I15
        j      loop           # I16
end:

```

Assuming a 5-stage MIPS pipeline and all elements in array *A* are positive odd integers, answer the following questions. You need to count until the last stage of instruction I16.

- (a) How many cycles does this code segment take to complete its execution in the first iteration (I1 to I16) in an ideal pipeline, that is, one with no delays?

For parts (b) to (d) below, given the assumption for each part, how many additional cycles does this code segment (I1 to I16) take to complete its execution in the first iteration as compared to an ideal pipeline? (For example, if part (a) takes 12 cycles and part (b) takes 20 cycles, you are to answer part (b) with the value 8 and not 20.)

- (b) Assuming without forwarding and branch decision is made at MEM stage (stage 4). No branch prediction is made and no delayed branching is used.
- (c) Assuming without forwarding and branch decision is made at ID stage (stage 2). No branch prediction is made and no delayed branching is used.
- (d) Assuming with forwarding and branch decision is made at ID stage (stage 2). Branch prediction is made where the branch is predicted not taken, and no delayed branching is used.
- (e) Assuming the setting in part (d) above and you are not allowed to modify any of the instructions, is it possible to reduce the additional delay cycles in part (d) by rearranging some instructions, and if possible, by how many cycles? Explain your answer. (Answer with no explanation will not be awarded any mark.)

**Answers:**

(a)  $16 + 5 - 1 = 20$  cycles.

Delays are highlighted under the columns (b), (c), (d) for parts (b),(c),(d) below respectively.

(b) **+24** cycles.

(c) **+20** cycles.

(d) **+4** cycles.

				(b)	(c)	(d)
	add	\$s5, \$0, \$0	# I1			
	add	\$t0, \$0, \$0	# I2			
loop:	slt	\$t8, \$t0, \$s2	# I3	+2	+2	
	beq	\$t8, \$0, end	# I4	+2	+2	+1
	sll	\$t1, \$t0, 2	# I5	+3	+1	
	add	\$t3, \$t1, \$s0	# I6	+2	+2	
	lw	\$s3, 0(\$t3)	# I7	+2	+2	
	andi	\$t9, \$s3, 1	# I8	+2	+2	+1
	beq	\$t9, \$0, skip	# I9	+2	+2	+1
	add	\$t4, \$t1, \$s1	# I10	+3	+1	
	lw	\$s4, 0(\$t4)	# I11	+2	+2	
	sub	\$s3, \$s3, \$s4	# I12	+2	+2	+1
	sw	\$s3, 0(\$t3)	# I13	+2	+2	
	addi	\$s5, \$s5, 1	# I14			
skip:	addi	\$t0, \$t0, 1	# I15			
	j	loop	# I16			
end:						
<b>Total:</b>				<b>+24</b>	<b>+20</b>	<b>+4</b>

(e) Other answers possible. Example:

Move I14 (addi \$s5, \$s5, 1) to between I11 (lw \$s4, 0(\$t4)) and I12 (sub \$s3, \$s3, \$s4) to remove the 1 cycle of delay at I12.

**CS2100 Computer Organisation**  
**Tutorial #11: Cache**  
(Week 13: 8 – 12 November 2021)

**LumiNUS Discussion Question**

D1. [CS2100 AY2007/8 Semester 2 Exam Question]

A machine with a word size of 16 bits and address width of 32 bits has a **direct-mapped cache** with 16 blocks and a block size of 2 words, initially empty.

- (a) Given a sequence of memory references as shown below, where each reference is given as a byte address in both decimal and hexadecimal forms, indicate whether the reference is a hit (H) or a miss (M).

Memory address		Hit (H) or Miss (M)?	(For reference)
(in decimal)	(in hexadecimal)		
4	0x4	<b>M</b>	0000 ... 0000 0100
92	0x5C		0000 ... 0101 1100
7	0x7		0000 ... 0000 0111
146	0x92		0000 ... 1001 0010
30	0x1E		0000 ... 0001 1110
95	0x5F		0000 ... 0101 1111
176	0xB0		0000 ... 1011 0000
93	0x5D		0000 ... 0101 1101
145	0x91		0000 ... 1001 0001
264	0x108		0000 ... 1 0000 1000
6	0x6		0000 ... 0000 0110

- (b) Given the above sequence of memory references, fill in the final contents of the cache. Use the notation  $M[i]$  to denote the word starting at memory address  $i$ , where  $i$  is in hexadecimal. If a block is replaced, cross out the content in the cache and write the new content over it.

Index	Tag value	Word 0	Word 1
0			
1			
2			
3			
4			
5			
6			
7			
8			
9			
10			
11			
12			
13			
14			
15			

### ***Tutorial Questions***

1. Here is a series of address references in decimal: 4, 16, 32, 20, 80, 68, 76, 224, 36, 44, 16, 172, 20, 24, 36, and 68 in a MIPS machine. Assuming a **direct-mapped cache** with 16 one-word blocks that is initially empty, label each address reference as a hit or miss and show the content of the cache.

You may write the data word starting at memory address X as M[X]. (For example, data word starting at memory address 12 is written as M[12]. This implies that the word includes the 4 bytes of data at addresses 12, 13, 14 and 15.) You may write the tag values as decimal numbers. If a block is replaced in the cache, cross out the corresponding content in the cache, and write the new content over it.

2. Use the series of references given in question 1 above: 4, 16, 32, 20, 80, 68, 76, 224, 36, 44, 16, 172, 20, 24, 36, and 68 in a MIPS machine. Assuming a **two-way set-associative cache** with two-word blocks and a total size of 16 words that is initially empty, label each address reference as a hit or miss and show the content of the cache. Assume **LRU** replacement policy.

You may write the data word starting at memory address X as M[X]. (For example, data word starting at memory address 12 is written as M[12]. This implies that the word includes the 4 bytes of data at addresses 12, 13, 14 and 15.) You may write the tag values as decimal numbers. If a block is replaced in the cache, cross out the corresponding content in the cache, and write the new content over it.

3. Although we use only data memory as example in the cache lecture, the principle covered is equally applicable to the instruction memory. This question takes a look at both the instruction cache and data cache.

The code below is from Tutorial 3 Question 1 (*palindrome checking*) with the following variable mappings:

low → \$s0, high → \$s1, matched → \$s3, base of string[] → \$s4, size → \$s5

#	Code	Comment
i0	[some instruction]	
i1	addi \$s0, \$zero, 0	# low = 0
i2	addi \$s1, \$s5, -1	# high = size-1
i3	addi \$s3, \$zero, 1	# matched = 1
	loop:	
i4	slt \$t0, \$s0, \$s1	# (low < high)?
i5	beq \$t0, \$zero, exit	# exit if (low >= high)
i6	beq \$s3, \$zero, exit	# exit if (matched == 0)
i7	add \$t1, \$s4, \$s0	# address of string[low]
i8	lb \$t2, 0(\$t1)	# t2 = string[low]
i9	addi \$t3, \$s4, \$s1	# address of string[high]
i10	lb \$t4, 0(\$t3)	# t4 = string[high]
i11	beq \$t2, \$t4, else	
i12	addi \$s3, \$zero, 0	# matched = 0
i13	j endW	# can be "j loop"
	else:	
i14	addi \$s0, \$s0, 1	# low++
i15	addi \$s1, \$s1, -1	# high--
	endW:	
i16	j loop	# end of while
	exit:	
i17	[some instruction]	

Parts (a) to (d) assume that instruction i0 is stored at memory address 0x0.

- (a) Instruction cache: **Direct mapped with 2 blocks of 16 bytes each** (i.e. each block can hold 4 consecutive instructions).

Starting with an empty cache, the fetching of instruction i1 will cause a cache miss. After the cache miss is resolved, we now have the following instructions in the instruction cache:

Instruction Cache Block 0	[i0, i1, i2, i3]
Instruction Cache Block 1	[empty]

Fetching of i2 and i3 are all cache hits as they can be found in the cache.

Assuming the string being checked is a palindrome. Show the instruction cache block content **at the end of the 1<sup>st</sup> iteration (i.e. up to instruction i16).**

(b) If the loop is executed for a total of 10 iterations, what is the total number of cache hits (i.e. after the 10<sup>th</sup> "j loop" is fetched)?

(c) Suppose we change the instruction cache to:

- **Direct mapped with 4 blocks of 8 bytes each** (i.e. each block can hold 2 consecutive instructions).

Assuming the string being checked is a palindrome. Show the instruction cache block content **at the end of the 1<sup>st</sup> iteration (i.e. up to instruction i16).**

Instruction Cache Block 0	
Instruction Cache Block 1	
Instruction Cache Block 2	
Instruction Cache Block 3	

(d) If the loop is executed for a total of 10 iterations, what is the total number of cache hits (i.e. after the 10<sup>th</sup> "j loop" is fetched)?

Let us now turn to the study of **data cache**. We will assume the following scenario for parts (e) to (g):

- The string being checked is **64-character long**. The first character is located at location **0x1000**.
- The string is a palindrome (i.e. it will go through 32 iterations of the code).

(e) Given a **direct mapped data cache with 2 cache blocks, each block is 8 bytes**, what is the final content of the data cache at the end of the code execution (after the code failed the beq at i5)? Use **s[X..Y]** to indicate the data **string[X]** to **string[Y]**.

Data Cache Block #0	
Data Cache Block #1	

(f) What is the hit rate of (e)? Give your answer in a fraction or a percentage correct to two decimal places.

(g) Suppose the string is now **72-character long**, the first character is still located at location **0x1000** and the string is still a palindrome, what is the hit rate at the end of the execution?

## CS2100 Computer Organisation

### Tutorial #11: Cache

#### Answers to Selected Questions

#### Tutorial Questions

2. Use the series of references given in question 1 above: 4, 16, 32, 20, 80, 68, 76, 224, 36, 44, 16, 172, 20, 24, 36, and 68 in a MIPS machine. Assuming a **two-way set-associative cache** with two-word blocks and a total size of 16 words that is initially empty, label each address reference as a hit or miss and show the content of the cache. Assume **LRU** replacement policy.

You may write the data word starting at memory address X as M[X]. (For example, data word starting at memory address 12 is written as M[12]. This implies that the word includes the 4 bytes of data at addresses 12, 13, 14 and 15.) You may write the tag values as decimal numbers. If a block is replaced in the cache, cross out the corresponding content in the cache, and write the new content over it.

#### Answer:

Since this is a MIPS machine, a word consists of 4 bytes or 32 bits.

Should first work out the tag, set index, and offset fields:

27 bits	2 bits	3
Tag	Set Index	Offset

```

4:    00...000  00  100 ← Miss
16:   00...000  10  000 ← Miss
32:   00...001  00  000 ← Miss
20:   00...000  10  100 ← Hit
80:   00...010  10  000 ← Miss
68:   00...010  00  100 ← Miss
76:   00...010  01  100 ← Miss
224:  00...111  00  000 ← Miss
36:   00...001  00  100 ← Miss
44:   00...001  01  100 ← Miss
16:   00...000  10  000 ← Hit
172:  00...101  01  100 ← Miss
20:   00...000  10  100 ← Hit
24:   00...000  11  000 ← Miss
36:   00...001  00  100 ← Hit
68:   00...010  00  100 ← Miss
    
```

Cache set	Valid bit	Tag	Word0	Word1	Valid bit	Tag	Word0	Word1
0	0 1	0 2 1	M[0] M[64] M[32]	M[4] M[68] M[36]	0 1	4 7 2	M[32] M[224] M[64]	M[36] M[228] M[68]
1	0 1	2 5	M[72] M[168]	M[76] M[172]	0 1	1	M[40]	M[44]
2	0 1	0	M[16]	M[20]	0 1	2	M[80]	M[84]
3	0 1	0	M[24]	M[28]	0			

3. Although we use only data memory as example in the cache lecture, the principle covered is equally applicable to the instruction memory. This question takes a look at both the instruction cache and data cache.

The code below is from Tutorial 8 Question 1 (*palindrome checking*) with the following variable mappings:

low → \$s0, high → \$s1, matched → \$s3, base of string[] → \$s4, size → \$s5

#	Code	Comment
i0	[some instruction]	
i1	addi \$s0, \$zero, 0	# low = 0
i2	addi \$s1, \$s5, -1	# high = size-1
i3	addi \$s3, \$zero, 1	# matched = 1
	loop:	
i4	slt \$t0, \$s0, \$s1	# (low < high)?
i5	beq \$t0, \$zero, exit	# exit if (low >= high)
i6	beq \$s3, \$zero, exit	# exit if (matched == 0)
i7	add \$t1, \$s4, \$s0	# address of string[low]
i8	lb \$t2, 0(\$t1)	# t2 = string[low]
i9	addi \$t3, \$s4, \$s1	# address of string[high]
i10	lb \$t4, 0(\$t3)	# t4 = string[high]
i11	beq \$t2, \$t4, else	
i12	addi \$s3, \$zero, 0	# matched = 0
i13	j endW	# can be "j loop"
	else:	
i14	addi \$s0, \$s0, 1	# low++
i15	addi \$s1, \$s1, -1	# high--
	endW:	
i16	j loop	# end of while
	exit:	
i17	[some instruction]	

Parts (a) to (d) assume that instruction i0 is stored at memory address 0x0.

- (a) Instruction cache: **Direct mapped with 2 blocks of 16 bytes each** (i.e. each block can hold 4 consecutive instructions).

Starting with an empty cache, the fetching of instruction i1 will cause a cache miss. After the cache miss is resolved, we now have the following instructions in the instruction cache:

Instruction Cache Block 0	[i0, i1, i2, i3]
Instruction Cache Block 1	[empty]

Fetching of i2 and i3 are all cache hits as they can be found in the cache.

Assuming the string being checked is a palindrome. Show the instruction cache block content **at the end of the 1<sup>st</sup> iteration (i.e. up to instruction i16)**.



**Answer:**

Instruction Cache Block 0	[i16, .....]
Instruction Cache Block 1	[i12, i13, i14, i15]

Working: Instructions executed = i1 to i11, i14 to i16

Block #0, Cache index = 0	[i0, i1, i2, i3]
Block #1, Cache index = 1	[i4, i5, i6, i7]
Block #2, Cache index = 0	[i8, i9, i10, i11]
Block #3, Cache index = 1	[i12, i13, i14, i15]
Block #4, Cache index = 0	[i16, other....]

(b) If the loop is executed for a total of 10 iterations, what is the total number of cache hits (i.e. after the 10<sup>th</sup> "j loop" is fetched)?

**Answer:**

Working (1<sup>st</sup> Iteration):

i1	i2	i3	i4	i5	i6	i7	i8	i9	i10	i11	i14	i15	i16
M	H	H	M	H	H	H	M	H	H	H	M	H	M

Working (2<sup>nd</sup> iteration onward):

i4	i5	i6	i7	i8	i9	i10	i11	i14	i15	i16
M	H	H	H	M	H	H	H	M	H	M

Total hits = 9 (1<sup>st</sup> iteration) + 7×9 (remaining 9 iterations) = **72**

(c) Suppose we change the instruction cache to:

- **Direct mapped with 4 blocks of 8 bytes each** (i.e. each block can hold 2 consecutive instructions).

Assuming the string being checked is a palindrome. Show the instruction cache block content **at the end of the 1<sup>st</sup> iteration (i.e. up to instruction i16)**.

**Answer:**

Instruction Cache Block 0	[i16, ...]
Instruction Cache Block 1	[i10, i11]
Instruction Cache Block 2	[i4, i5]
Instruction Cache Block 3	[i14, i15]

Working:

First, find out the block information for the full code:

Block #0, Cache index = 0	[i0, i1]
Block #1, Cache index = 1	[i2, i3]
Block #2, Cache index = 2	[i4, i5]
Block #3, Cache index = 3	[i6, i7]
Block #4, Cache index = 0	[i8, i9]
Block #5, Cache index = 1	[i10, i11]
Block #6, Cache index = 2	[i12, i13]
Block #7, Cache index = 3	[i14, i15]
Block #8, Cache index = 0	[i16, ...]

Second, use the execution pattern to find out what is accessed, since we execute i1 to i11 (Block #0 to Block #5) then i14 to i16 (Block #7 and Block #8), we get the final cache content as shown. You should note that Block #6 [i12, i13] is not accessed in this particular execution.

- (d) If the loop is executed for a total of 10 iterations, what is the total number of cache hits (i.e. after the 10<sup>th</sup> "j loop" is fetched)?

**Answer:**

Working (1<sup>st</sup> Iteration):

i1	i2	i3	i4	i5	i6	i7	i8	i9	i10	i11	i14	i15	i16
M	M	H	M	H	M	H	M	H	M	H	M	H	M

Working (2<sup>nd</sup> iteration onward):

i4	i5	i6	i7	i8	i9	i10	i11	i14	i15	i16
H	H	M	H	M	H	H	H	M	H	M

**Total hits = 6 (1<sup>st</sup> iteration) + 7×9 (remaining 9 iterations) = 69**

Let us now turn to the study of **data cache**. We will assume the following scenario for parts (e) to (g):

- The string being checked is **64-character long**. The first character is located at location **0x1000**.
- The string is a palindrome (i.e. it will go through 32 iterations of the code).

- (e) Given a **direct mapped data cache with 2 cache blocks, each block is 8 bytes**, what is the final content of the data cache at the end of the code execution (after the code failed the beq at i5)? Use **s[X..Y]** to indicate the data **string[X]** to **string[Y]**.

**Answer:**

Data Cache Block #0	<b>s[32..39]</b>
Data Cache Block #1	<b>s[24..31]</b>

Access patterns = s[0], s[63], s[1], s[62], ..., s[31], s[32]

Blocks information (blocks that can go into the same cache location are listed together):

Cache index = 0	s[0..7] [16..23] [32..39] [48..55]
Cache index = 1	s[8..15] [24..31] [40..47] [56..63]

- (f) What is the hit rate of (e)? Give your answer in a fraction or a percentage correct to two decimal places.

**Answer:**

Observation: the access pattern nicely alternates between Block0-Block1 and Block1-Block0. So, in general, other than the first miss to bring in a block, the remaining 7 accesses on the block are all hits.

Hence, hit rate = **7/8** or **87.50%**

- (g) Suppose the string is now **72-character long**, the first character is still located at location **0x1000** and the string is still a palindrome, what is the hit rate at the end of the execution?

**Answer:**

Access patterns = s[0], s[71], s[1], s[70], ..., s[35], s[36]

Blocks information (blocks that can go into the same cache location are listed together):

Cache index = 0	s[0..7] [16..23] [32..39] [48..55] [64..71]
Cache index = 1	s[8..15] [24..31] [40..47] [56..63]

Observation: the access pattern is either Block0-Block0 or Block1-Block1. So, every access is a miss, except the last block [32..39]! This is an example of **cache thrashing** (you can imagine the cache is "beaten up" pretty badly ☺).

Hence, hit rate = **7/72** (the last 7 accesses on block [32..39]) or **9.72%**